

# ARTIFICIAL INTELLIGENCE NEEDS REAL INTELLIGENCE

## Clustering II

Professor: Tianyuan ZHANG  
tianyuan.zhang@kedgebs.com



kedge.edu

EFMD EQUIS ACCREDITED AACSB ACCREDITED AMBA ACCREDITED

28/11/2023

Machine Learning Part I by Tianyuan ZHANG

1

# Recap of Previous Session

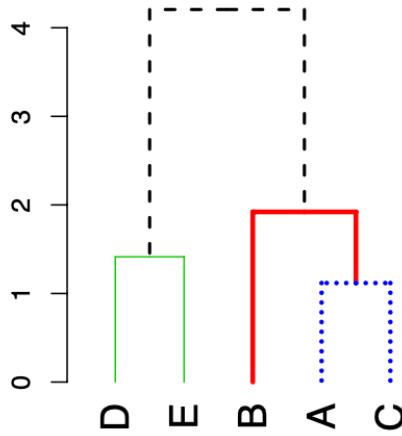
- Unsupervised Learning
  - Learn **hidden patterns or structures** exclusively from **unlabeled data**
  - Three types of common unsupervised learning algorithms
    - **Clustering:**
      - Natural groups → Clusters
    - **Association rule mining:**
      - Frequent co-occurrence → Association rules
    - **Dimensionality reduction:**
      - Simplifications → Compressed data

# Recap of Previous Session

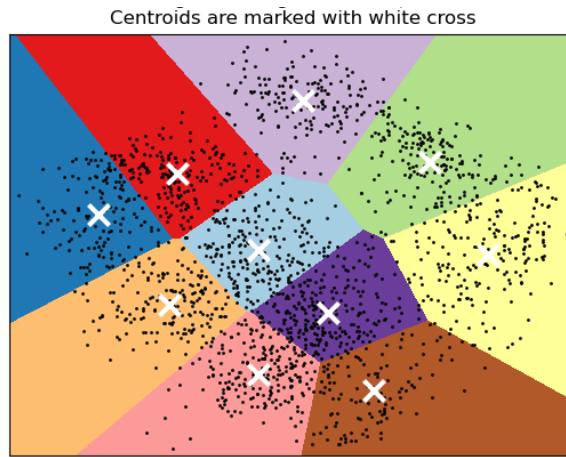
- Clustering
  - Grouping **unlabeled** data into different natural **clusters**
  - Data points within the same cluster are similar.
  - Data points in different clusters are different.
  - The clusters are not pre-defined by humans before clustering
  - We need to interpret the clustering results after clustering

# Recap of Previous Session

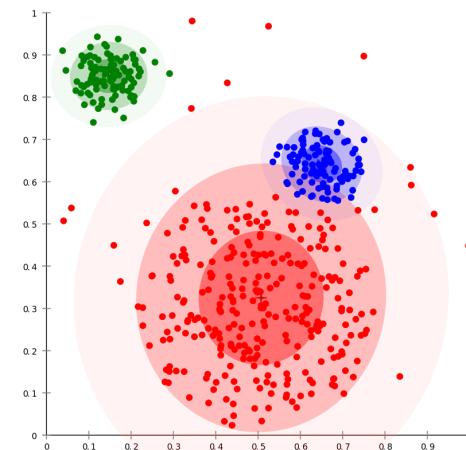
- Different types of clustering algorithms



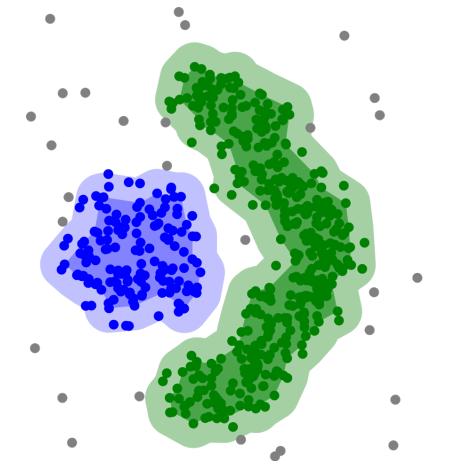
Connectivity-based  
Hierarchical clustering



Centroid-based clustering  
Partitioning method



Distribution-based  
clustering



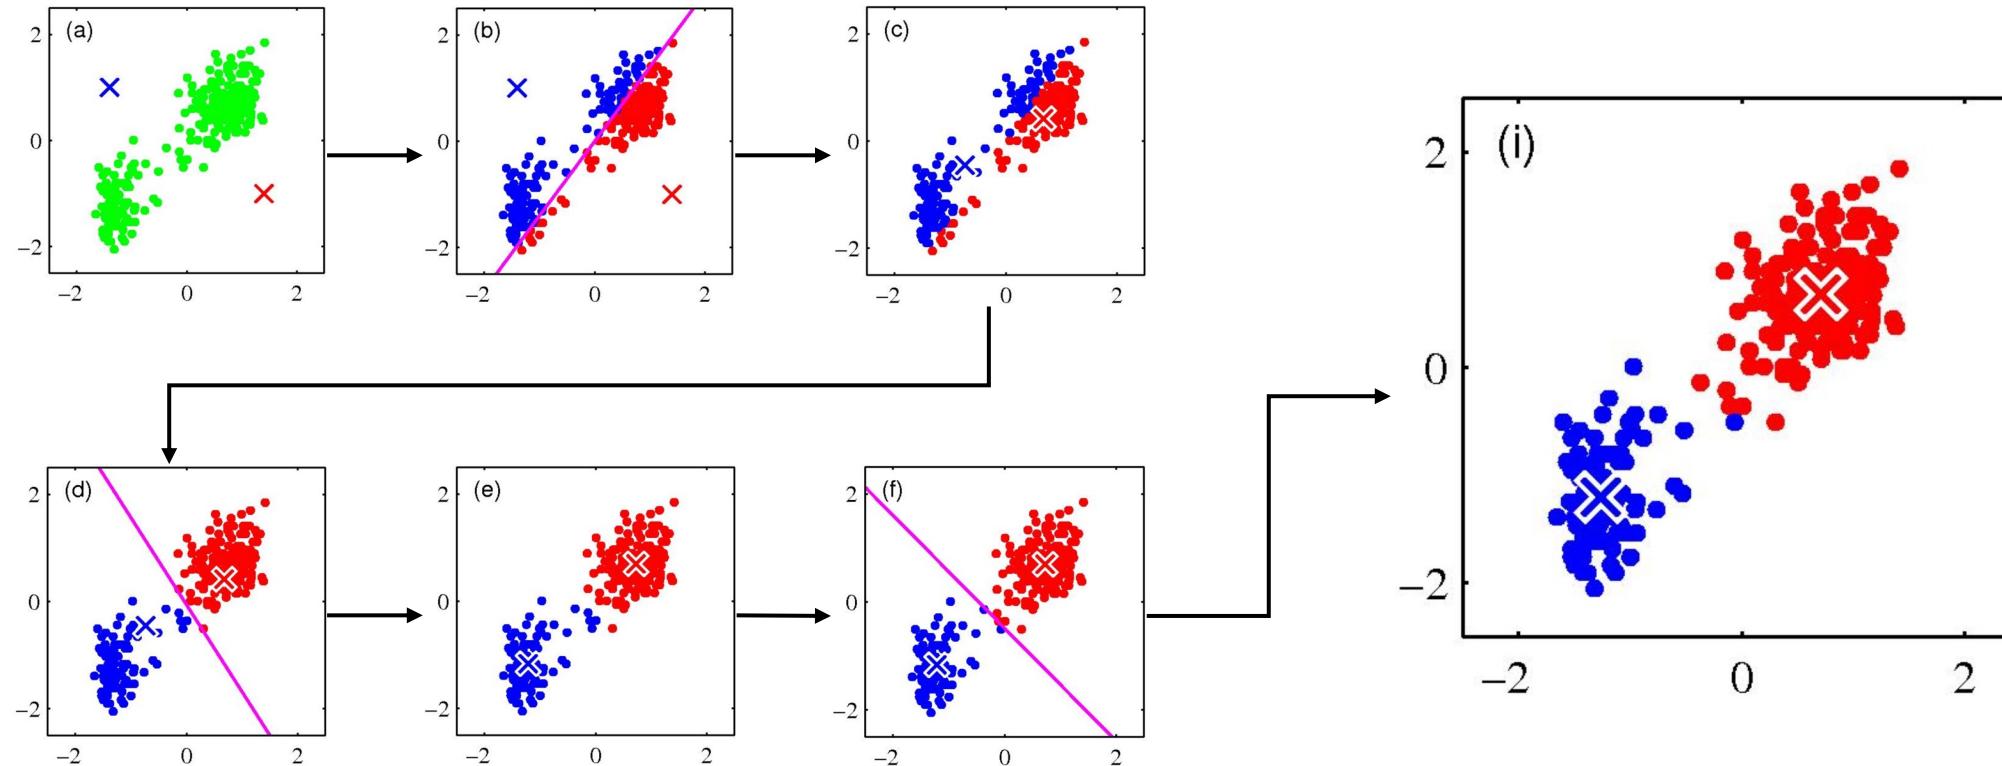
Density-based  
clustering

# Recap of Previous Session

- K-Means Clustering
  - An **iterative centroid-based** clustering algorithm
  - A **partitioning** method that divides the feature space into K distinct, non-overlapping subsets, so-called clusters
  - Each cluster is represented by a **centroid**
    - Which cluster a point belongs to depends on which centroid it is closest to
  - **Steps**
    1. Initialize K centroids randomly
    2. Assign each data point to the nearest centroid, forming K clusters
    3. Recalculate the centroids as the mean of all points in each cluster
    4. Repeat steps 2 and 3 until convergence

# Recap of Previous Session

- K-Means Clustering

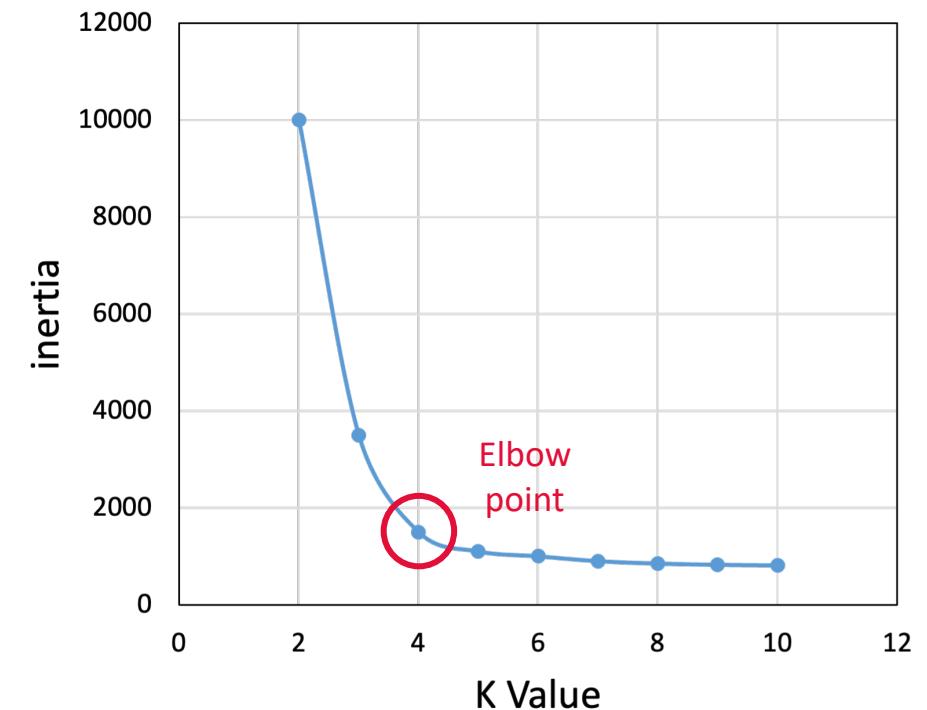


# Recap of Previous Session

- K-Means Clustering
  - Requirements
    - The clusters need to be convex and isotropic.
    - The clusters should have equal variance.
    - The clusters should have similar size.
    - Perform feature scaling before using K-Means algorithm.
    - Select a proper K value

# Recap of Previous Session

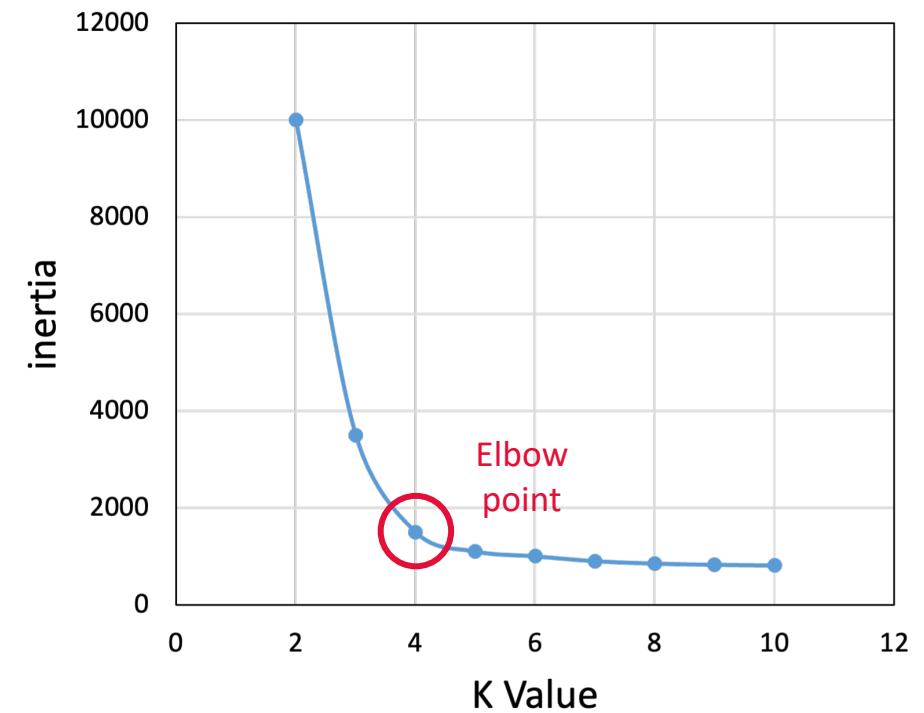
- K-Means Clustering
  - Select optimal K value
    - Elbow methods
    - $\text{inertia} = \sum_{i=0}^n \min_{\mu_j \in C} (\|x_i - \mu_j\|^2)$
    - If there is no clear elbow point, use other metrics
      - Silhouette Coefficient
      - Calinski-Harabasz Index
      - Davies-Bouldin Index



# Recap of Previous Session

- K-Means Clustering
  - Select optimal K value
    - Elbow methods
    - $\text{inertia} = \sum_{i=0}^n \min_{\mu_j \in C} (\|x_i - \mu_j\|^2)$
    - If there is no clear elbow point, use other metrics
      - Silhouette Coefficient
      - Calinski-Harabasz Index
      - Davies-Bouldin Index

*Lower value is better*

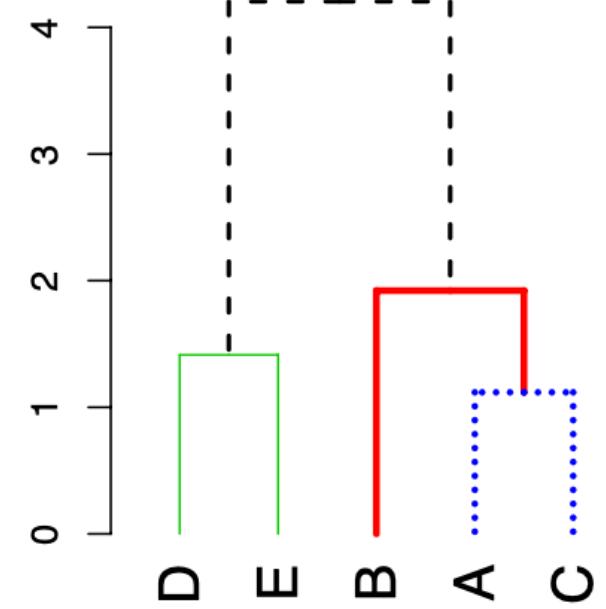
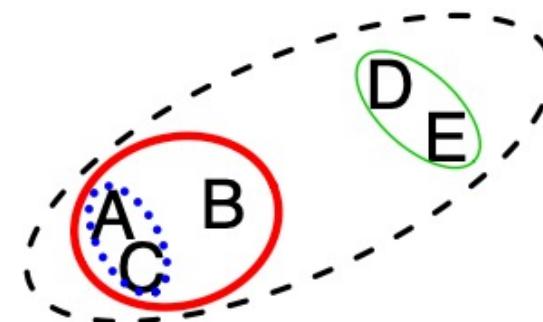


# Outline

- Hierarchical Clustering

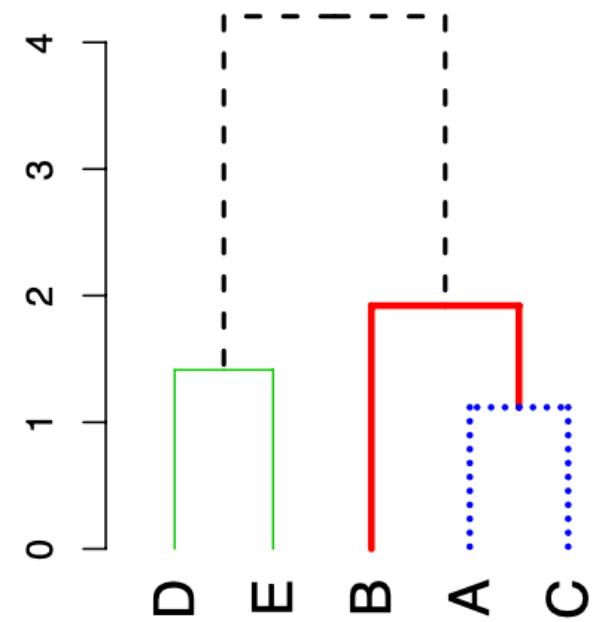
# Hierarchical Clustering

- **Connectivity-based** clustering algorithm
- Cluster is defined as sets of points that are closely connected to each other
- Build **nested** clusters by merging or splitting them successively
- Result is a hierarchy of clusters, represented by a **dendrogram** or a tree



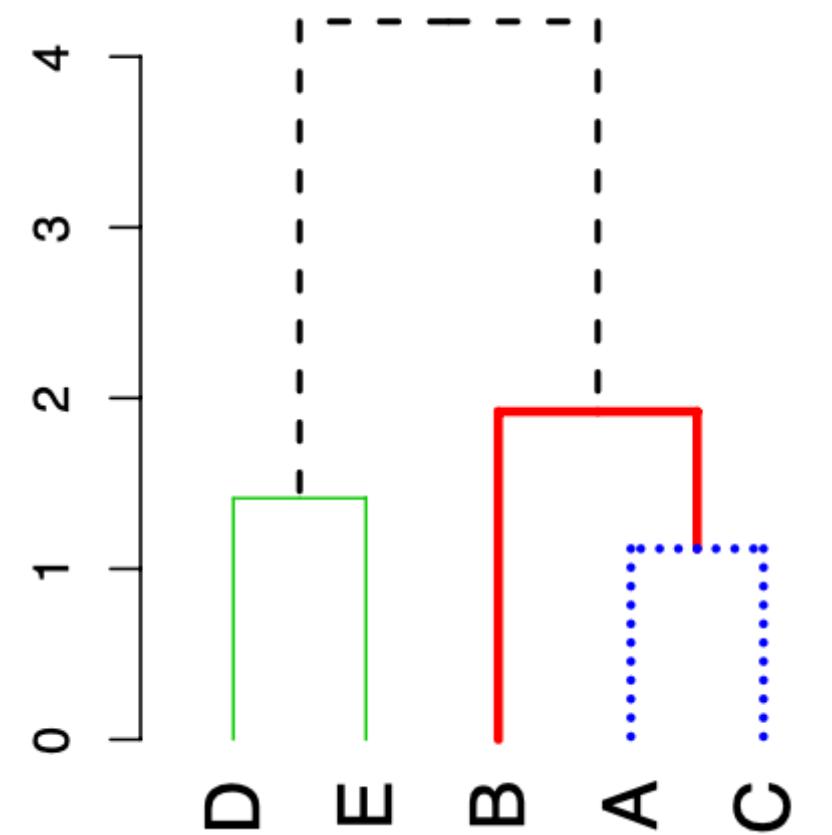
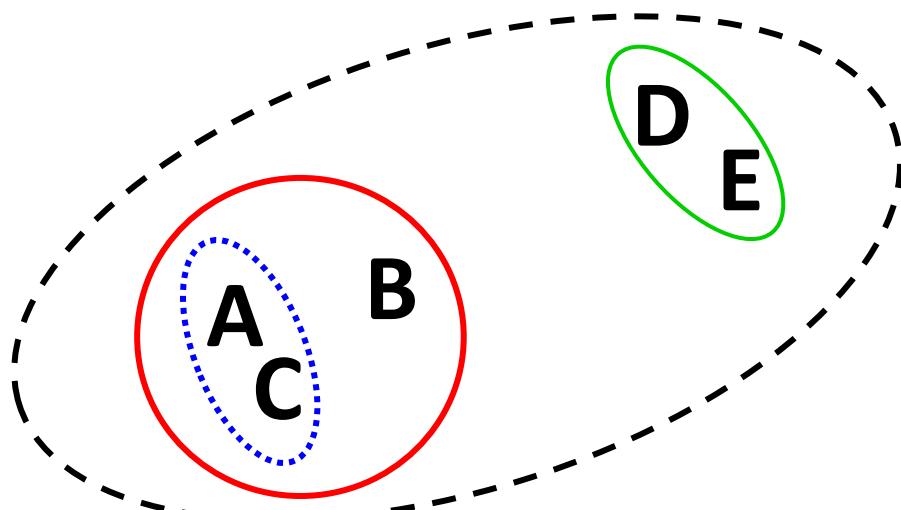
# Hierarchical Clustering

- Two types of hierarchical clustering algorithm
  - Agglomerative clustering
    - Bottom-up approach
    - Every point starts in its own cluster
    - Pairs of clusters are merged as one moves up the hierarchy
  - Divisive clustering
    - Top-down approach
    - All points start in one cluster
    - Splits are performed recursively as one moves down the hierarchy
  - We will focus on agglomerative clustering only



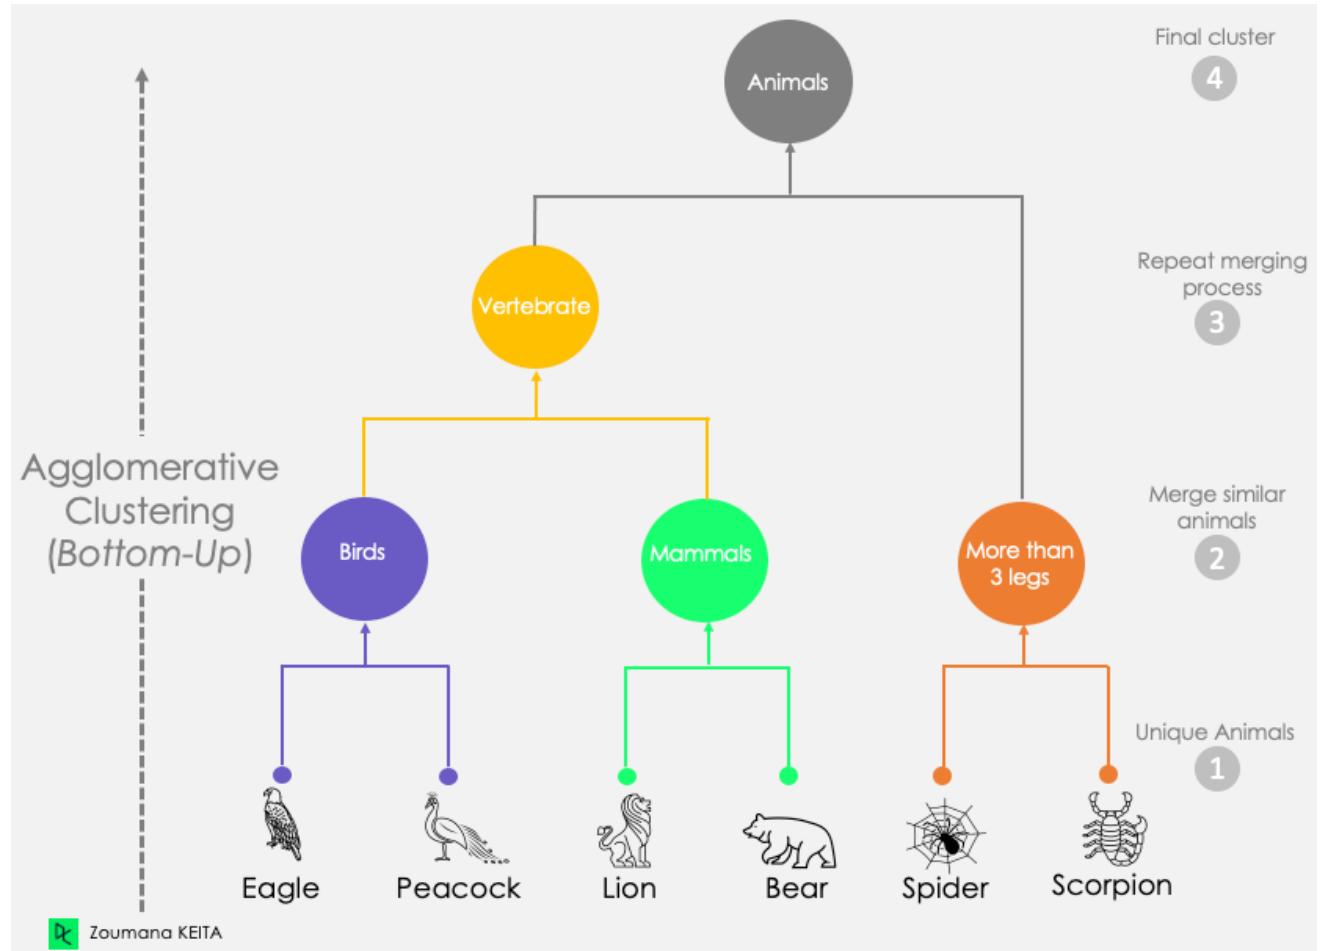
# Hierarchical Clustering

- Agglomerative clustering
  1. Start with each point in its own cluster
  2. Identify the two closest clusters. Merge them.
  3. Repeat until all points are in a single cluster.



# Hierarchical Clustering

- Agglomerative clustering



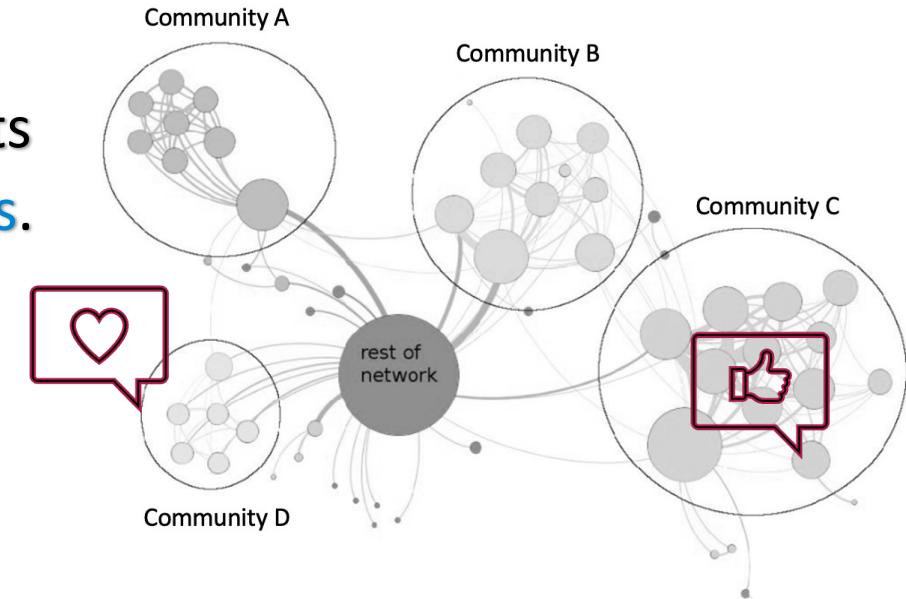
# Hierarchical Clustering

- Agglomerative clustering is based on **distance**
  - Close points are similar to each other → Small distance
  - Distance represents the **dissimilarity** between points
  - **Distance metrics**
    - Euclidean distance (default in sklearn)
    - Other options

metric	Function
'cityblock'	metrics.pairwise.manhattan_distances
'cosine'	metrics.pairwise.cosine_distances
'euclidean'	metrics.pairwise.euclidean_distances
'haversine'	metrics.pairwise.haversine_distances
'l1'	metrics.pairwise.manhattan_distances
'l2'	metrics.pairwise.euclidean_distances
'manhattan'	metrics.pairwise.manhattan_distances
'nan_euclidean'	metrics.pairwise.nan_euclidean_distances

# Hierarchical Clustering

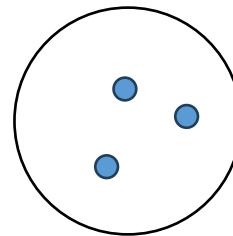
- Agglomerative clustering is based on distance
  - Distance represents the dissimilarity between points
  - Some special dataset needs special distance metrics.
    - For example, if we have a social media network dataset records the interactions between different users.
    - The more frequently two users interact, the closer we assume them to be.



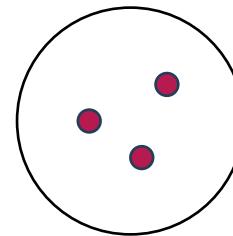
$$distance = \frac{1}{Frequency\ of\ interaction}$$

# Hierarchical Clustering

- How to calculate the distance between two clusters?
- The distance (or dissimilarity) between two clusters is called the linkage.
- **Linkage**
  - A dissimilarity measure
  - Given two clusters, a linkage tells us how different the points in these clusters are



*Cluster S*



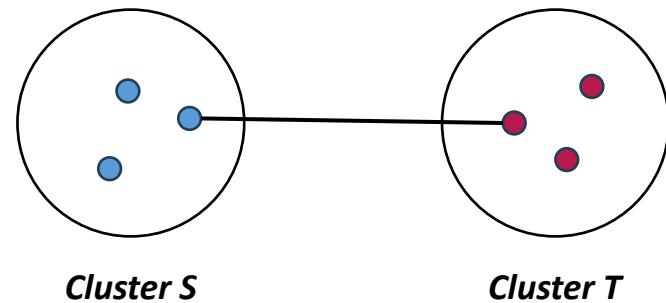
*Cluster T*

# Hierarchical Clustering

- Common linkage types

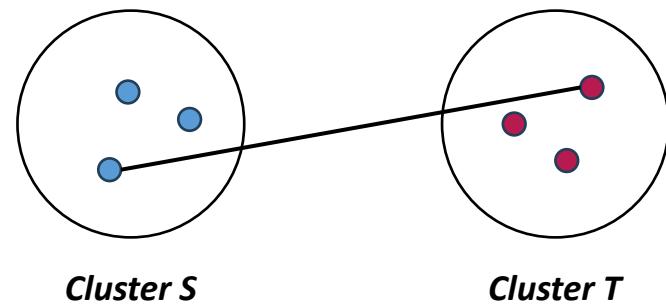
- Single linkage

- Minimal distance between the closest points of the pair of clusters
- $d(S, T) = \min\{d(x, y) : x \in S, y \in T\}$
- $d(x, y)$  is the distance between elements  $x$  and  $y$



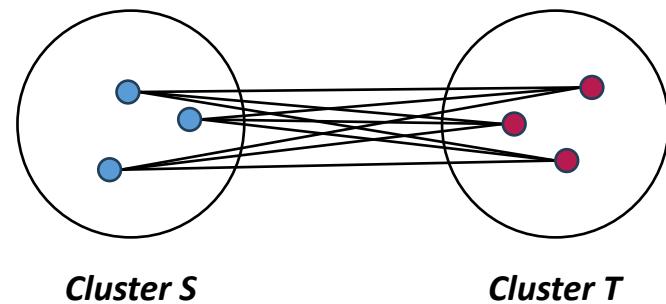
# Hierarchical Clustering

- Common linkage types
  - Single linkage
  - Complete linkage
    - Longest distance between any two points in the clusters
    - $d(S, T) = \max\{d(x, y) : x \in S, y \in T\}$
    - $d(x, y)$  is the distance between elements  $x$  and  $y$



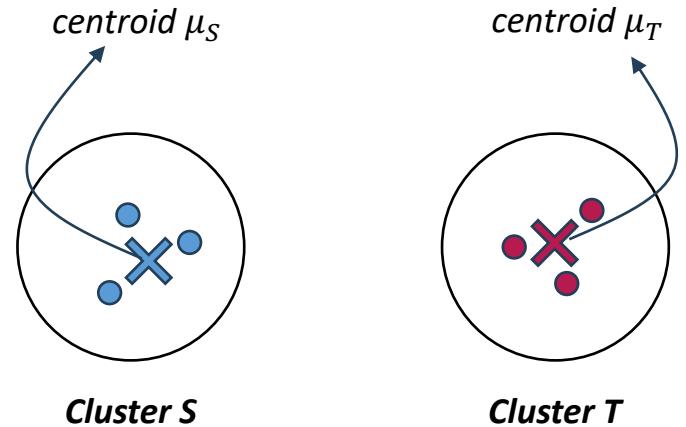
# Hierarchical Clustering

- Common linkage types
  - Single linkage
  - Complete linkage
  - Average linkage
    - Average distance between all pairs of points in the two clusters
    - $d(S, T) = \frac{1}{|S| \times |T|} \sum_{x \in S} \sum_{y \in T} d(x, y)$
    - $d(x, y)$  is the distance between elements  $x$  and  $y$
    - $|S|$  and  $|T|$  are the sizes of clusters  $S$  and  $T$



# Hierarchical Clustering

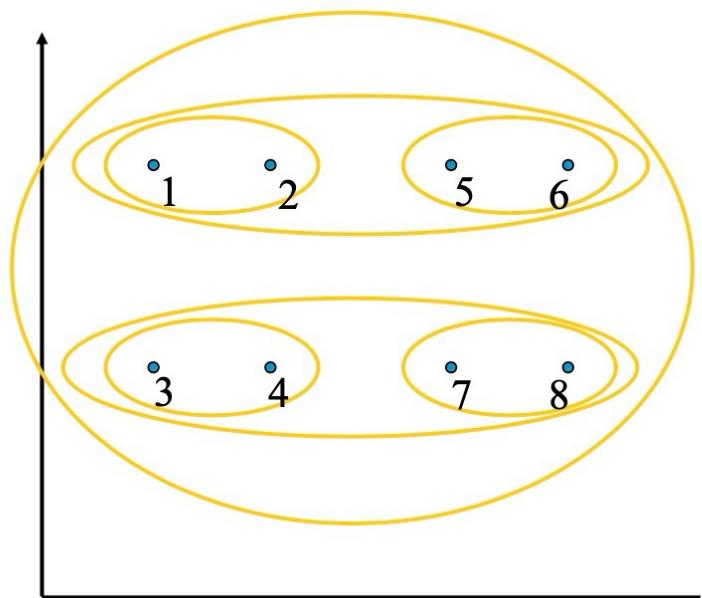
- Common linkage types
  - Single linkage
  - Complete linkage
  - Average linkage
  - Ward linkage
    - The closest pair of clusters should have minimal increase in total within-cluster variance after merging the two



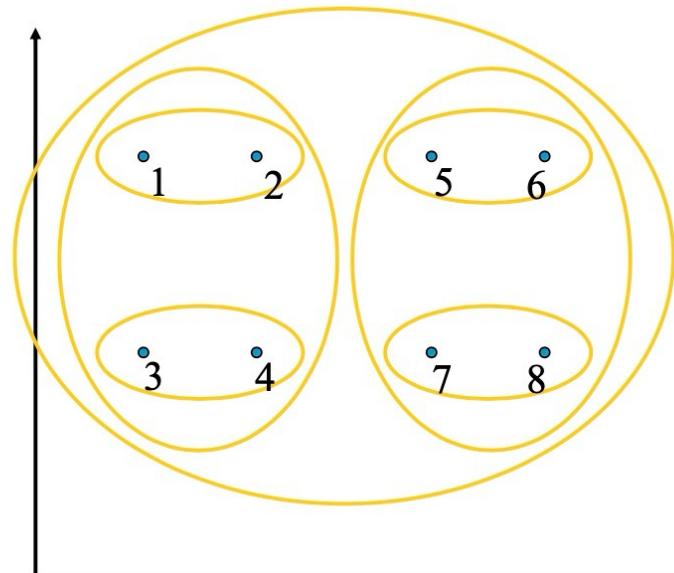
$$d(S, T) = \underbrace{\sum_{x \in S \cup T} \|x - \mu_{S \cup T}\|^2}_{\text{within-cluster variance after merging cluster } S \text{ and } T} - \underbrace{\sum_{x \in S} \|x - \mu_S\|^2}_{\text{within-cluster variance of cluster } S} - \underbrace{\sum_{x \in T} \|x - \mu_T\|^2}_{\text{within-cluster variance of cluster } T}$$

# Hierarchical Clustering

- Different types of linkage will lead to different results.



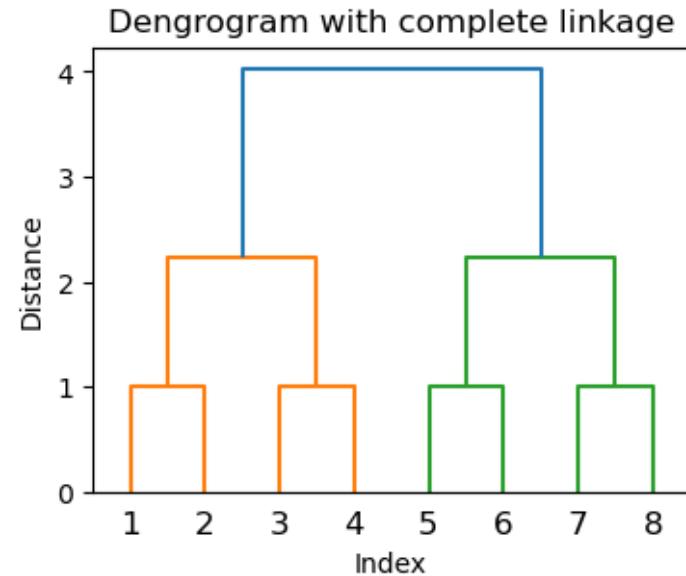
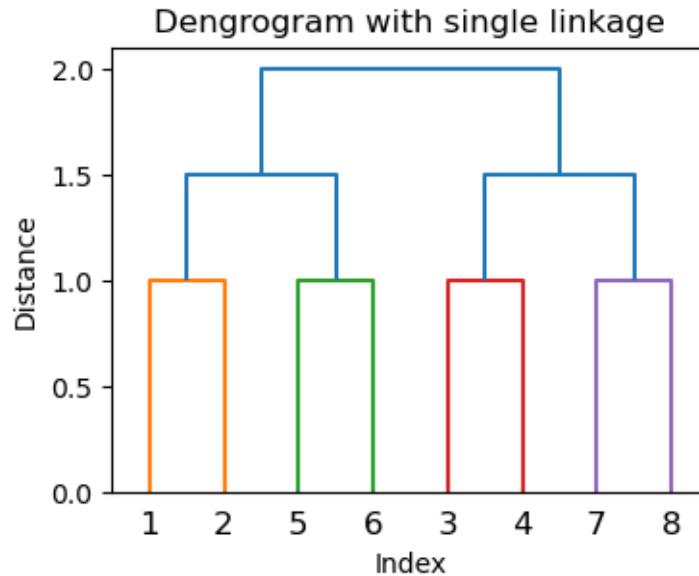
*Single linkage*



*Complete linkage*

# Hierarchical Clustering

- Different types of linkage will lead to different results.



# Hierarchical Clustering

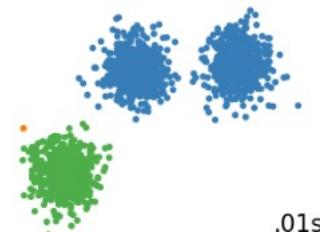
- Different types of linkage will lead to different results.
  - Single linkage
    - Tends to produce long clusters as chains.
    - Can handle non-elliptical shapes.
    - Have problems when clusters are close to each other



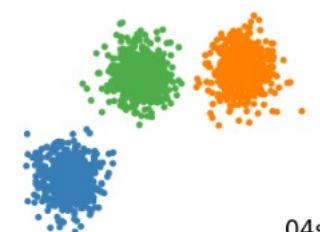
*Single  
Linkage*



*Complete  
Linkage*



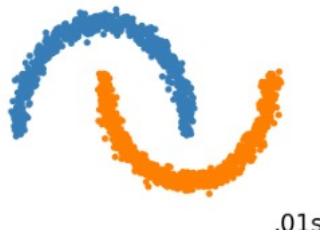
*Single  
Linkage*



*Complete  
Linkage*

# Hierarchical Clustering

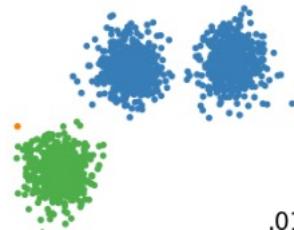
- Different types of linkage will lead to different results.
  - Complete linkage
    - Tends to produce compact, tightly bound clusters.
    - Struggle with elongated shapes.
    - Can handle compact clusters that are close to each other.



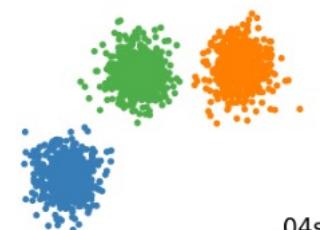
*Single  
Linkage*



*Complete  
Linkage*



*Single  
Linkage*



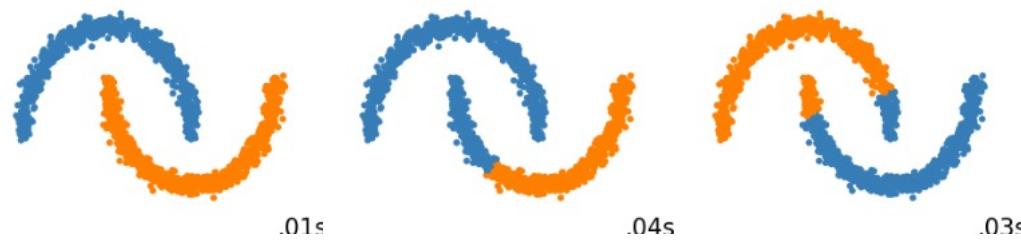
*Complete  
Linkage*

# Hierarchical Clustering

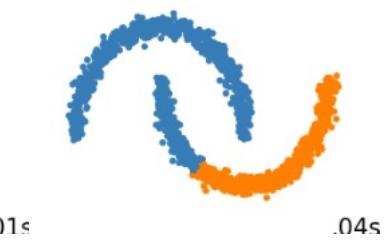
- Different types of linkage will lead to different results.

- Average linkage

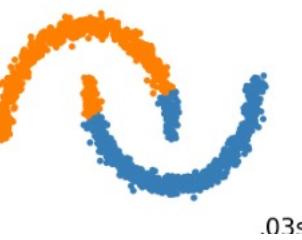
- Strikes a balance between single and complete linkage
- Sometime can produce more natural clusters, but not always
- Good to use when there is no prior knowledge about the distribution of the data



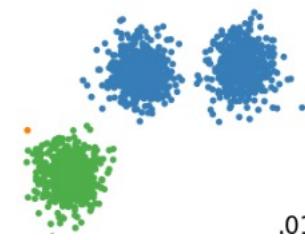
*Single  
Linkage*



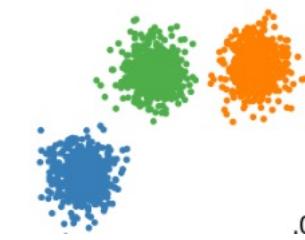
*Average  
Linkage*



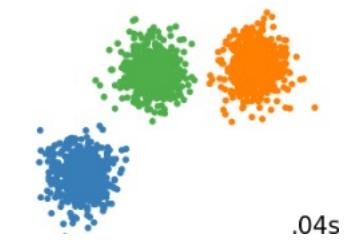
*Complete  
Linkage*



*Single  
Linkage*



*Average  
Linkage*



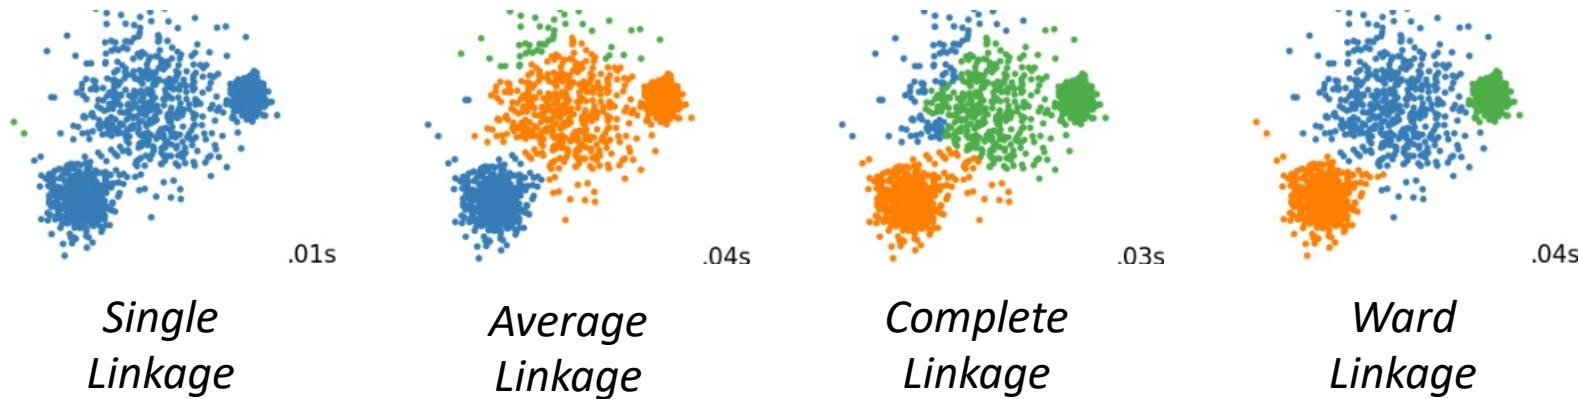
*Complete  
Linkage*

# Hierarchical Clustering

- Different types of linkage will lead to different results.

- Ward linkage

- Minimize the within-cluster variance
  - Tend to produce spherical clusters.
  - Can be more computationally intensive.

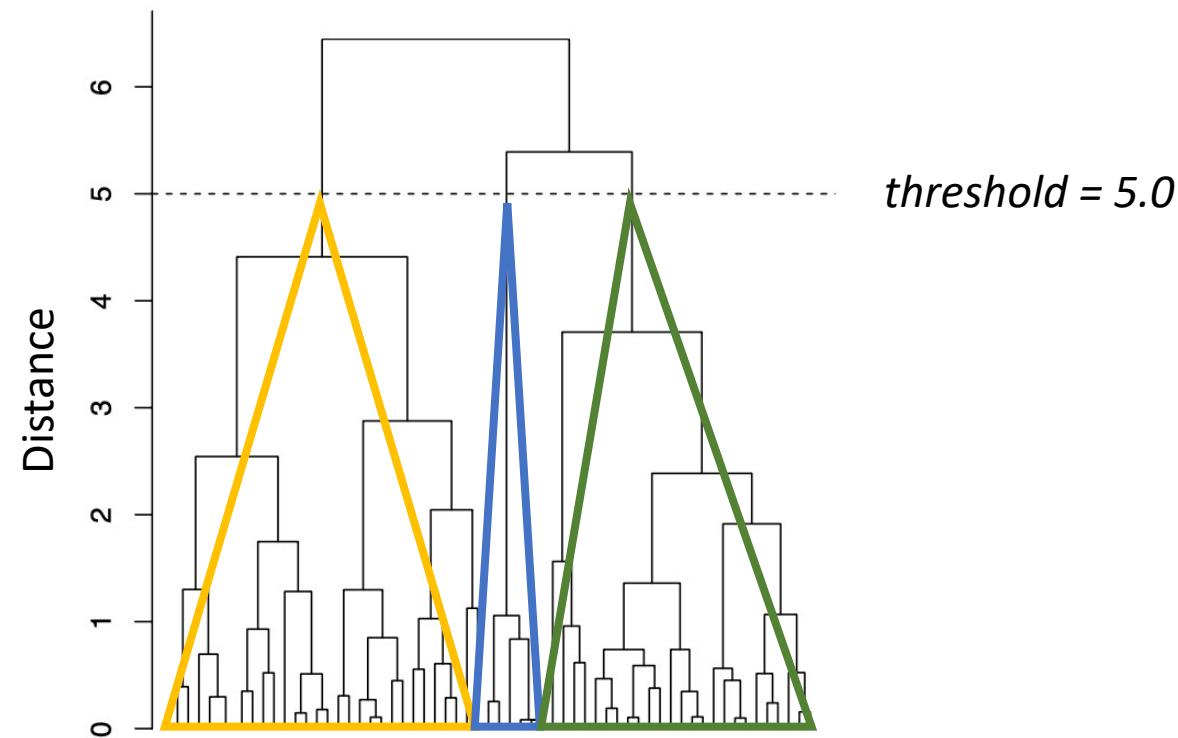
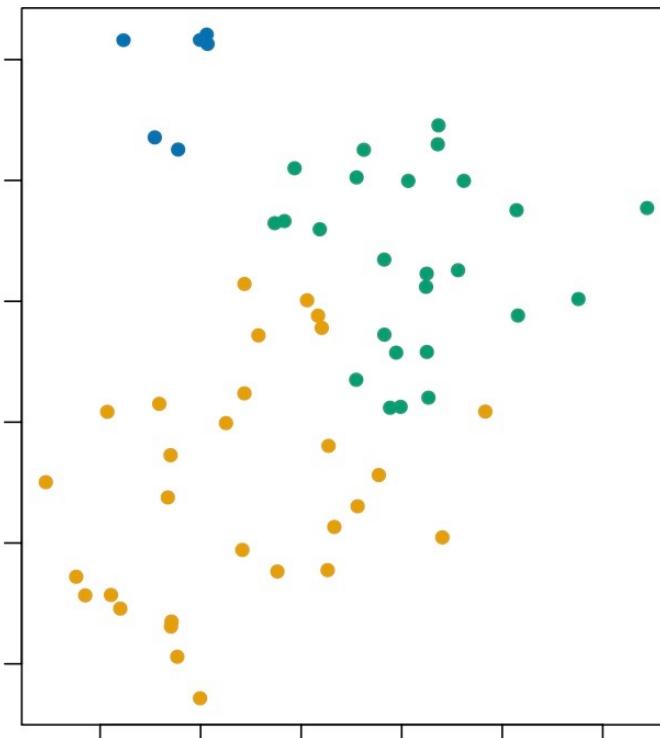


# Hierarchical Clustering

- How to get the clustering assignments from the hierarchy of clusters?
- We need to set a distance threshold for cutting the tree
  - Clusters with distance smaller than the threshold will be merged.
  - Clusters with distance larger than the threshold won't be merged.

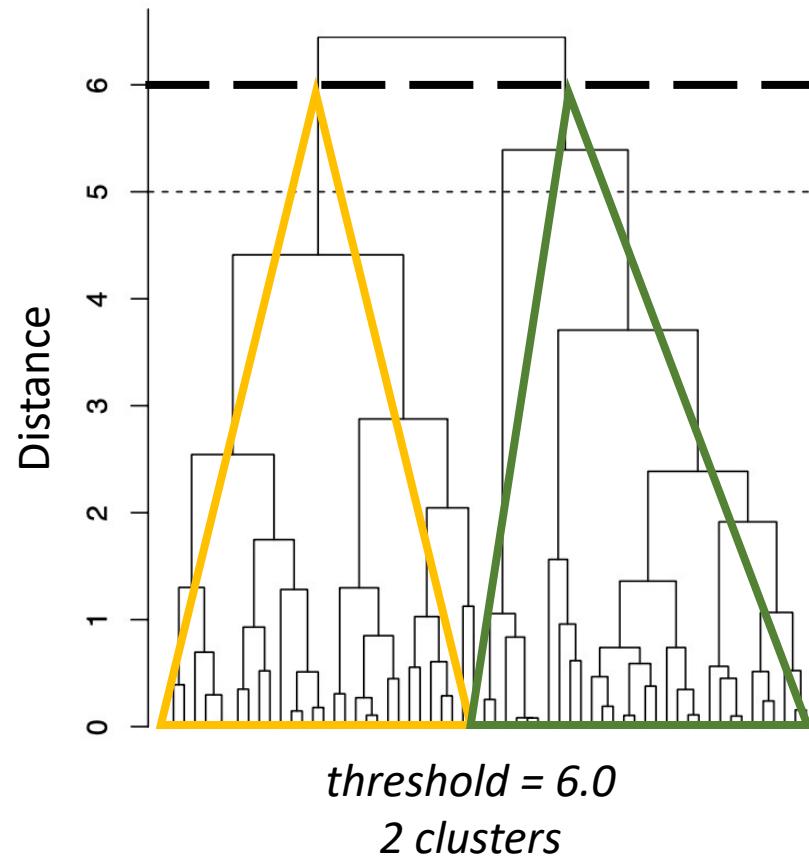
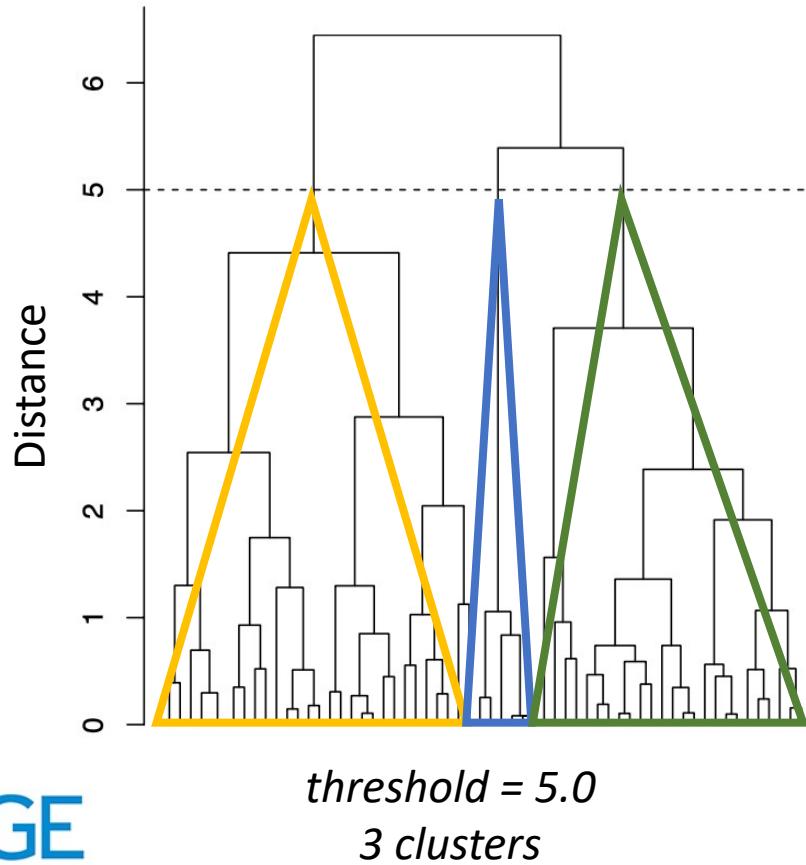
# Hierarchical Clustering

- Set a distance threshold for cutting the tree



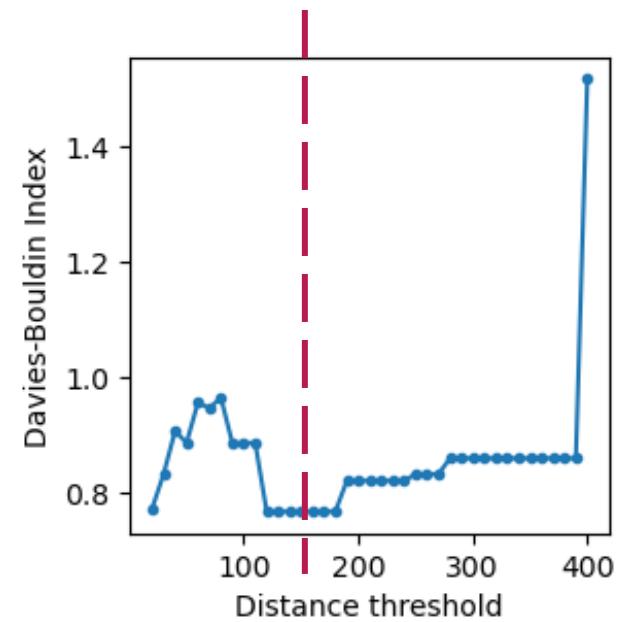
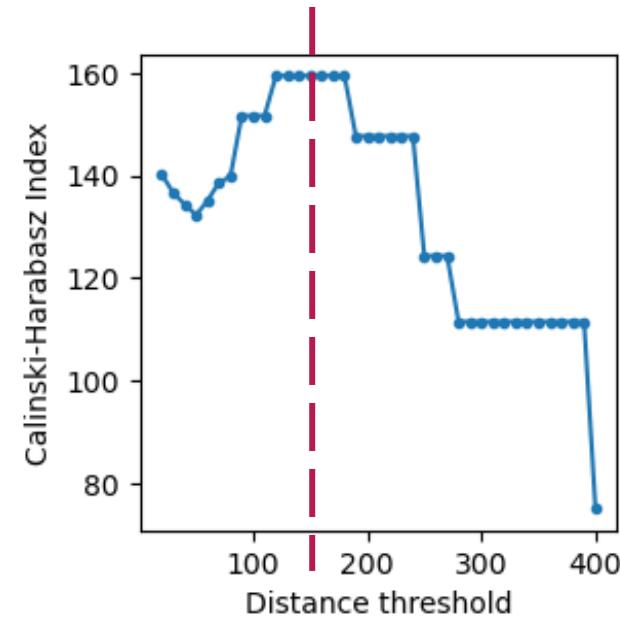
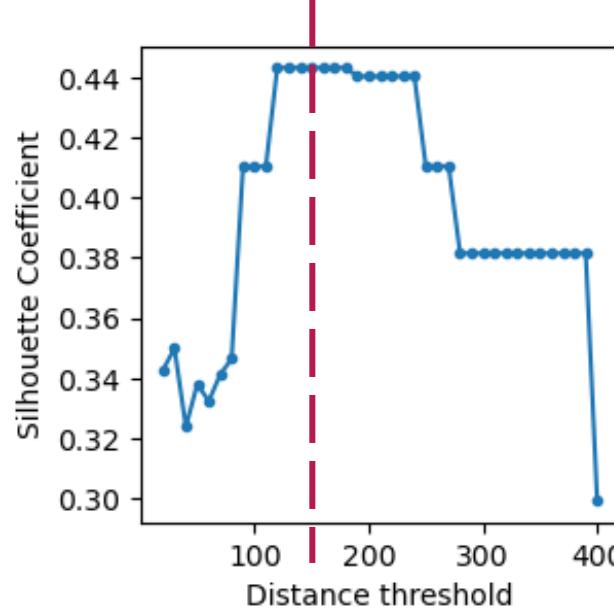
# Hierarchical Clustering

- Set a distance threshold for cutting the tree



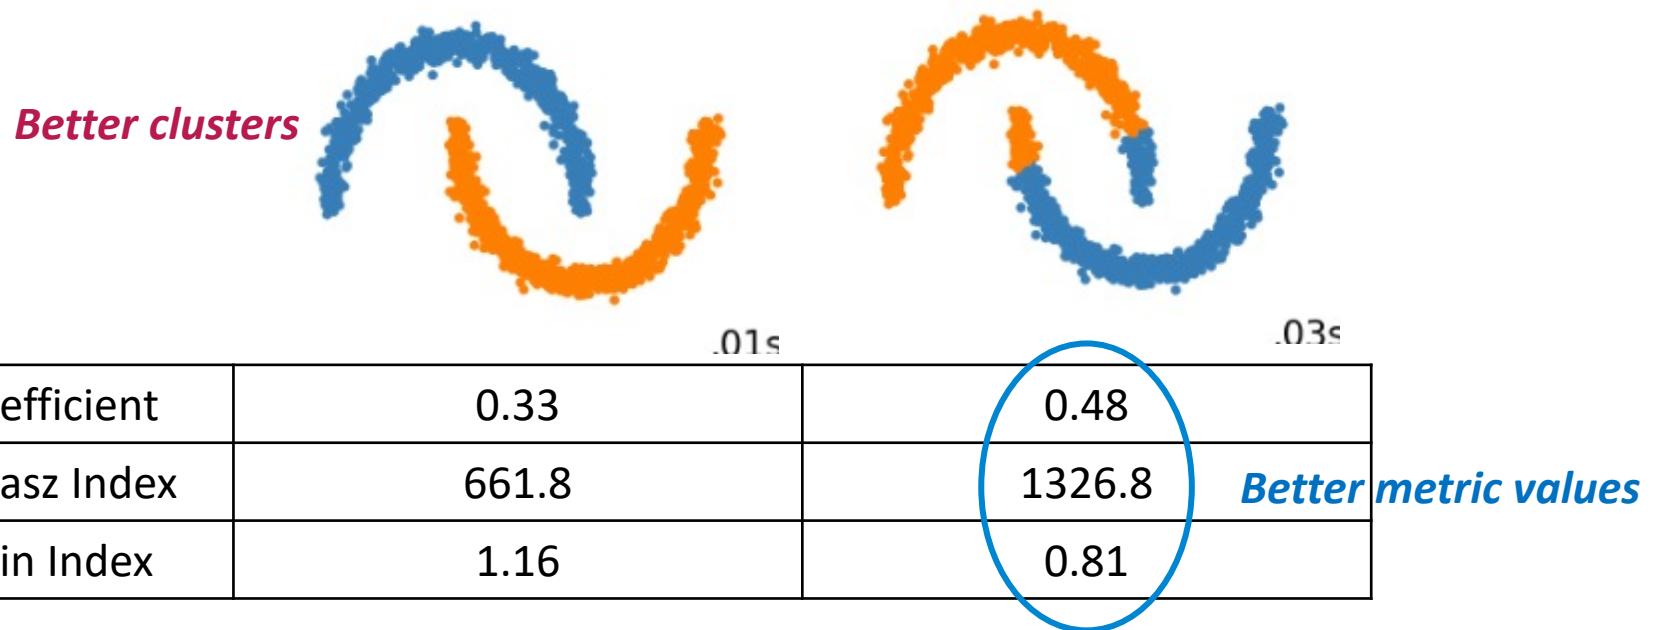
# Hierarchical Clustering

- How to select the distance threshold?
  - Use clustering quality evaluation metrics as reference



# Hierarchical Clustering

- How to select the distance threshold?
  - Use clustering quality evaluation metrics as reference
  - Good metric values can't guarantee the meaningful results



# Hierarchical Clustering

- Instead of selecting a distance threshold, we can directly specify the number of clusters we want

## sklearn.cluster.AgglomerativeClustering

```
class sklearn.cluster.AgglomerativeClustering(n_clusters=2, *,  
affinity='deprecated', metric=None, memory=None, connectivity=None,  
compute_full_tree='auto', linkage='ward', distance_threshold=None,  
compute_distances=False) ¶
```

[source]

**Parameters:** **n\_clusters : int or None, default=2**

The number of clusters to find. It must be `None` if `distance_threshold` is not `None`.

# Hands-on Exercise

- Exercise 08 Clustering II