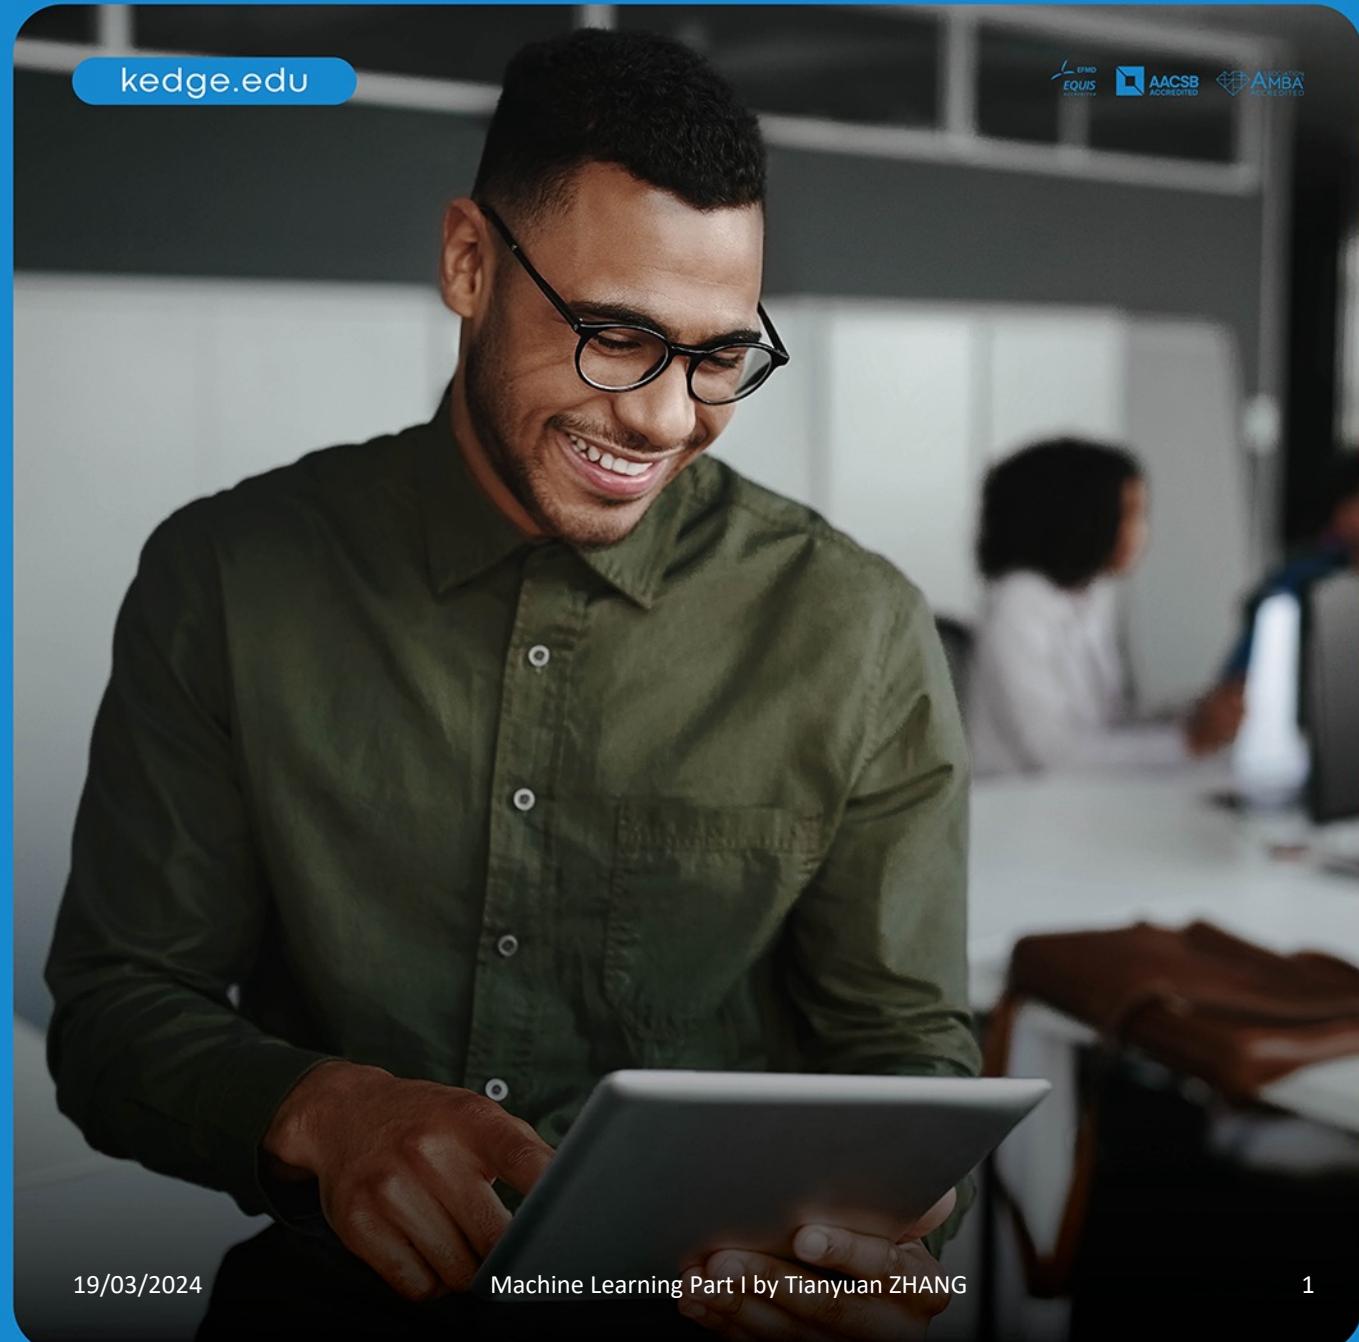


ARTIFICIAL INTELLIGENCE NEEDS REAL INTELLIGENCE

Linear Neural Network I

Professor: Tianyuan ZHANG
tianyuan.zhang@kedgebs.com



Recap of Ensemble Learning

- Ensemble Learning
 - A machine learning paradigm where multiple models (often called “**weak learners**”) are trained to solve the same problem and combined to get better results.
- Rationale
 - No single model can capture all the patterns in the data perfectly.
 - Leverage the strength and compensate for the weaknesses of individual models.
 - A group of “weak learners” can come together to form a “strong learner”.
 - **Wisdom of the crowd**

Recap of Ensemble Learning

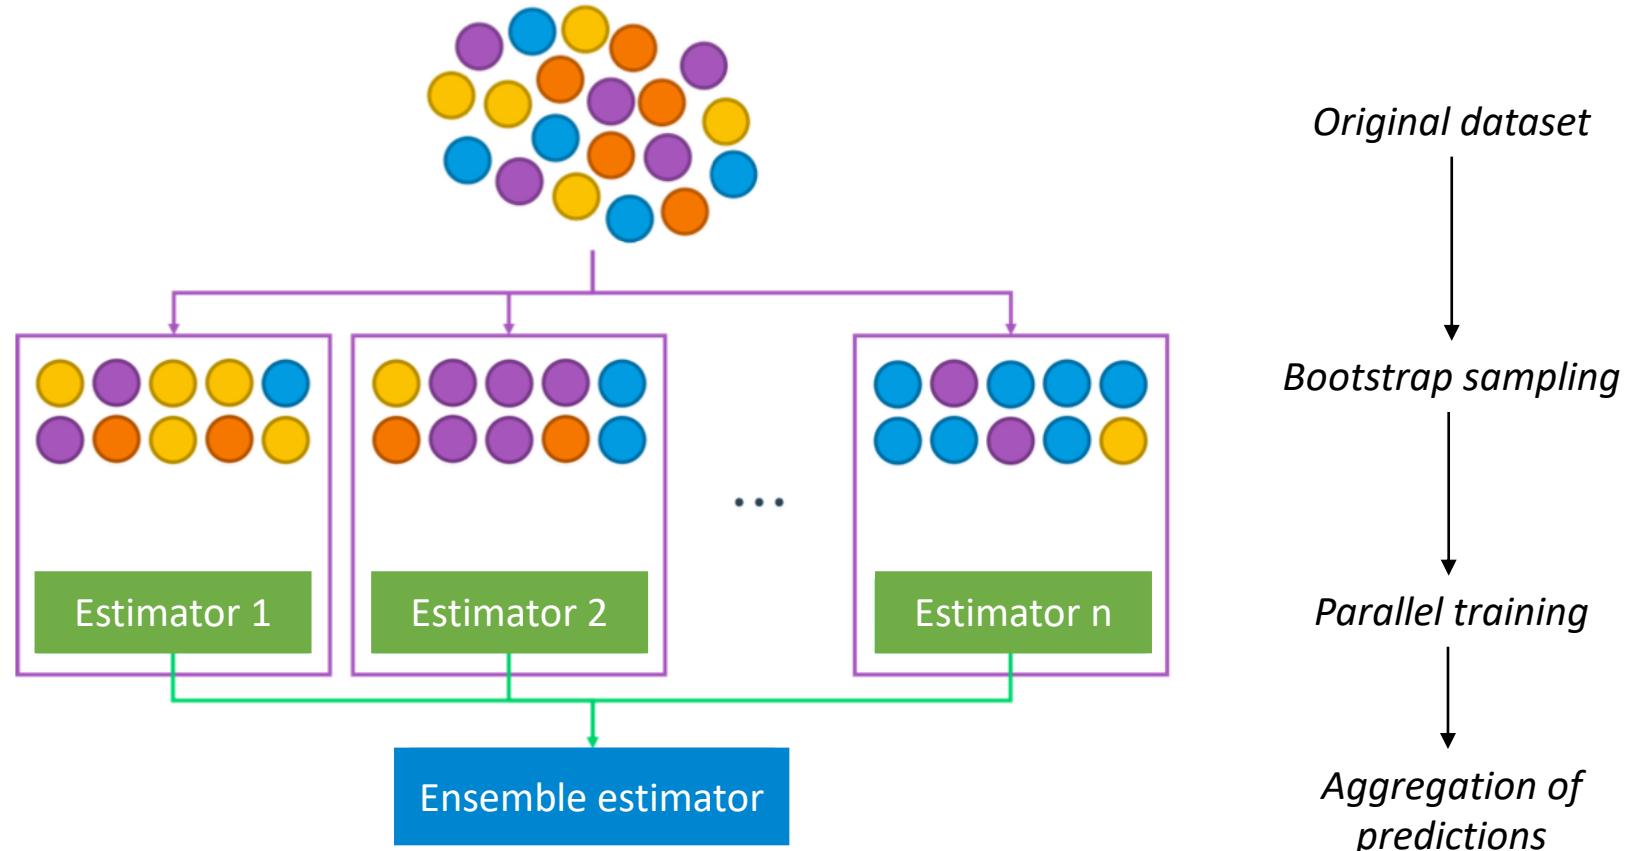
- Two questions to answer when designing ensemble learning methods
 1. How to generate multiple base models (estimators)?
 2. How to integrate / combine their predictions to make final prediction?
- Different decisions lead to different **types** of ensemble learning methods
 - Bagging
 - Boosting
 - Stacking

Recap of Ensemble Learning

- **Bootstrap Aggregating → Bagging**
- An ensemble learning technique
 - **Parallel model training**
 - Multiple models / estimators are trained simultaneously on different subsets of the original training dataset.
 - **Bootstrap sampling**
 - Each subset of training data is generated by randomly sampling with replacement.
 - **Aggregation of predictions**
 - Regression: Averaging the predictions
 - Classification: Majority voting

Recap of Ensemble Learning

- **Bagging**



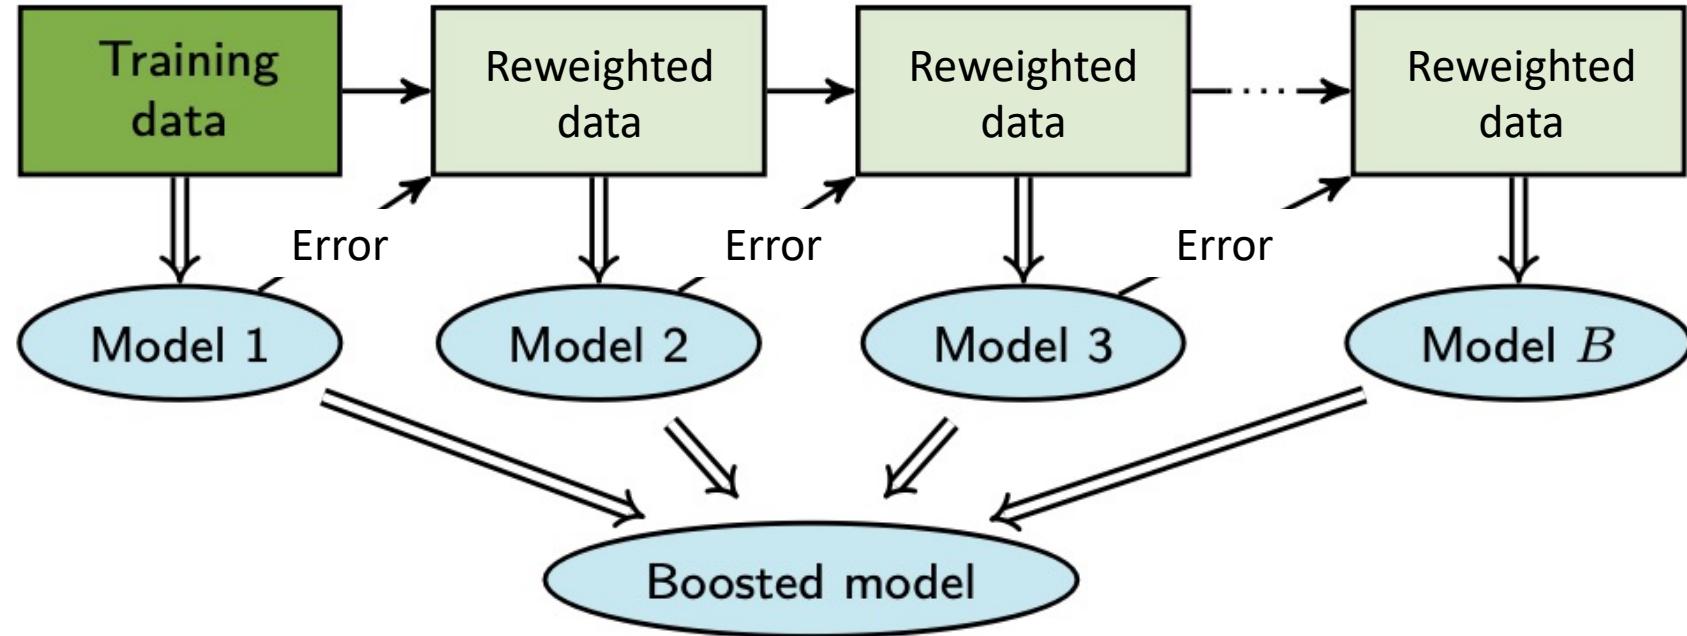
Recap of Ensemble Learning

- **Boosting**

- A type of ensemble learning methods that boosts weak learners to strong learners
- Key features
 - **Sequential model training**
 - Each model in the sequence is trained on the dataset with a focus on the instances that previous models misclassified.
 - **Error correction focus**
 - The algorithm identifies the errors or mispredictions of the previous model and increases their weight, making them more significant in the training of the next model.
 - **Cumulative learning**
 - Unlike parallel ensemble methods, boosting involves cumulative learning where each model builds on the knowledge of its predecessors.

Recap of Ensemble Learning

- **Boosting**



- An **iterative** process, where the training set is **reweighted** in each iteration

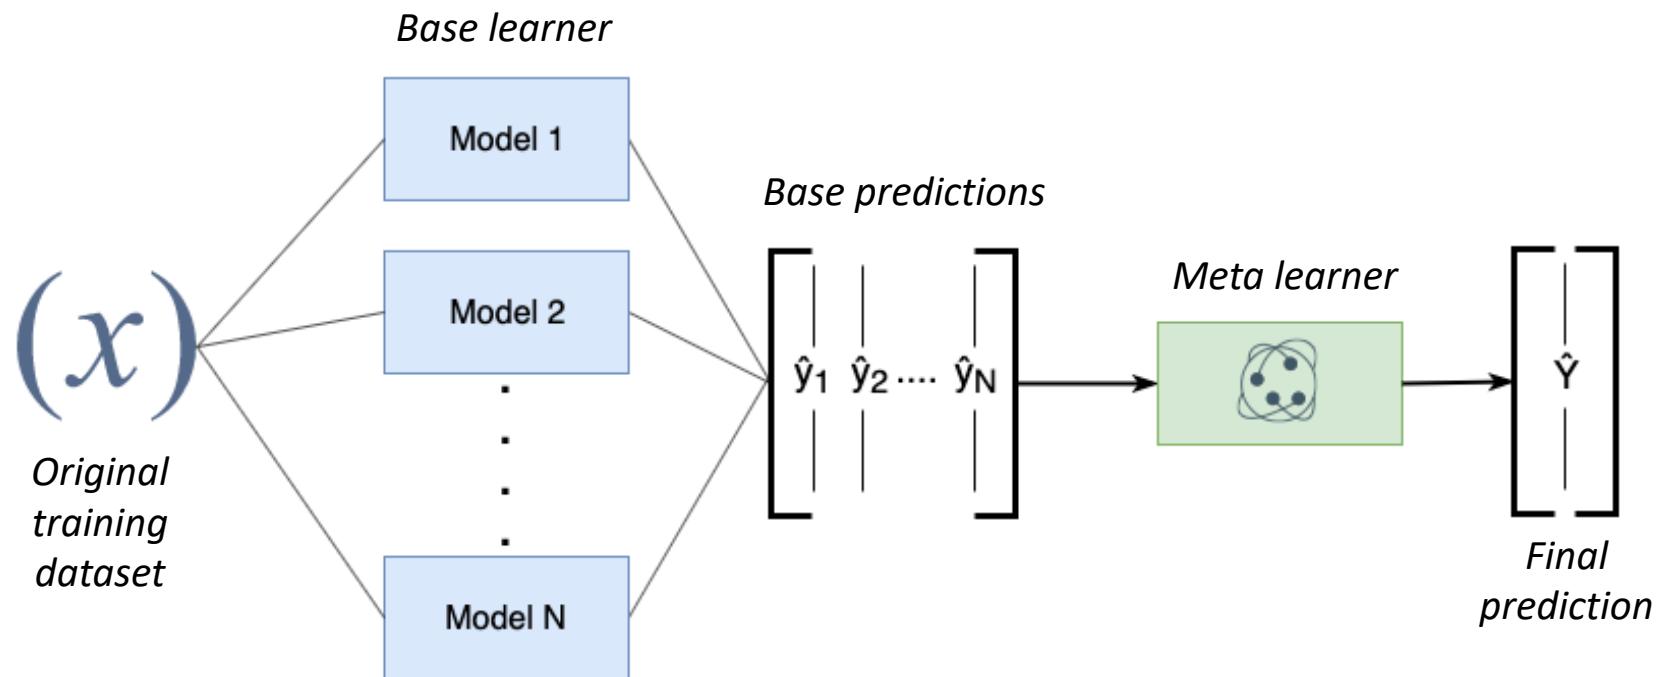
Recap of Ensemble Learning

- **Stacking**

- Same training set + different algorithms → different weak learners
- Integrate base predictions using a meta-learner
- Process
 1. Choose M different algorithms, train M different base learners $\{\hat{y}^j(\mathbf{x})\}_{j=1}^n$ on the training data set $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$
 2. Form a new training data set using the base predictions
 - New training data set $\{(\mathbf{x}'_i, y_i)\}_{i=1}^n$
 - $\mathbf{x}'_i = \{\hat{y}_i^1(\mathbf{x}_i), \dots, \hat{y}_i^M(\mathbf{x}_i)\}$
 3. Train a meta-learner $\hat{Y}(\mathbf{x})$ on the new training data set $\{(\mathbf{x}'_i, y_i)\}_{i=1}^n$

Recap of Ensemble Learning

- **Stacking**



Recap of Ensemble Learning

- Ensemble Learning
 - A way to combine a group of “weak learners” into one “strong learner”
 - Reduce the prediction errors of a single model
 - Bias: under-fitting model
 - Variance: over-fitting model
 - If the available machine learning algorithms fail to solve the complex problem, ensemble learning can help.
- What if **a single model** can be **complex enough** to avoid high bias and **generalized well** to avoid high variance?
 - **Human being**
 - **Artificial Neural Network → Deep Learning**

Outline

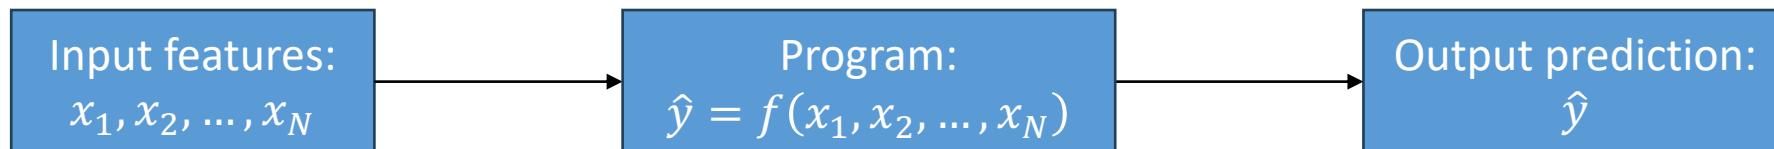
- **Classical Machine Learning vs. Deep Learning**
- Artificial Neural Network
- Implement an ANN
- Preliminaries with PyTorch

Classical Machine Learning vs. Deep Learning

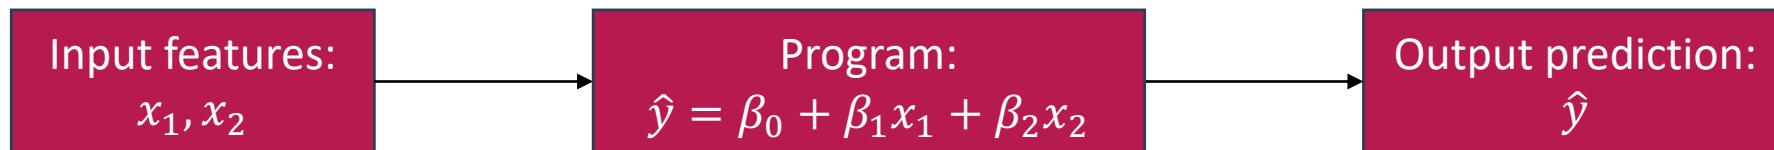
- Break down classical ML into several key concepts

- **Parameter**

- The **behavior** of a **flexible program** is determined by a number of parameters



- Example: Linear regression



- $\beta_0, \beta_1, \beta_2$ are parameters that determine the behavior of the program
 - Given the same inputs, different values of parameters lead to different outputs

Classical Machine Learning vs. Deep Learning

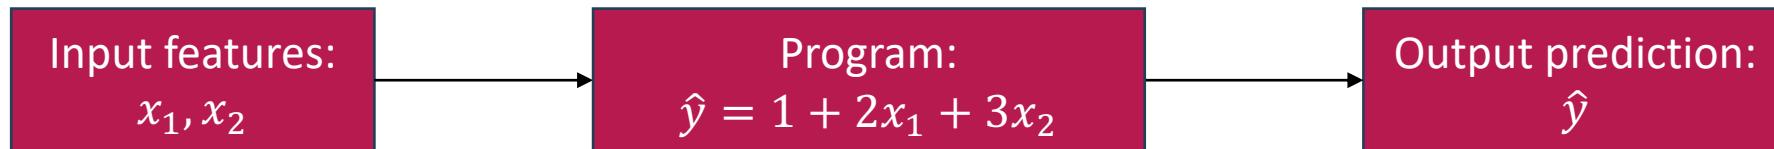
- Break down classical ML into several key concepts

- **Parameter**

- The behavior of a flexible program is determined by a number of parameters

- **Model**

- Once the **parameters are fixed**, the flexible program becomes a model
 - Example: Linear regression



- $\beta_0, \beta_1, \beta_2$ are fixed at 1, 2, 3
 - Then we can use this program to calculate the outputs given the inputs
 - It becomes a ML model that can make predictions given the inputs

Classical Machine Learning vs. Deep Learning

- Break down classical ML into several key concepts
 - **Parameter**
 - The behavior of a flexible program is determined by a number of parameters
 - **Model**
 - Once the parameters are fixed, the flexible program becomes a model
 - **A family of model**
 - The **set of all distinct models (input-output mapping)** that we can produce by manipulating the parameters of the same flexible program
 - Example: Linear regression
 - Model 1 $\rightarrow \hat{y} = 1 + 2x_1 + 3x_2$
 - Model 2 $\rightarrow \hat{y} = 4 - 1x_1 + 2x_2$
 - Model 3 $\rightarrow \hat{y} = 0 + 7x_1 + 5x_2$

Classical Machine Learning vs. Deep Learning

- Break down classical ML into several key concepts
 - **Parameter**
 - The behavior of a flexible program is determined by a number of parameters
 - **Model**
 - Once the parameters are fixed, the flexible program becomes a model
 - **A family of model**
 - The **set of all distinct models (input-output mapping)** that we can produce by manipulating the parameters of the same flexible program
 - **Learning process**
 - The “meta-program’ that uses training data to **choose the best values of parameters**, which will **produce a model**
 - Example: Linear regression, use ordinary least squares method to estimate $\beta_0, \beta_1, \beta_2$

Classical Machine Learning vs. Deep Learning

- Break down classical ML into several key concepts
 - **Parameter**
 - **Model**
 - **A family of model**
 - **Learning process**
 - The “meta-program’ that uses training data to choose the best values of parameters, which will produce a model
 - **Objective function**
 - A formal measure of how good (or bad) the models are.
 - Example: RMSE for regression, Accuracy for classification
 - **Optimization algorithm**
 - Search for the best values of parameters for minimizing / maximizing the objective function

Classical Machine Learning vs. Deep Learning

- Classical ML → Deep Learning
 - Parameter → Parameter of **artificial neuron**
 - Model → **Artificial neural networks** (consists of artificial neurons)
 - Learning process
 - Objective function → **Loss function**
 - **Conventionally lower is better**
 - Optimization algorithm → **Any optimization algorithm**
 - The dominant approach is based on **gradient descent**
- **Advantages** of deep learning
 - Artificial neural network and its variants can form very complex models (**deep**)
 - Various techniques to avoid over-fitting and ensure generalization

Classical Machine Learning vs. Deep Learning

- Examples of deep learning applications

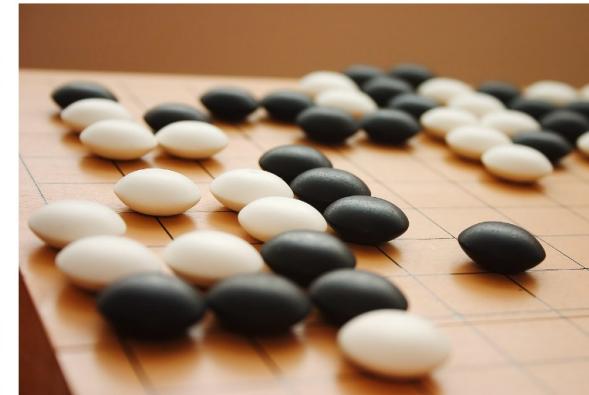
*Speech recognition
Natural language processing*



*Computer vision
Object recognition
Semantic segmentation*



Deep reinforcement learning



*Generative AI
Style transfer*

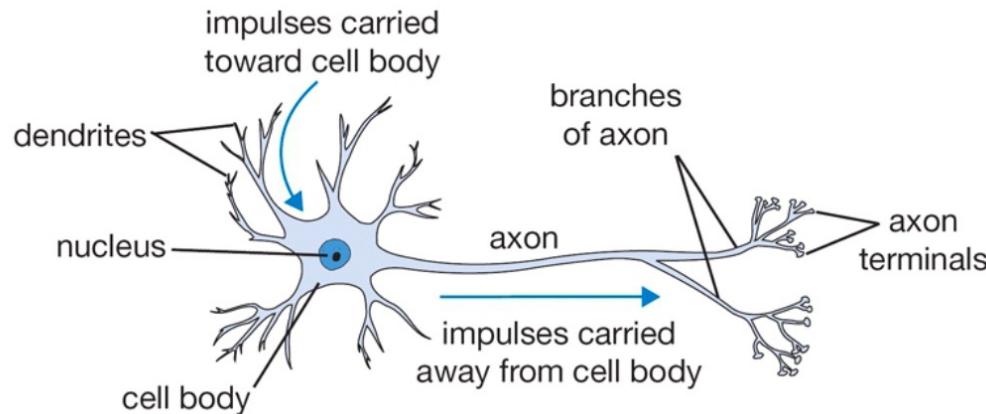


Outline

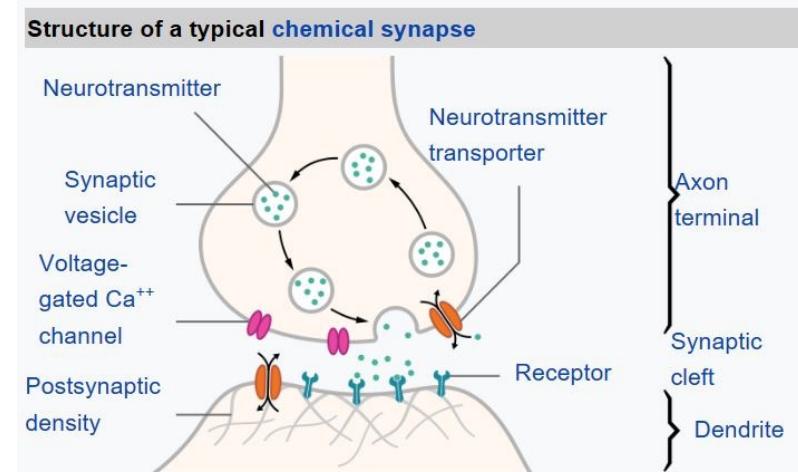
- Classical Machine Learning vs. Deep Learning
- **Artificial Neural Network**
- Implement an ANN
- Preliminaries with PyTorch

Artificial Neural Network

- Inspiration: The Brain
 - Human brain has $\sim 10^{11}$ neurons, each of which communicates (is connected) to $\sim 10^4$ other neurons



*The basic computation unit
of the brain: Neuron*

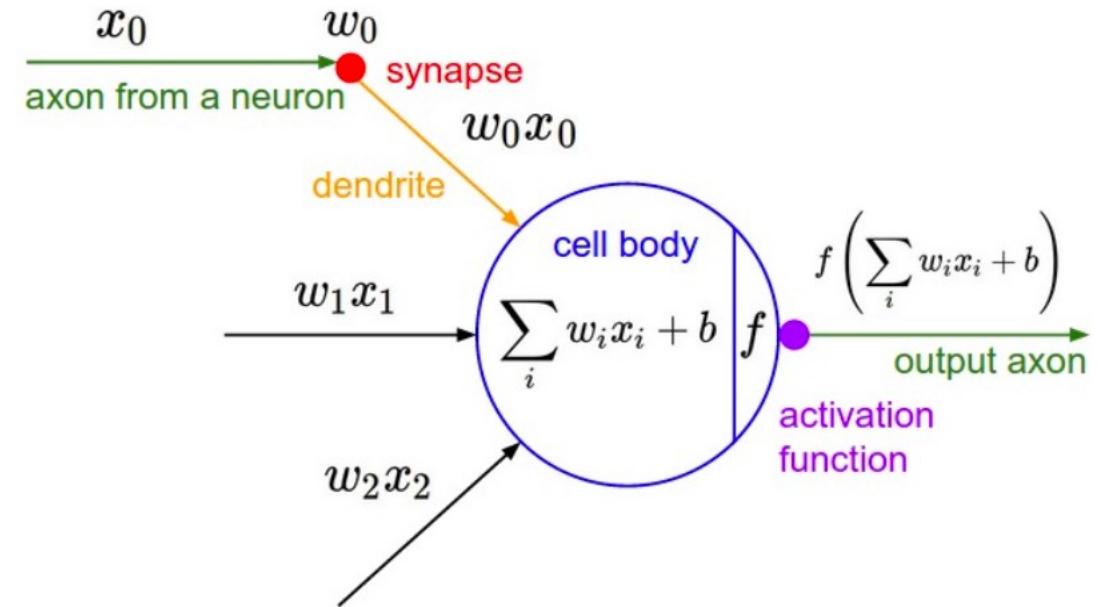


*Connections between axons
and dendrites*

Artificial Neural Network

- **Artificial Neuron**

- A mathematical model of a neuron
 - $f(\sum_{i=1}^n w_i x_i + b)$ is the **output** of a neuron
 - (x_1, \dots, x_n) are an **input** data point with n features
 - (w_1, \dots, w_n) are the **weights** of each feature
 - b is the **bias** added to the weighted sum
 - $(\sum_{i=1}^n w_i x_i + b)$ can be seen as the **impulse received** by the neuron
 - f is the **activation function**, determining whether the neuron is activated by the received impulse or not



Artificial Neural Network

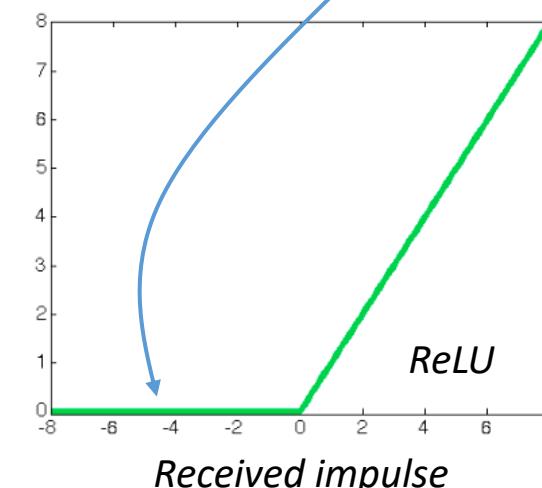
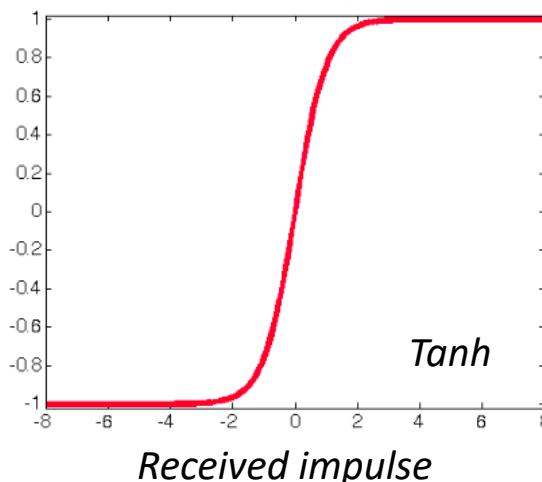
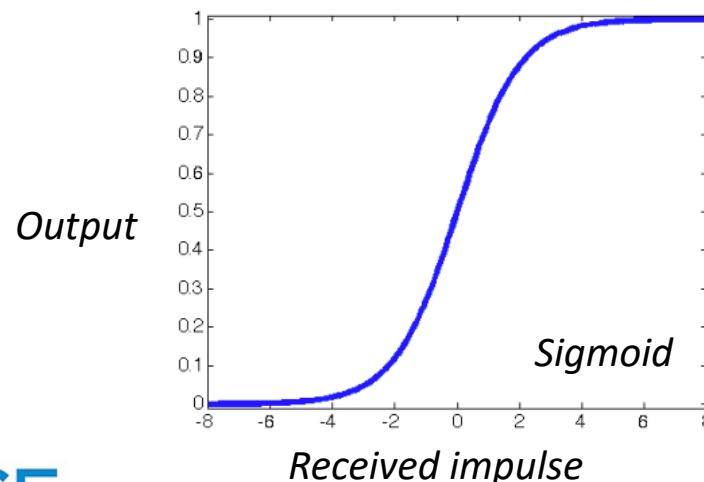
- Commonly used **activation functions**

- Linear: $f(z) = z$

- Non-linear:

- Sigmoid: $\sigma(z) = 1/(1 + e^{-z})$
- Tanh: $\tanh(z) = (e^z - e^{-z})/(e^z + e^{-z})$
- ReLU (Rectified Linear Unit): $\text{ReLU}(z) = \max(0, z)$

If the output of a neuron is 0 or close to 0, we can say the neuron is not activated



Artificial Neural Network

- Example Python scripts of a neuron with a sigmoid activation function



```
class Neuron(object):
    # ...
    def forward(inputs):
        """assume inputs and weights are 1-D numpy arrays"""
        """assume bias is a number"""
        received_impulses = np.sum(inputs * self.weights) + self.bias
        output = 1.0 / (1.0 + math.exp(-received_impulses)) # sigmoid activation function
        return output
```

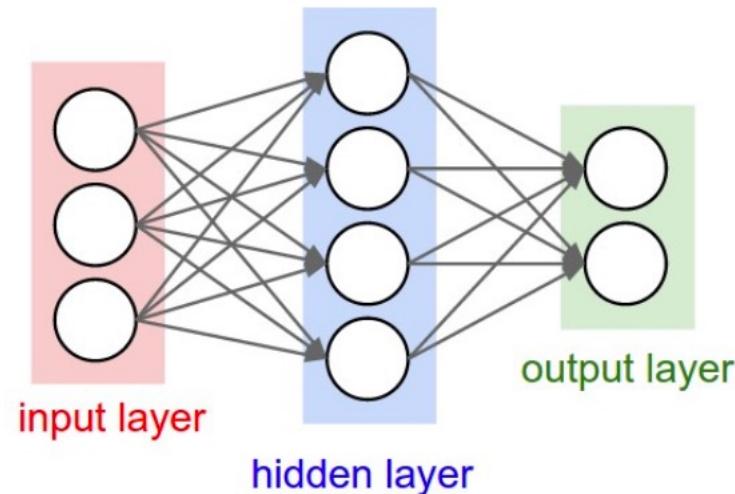
Forward means the artificial neuron is performing **inference**

- Take features of an instance as inputs
- Compute the received impulse using weights and bias
- Transform the impulse into output through the activation function

Artificial Neural Network

- **Artificial Neural Network (ANN)**

- A network consists of multiple connected artificial neurons arranged into different layers



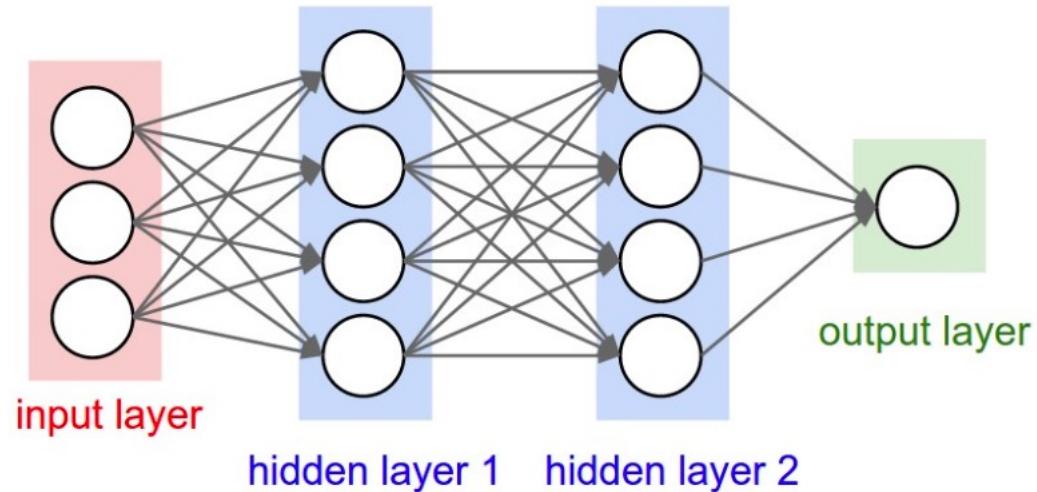
Terminology:

- input nodes
- output nodes
- hidden layers
- connections
- weights
- activation function

Artificial Neural Network

- **Artificial Neural Network (ANN)**

- A network consists of multiple connected artificial neurons arranged into different layers



A 3-layer neural network with
two hidden layers

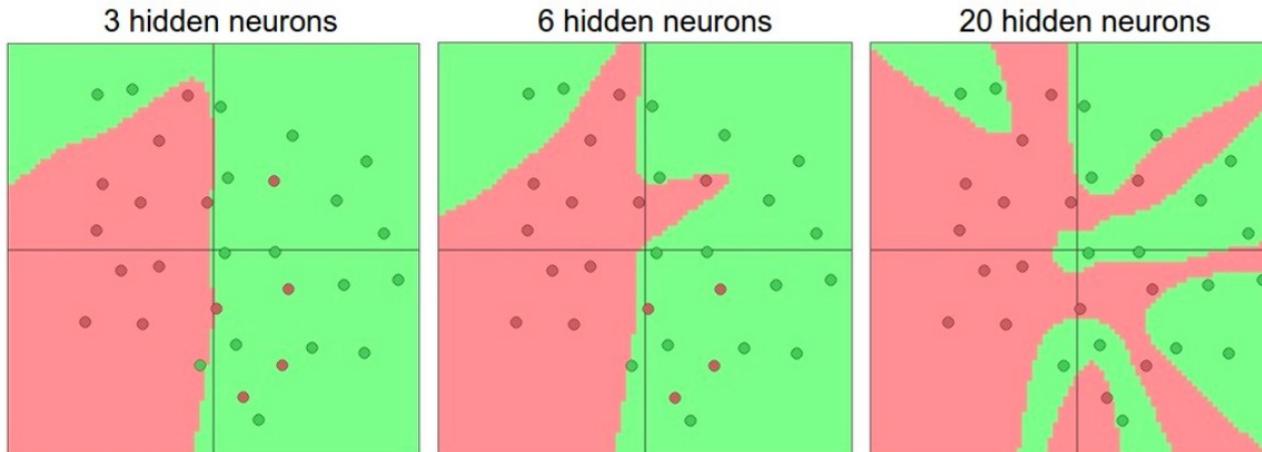
Naming conventions: a N-layer neural network

- N-1 hidden layers
- 1 output layer
- We do not count the input layer

Artificial Neural Network

- **Artificial Neural Network (ANN)**

- ANN is a family of model with excellent **representational power**
 - Depends on the number of hidden neurons and hidden layers, the network can be very simple or complex
- A 2-layer ANN can be a **universal approximator**
 - By adjusting the number of neurons in the hidden layer and their parameters, it can represent any function

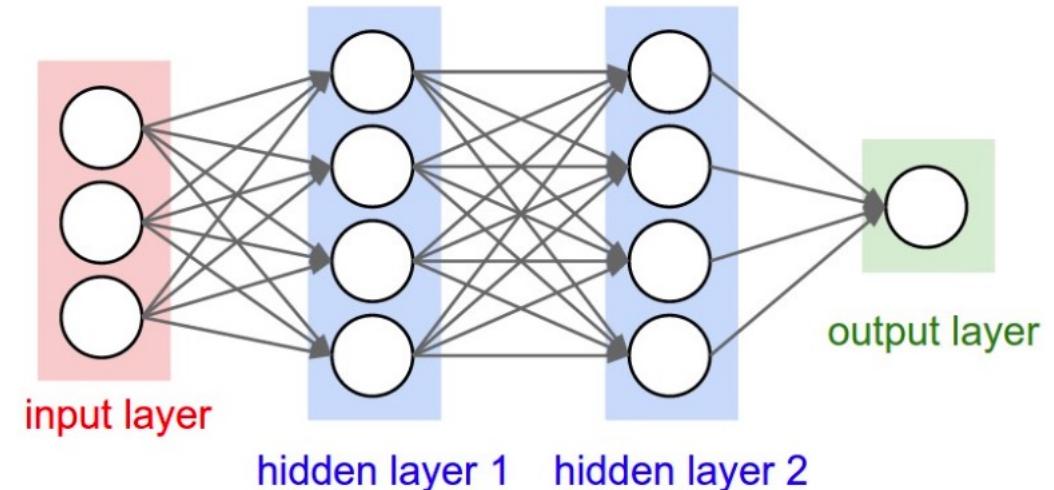


Artificial Neural Network

- Artificial Neural Network (ANN)

- Hyper-parameters

- The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron
- *The first three determined the structure of the ANN*
 - *Pre-defined when we create an ANN*
 - *Remain unchanged during the training process*



Artificial Neural Network

- **Artificial Neural Network (ANN)**

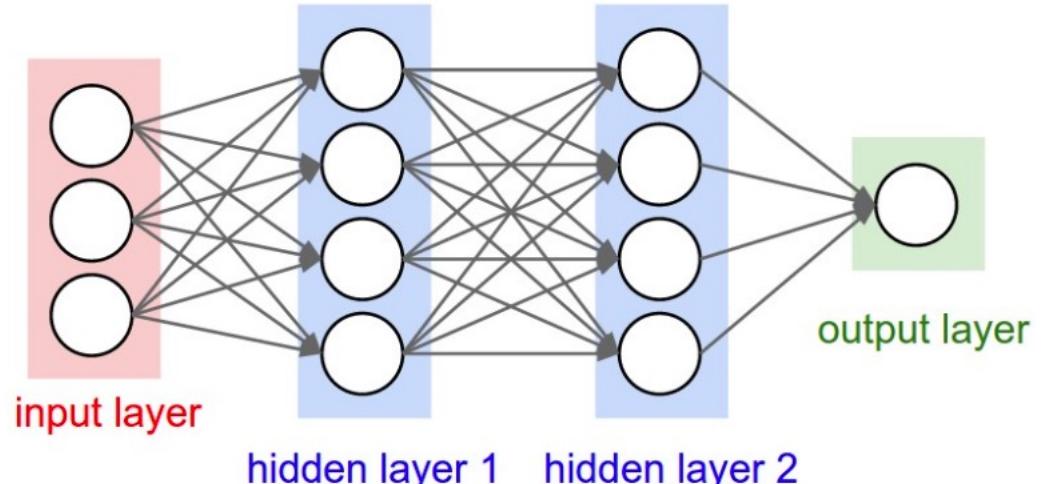
- Hyper-parameters

- The number of layers
- The number of neurons in each layer
- The connections between neurons
- The activation function of each neuron

- **Parameters → Learning parameters**

- The weight of each connection
- The bias of each neuron added to the weighted sum

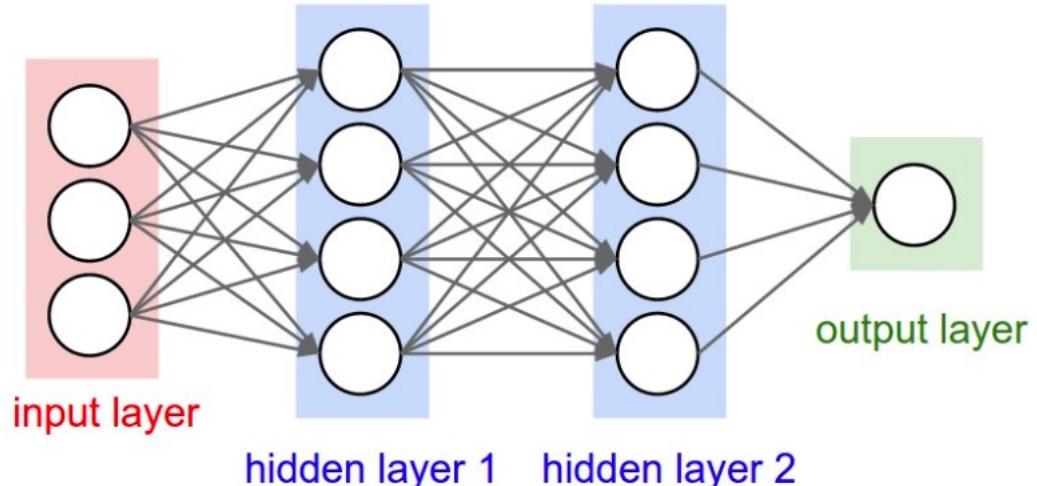
- *Adjusted during the training process*
- *The values of parameters are learned from the training data*
- *More connections will have more weights*
- *More neurons will have more bias*



Artificial Neural Network

- **Artificial Neural Network (ANN)**

- Hyper-parameters
 - The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron
- Parameters → Learning parameters
 - The weight of each connection
 - The bias of each neuron added to the weighted sum
- The **capacity** (complexity, representational power) of a network can be assessed by **the number of learning parameters**



Outline

- Classical Machine Learning vs. Deep Learning
- Artificial Neural Network
- **Implement an ANN**
- Preliminaries with PyTorch

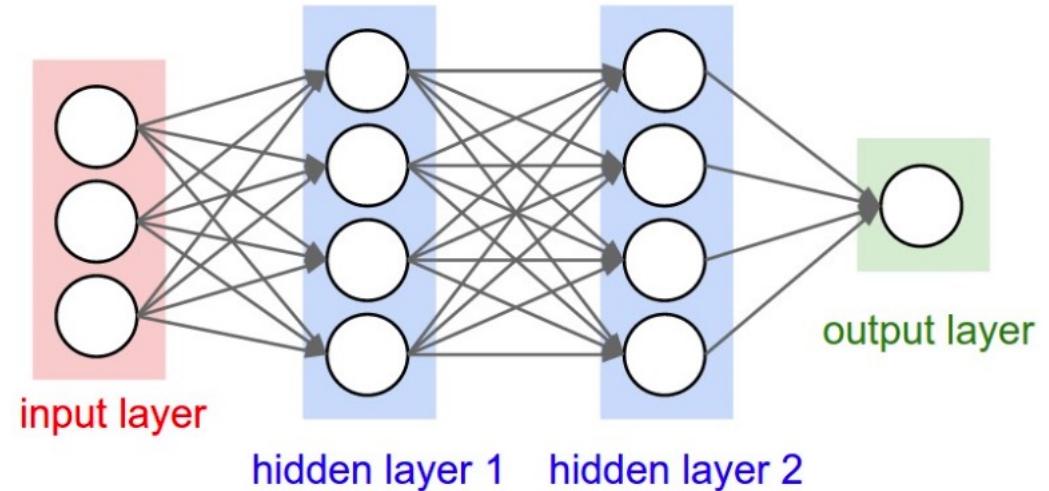
Implement an ANN

- How to **create** an ANN?
- How does ANN perform **inference**?
- How to **train** an ANN?

Implement an ANN

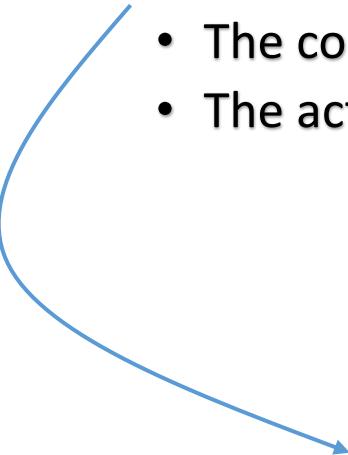
- How to **create** an ANN?
 - **Declare the hyper-parameters**
 - The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron

- *3 layers (2 hidden layer + 1 output layer)*
- *3 neurons in the input layer*
- *4 neurons in the 1st hidden layer, 4 in the 2nd*
- *1 neuron in the output layer*
- *Fully connected between layers*
- *Sigmoid activation function*

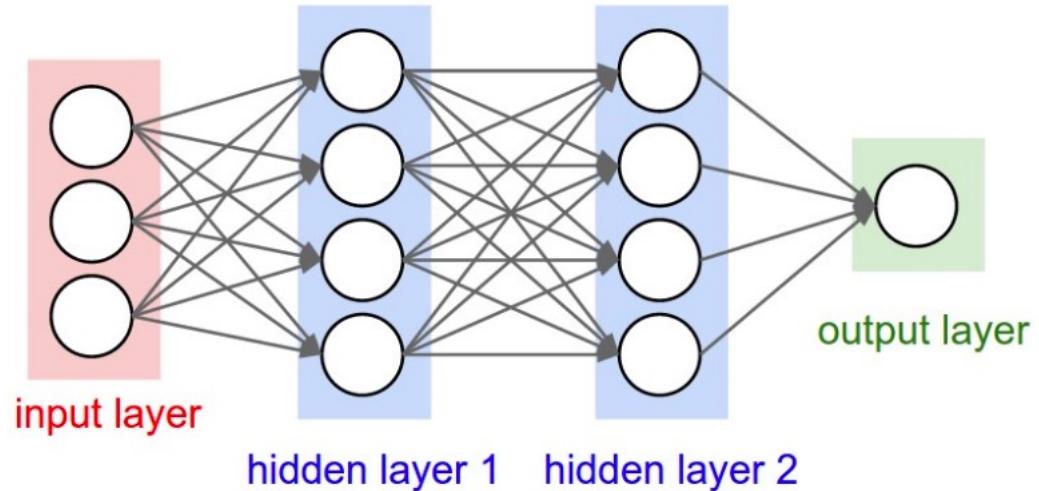


Implement an ANN

- How to **create** an ANN?
 - **Declare the hyper-parameters**
 - The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron

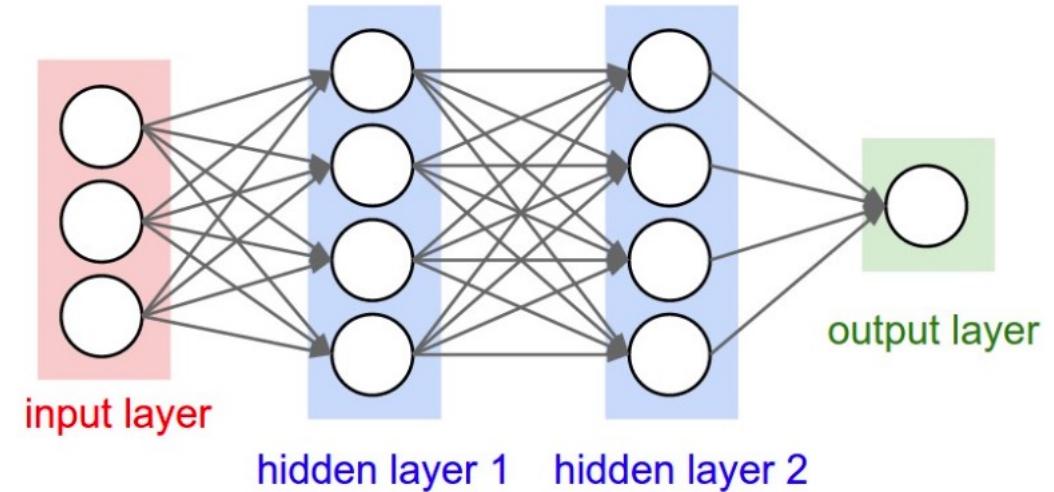
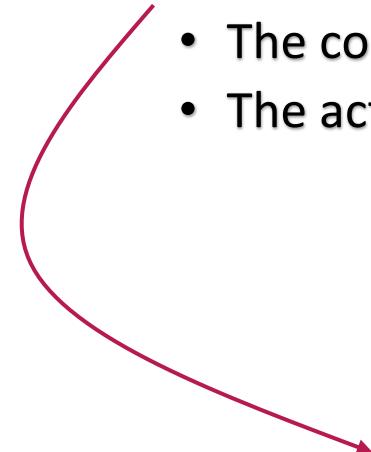


- *The structure of the ANN determines its capacity:*
 - Be careful to ensure your ANN having enough capacity given the problem to be solved
 - If your ANN is **under-fitted**, consider to **increase the capacity**
 - Add more layers → Deeper
 - Add more neurons
 - Add more connections



Implement an ANN

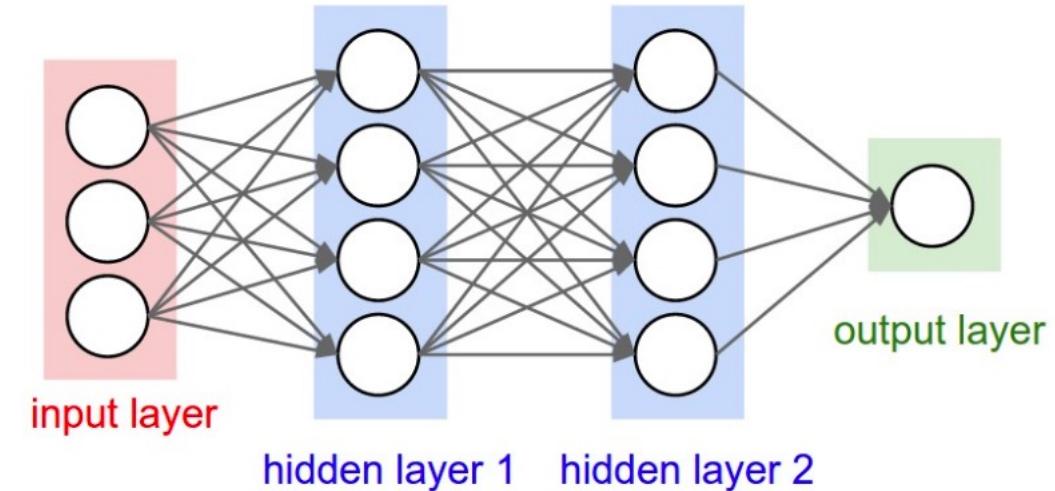
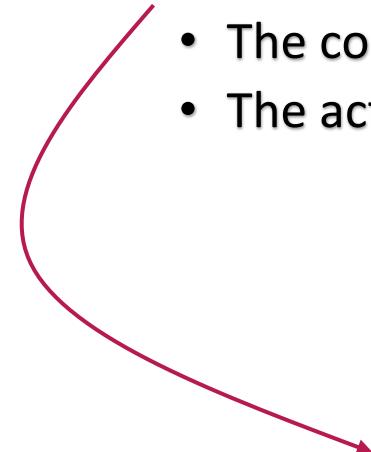
- How to **create** an ANN?
 - **Declare the hyper-parameters**
 - The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron



- *The structure of the ANN determines its capacity:*
 - Be careful to ensure your ANN having enough capacity given the problem to be solved
 - If your ANN is **over-fitted**, consider to **decrease the capacity**
 - Delete some layers → Shallower
 - Delete (or deactivate) some neurons
 - Delete some connections (from fully connected to sparse connected)

Implement an ANN

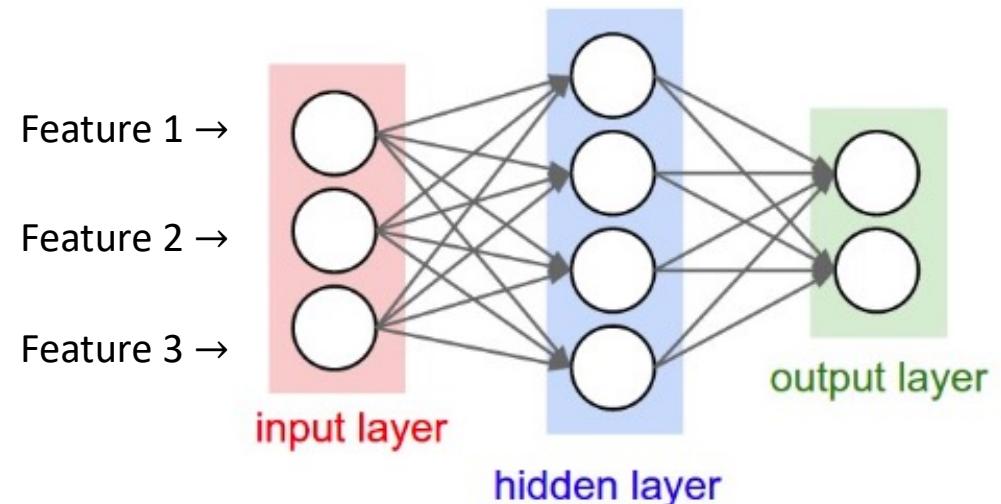
- How to **create** an ANN?
 - **Declare the hyper-parameters**
 - The number of layers
 - The number of neurons in each layer
 - The connections between neurons
 - The activation function of each neuron



- *The structure of the ANN determines its capacity:*
 - Be careful to ensure your ANN having enough capacity given the problem to be solved
 - In practice:
 - Always start with an over-fitted model which ensure your ANN having enough capacity
 - Then try to decrease the capacity by simplify the ANN
 - Or in most cases, applying other technique to prevent the ANN from overfitting
 - For example, early stopping

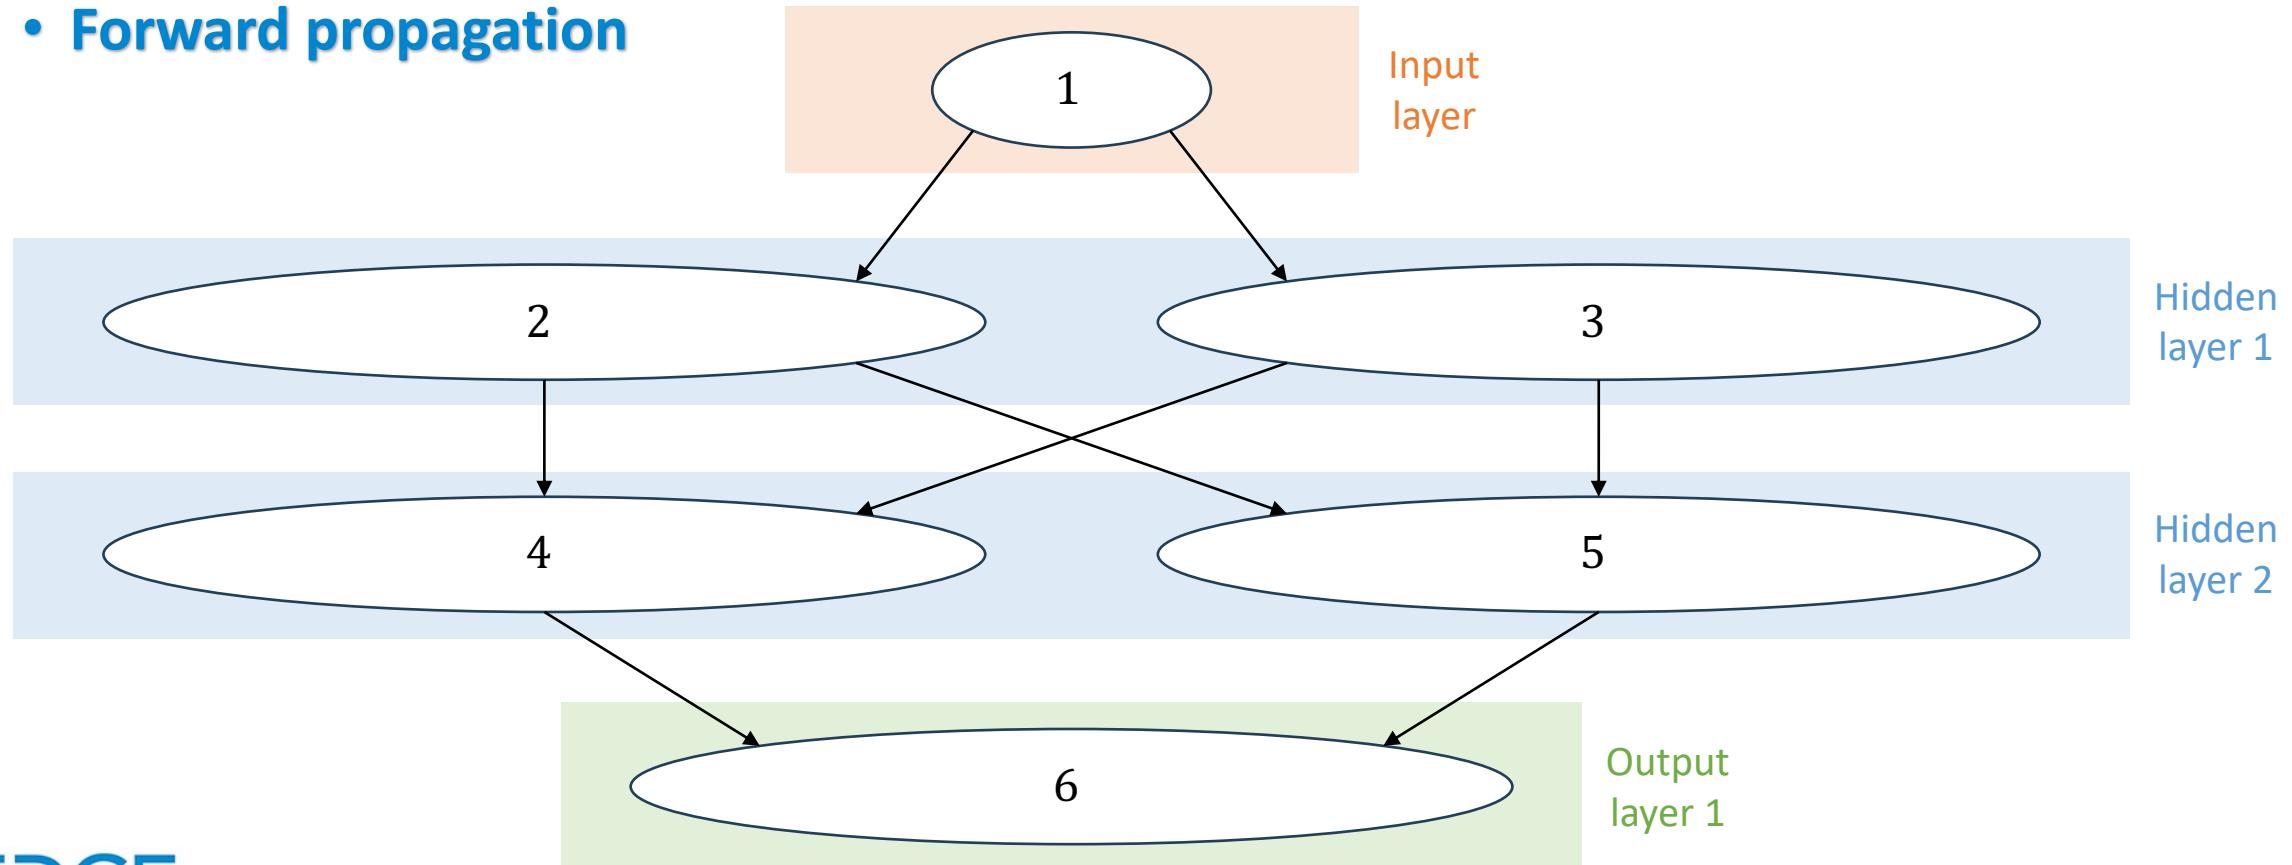
Implement an ANN

- How does ANN perform **inference**?
 - Given an input instance
 - Feed the n features to the n input nodes
 - **The number of features should equal to the number of input nodes**
 - Calculate the output of each neuron
 - Layer by layer
 - Weights, bias, activation function
 - The output of the output node is the prediction made by the ANN
 - **Forward propagation**



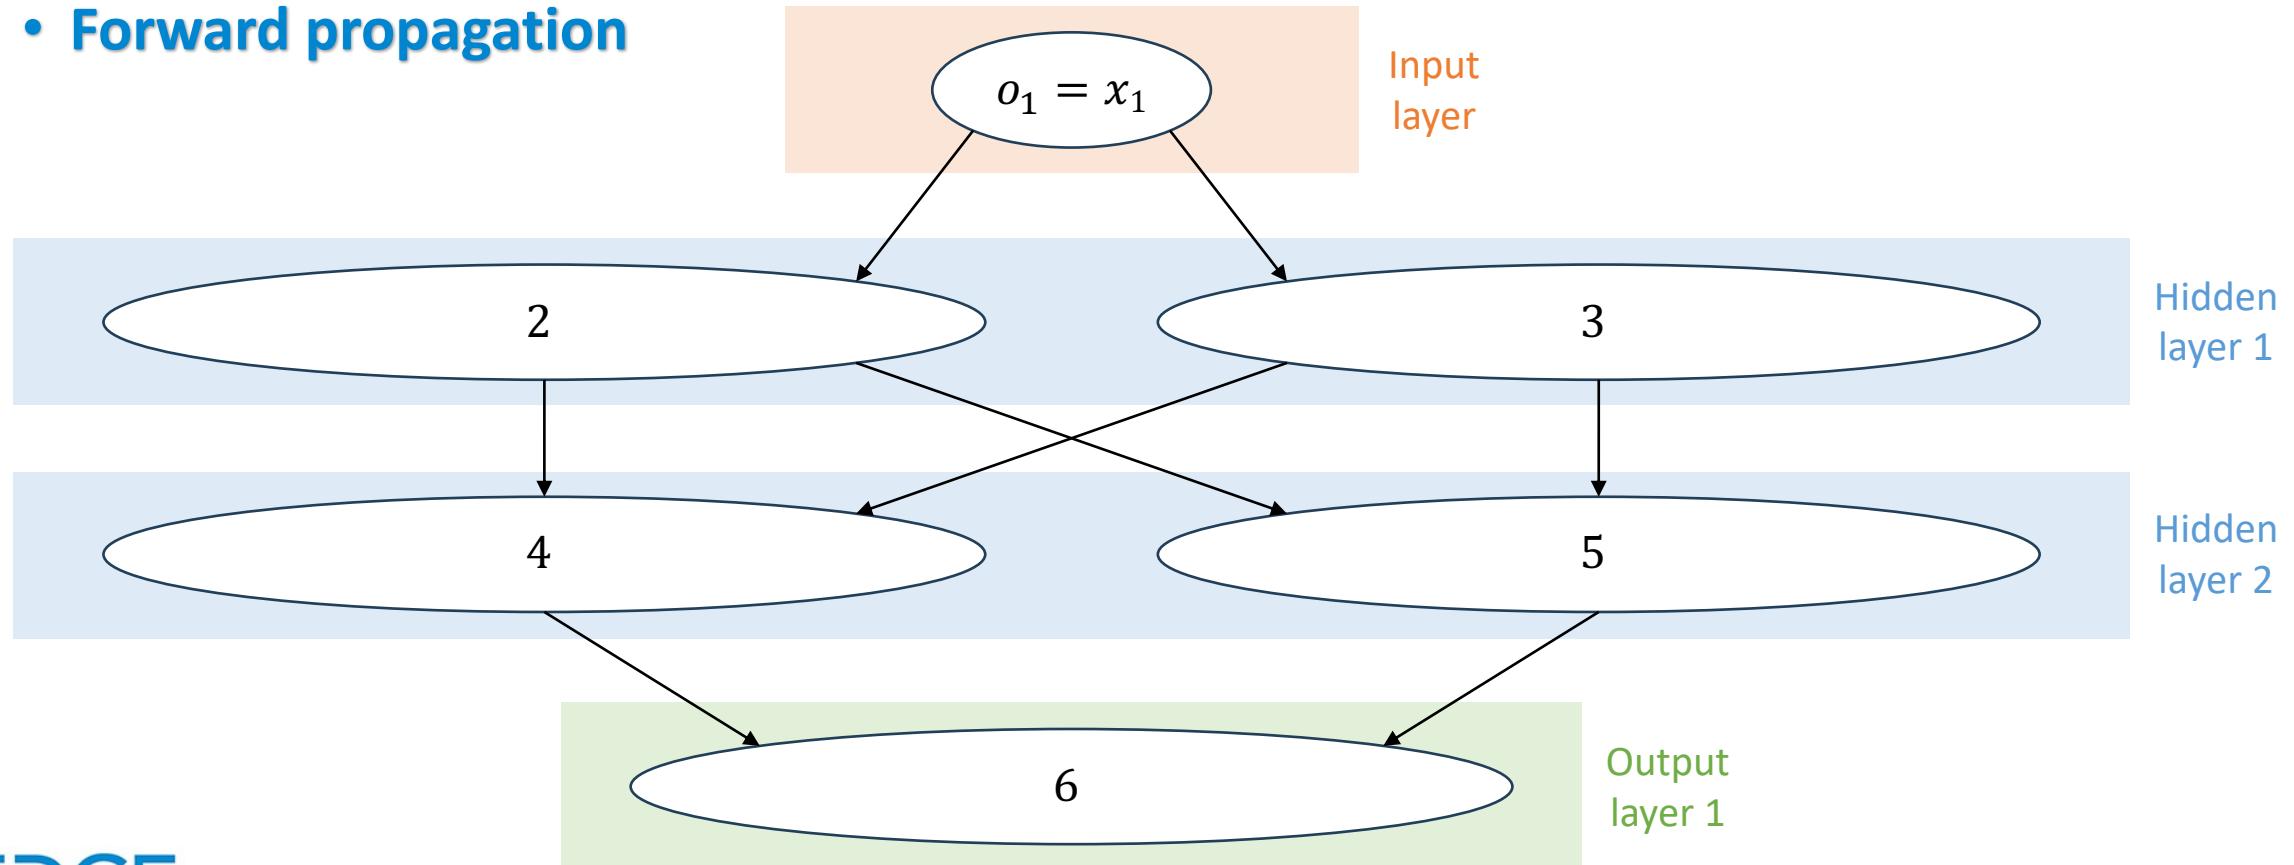
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



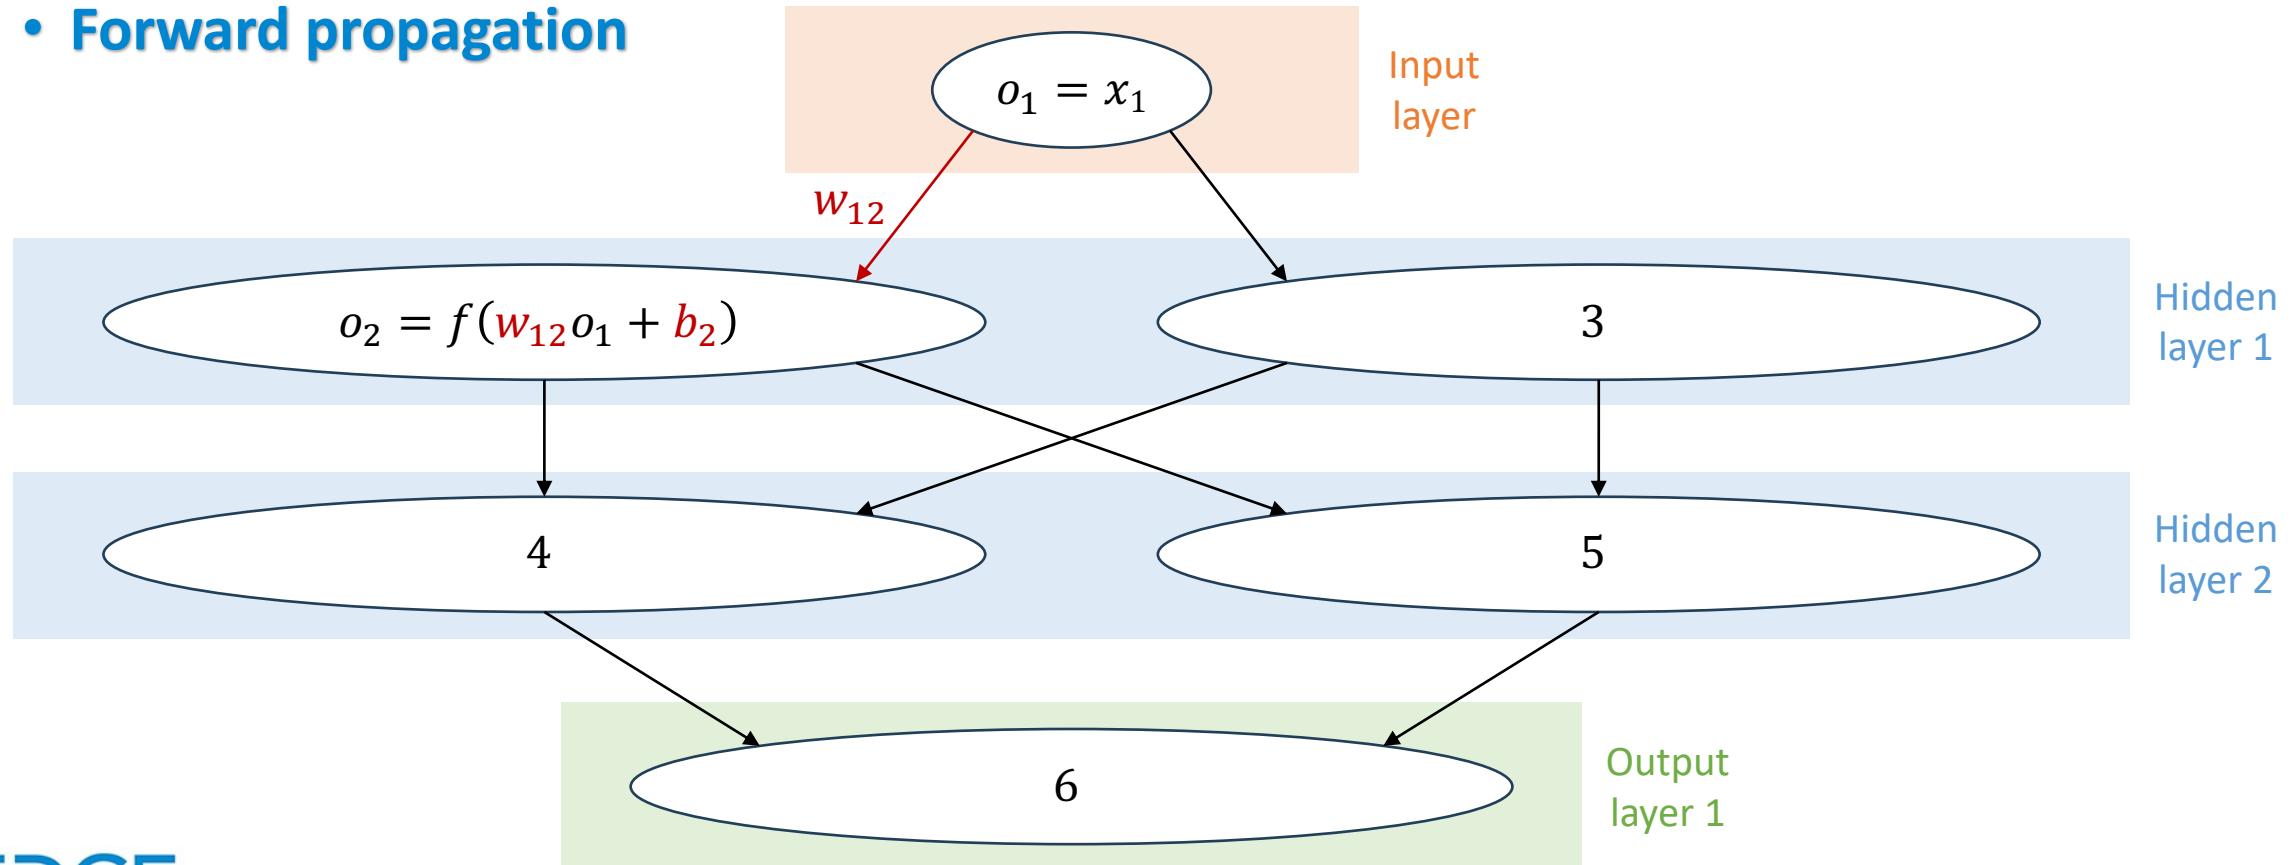
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



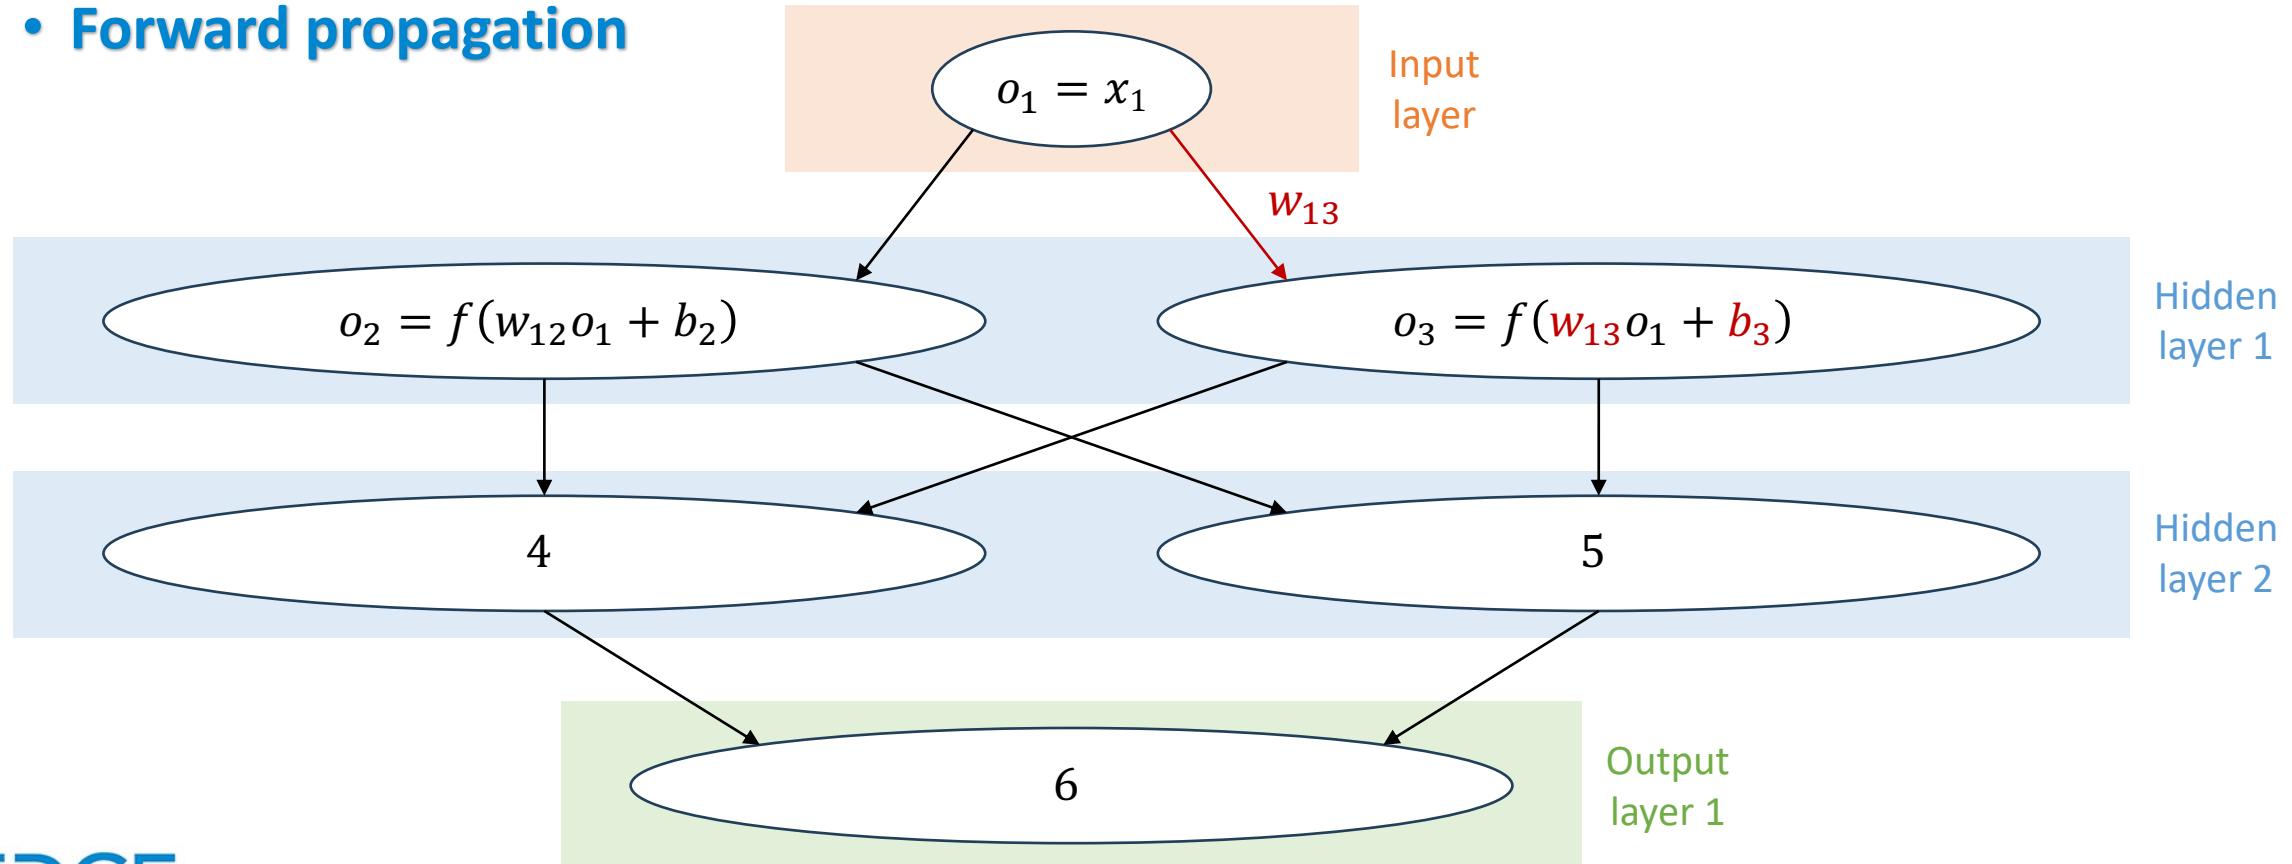
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



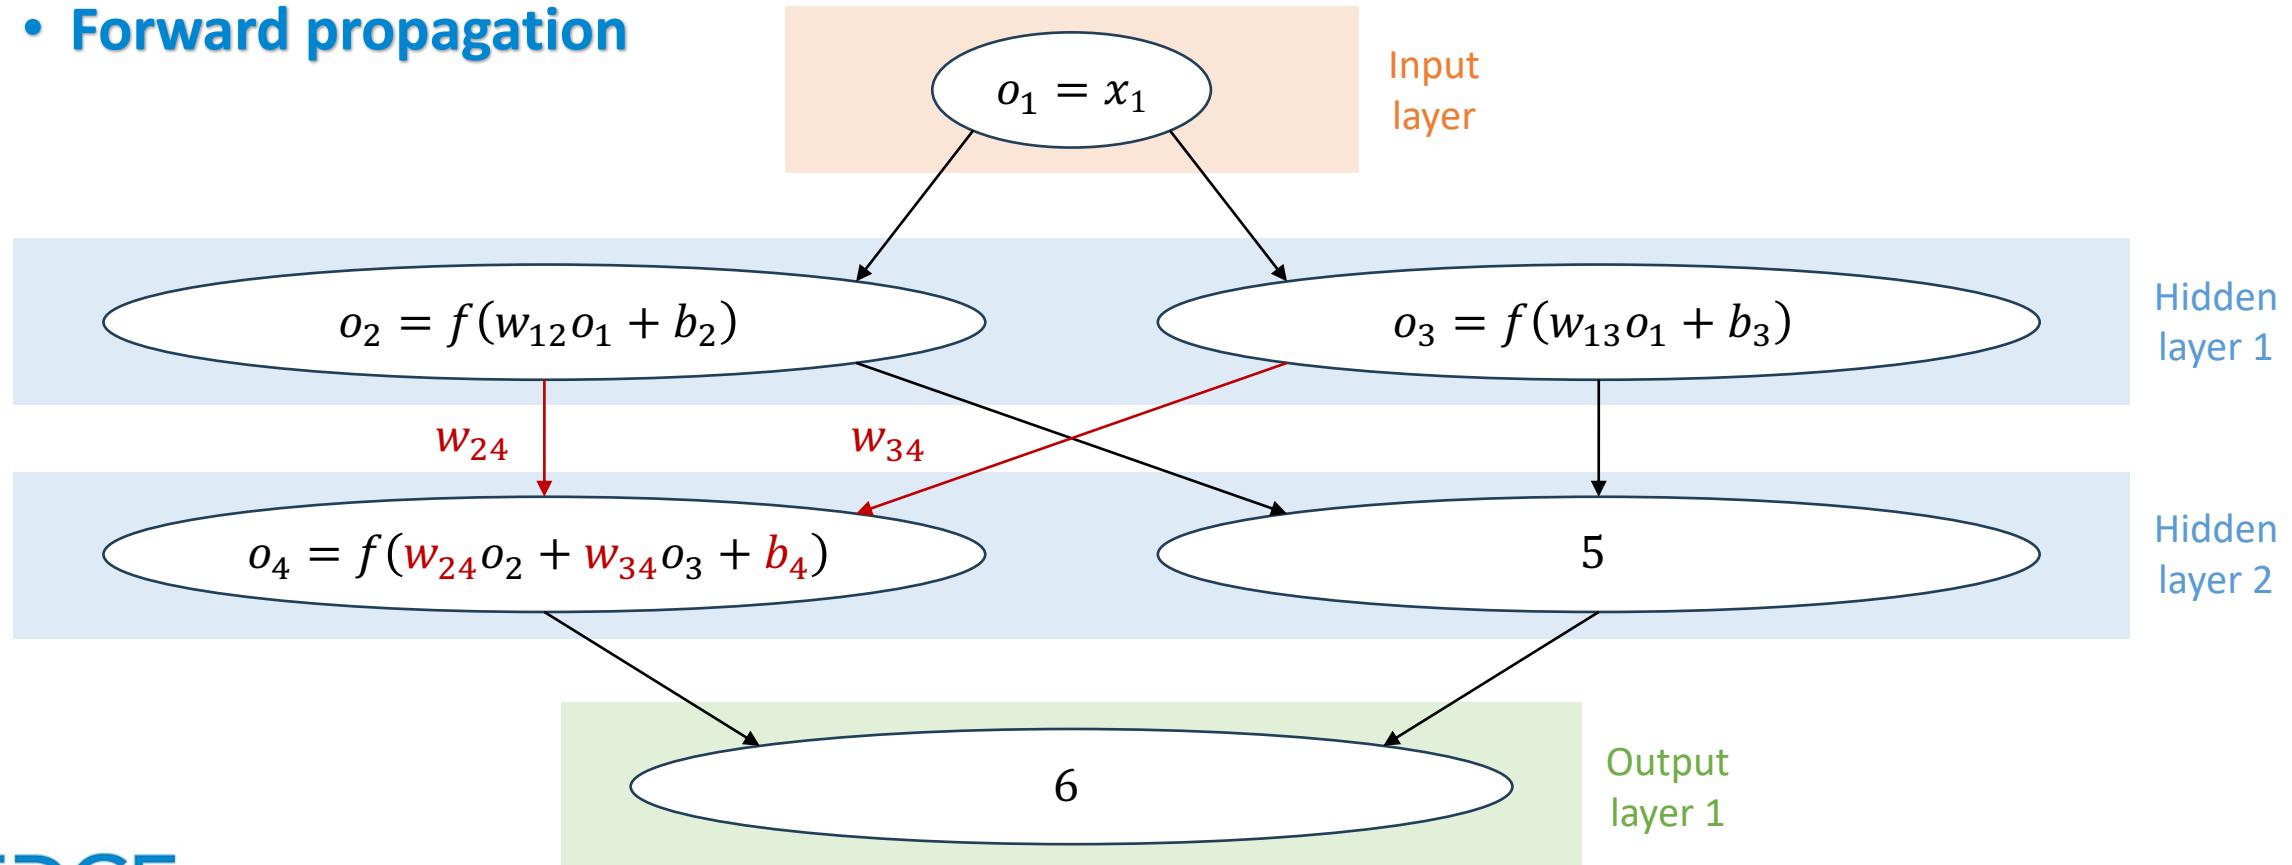
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



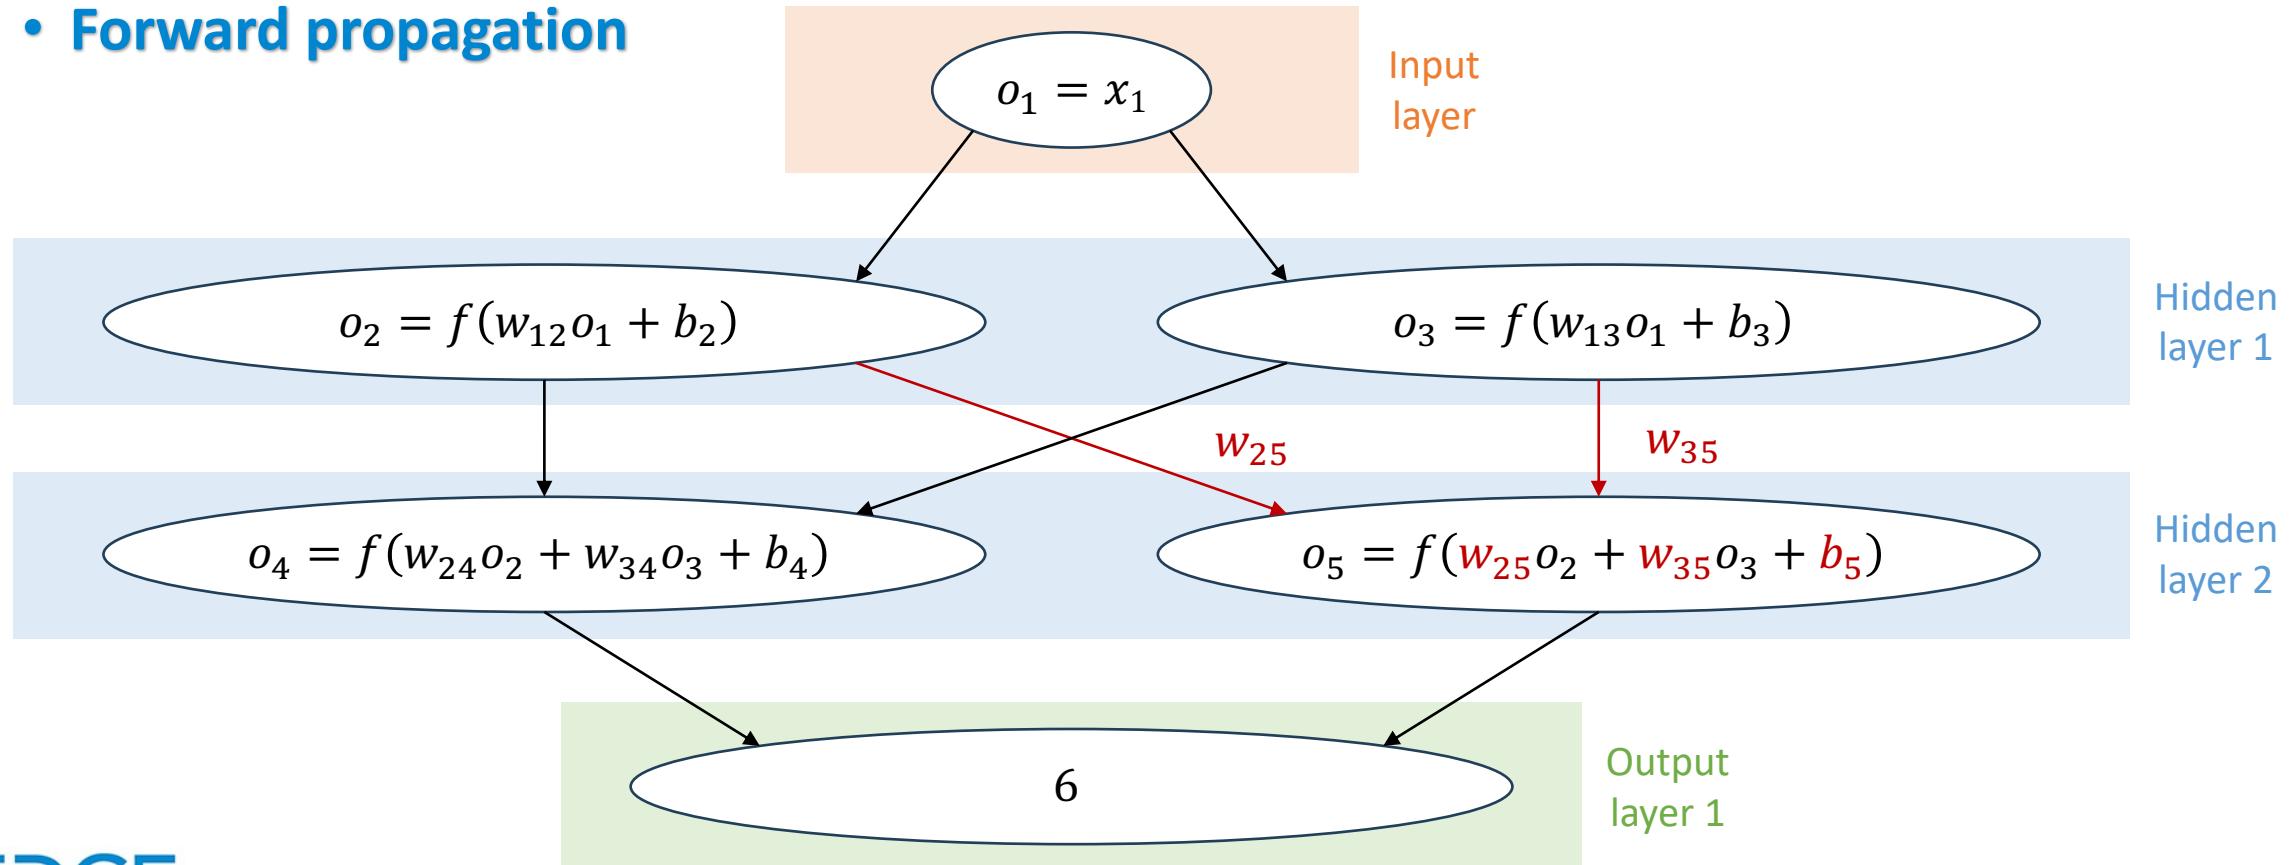
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



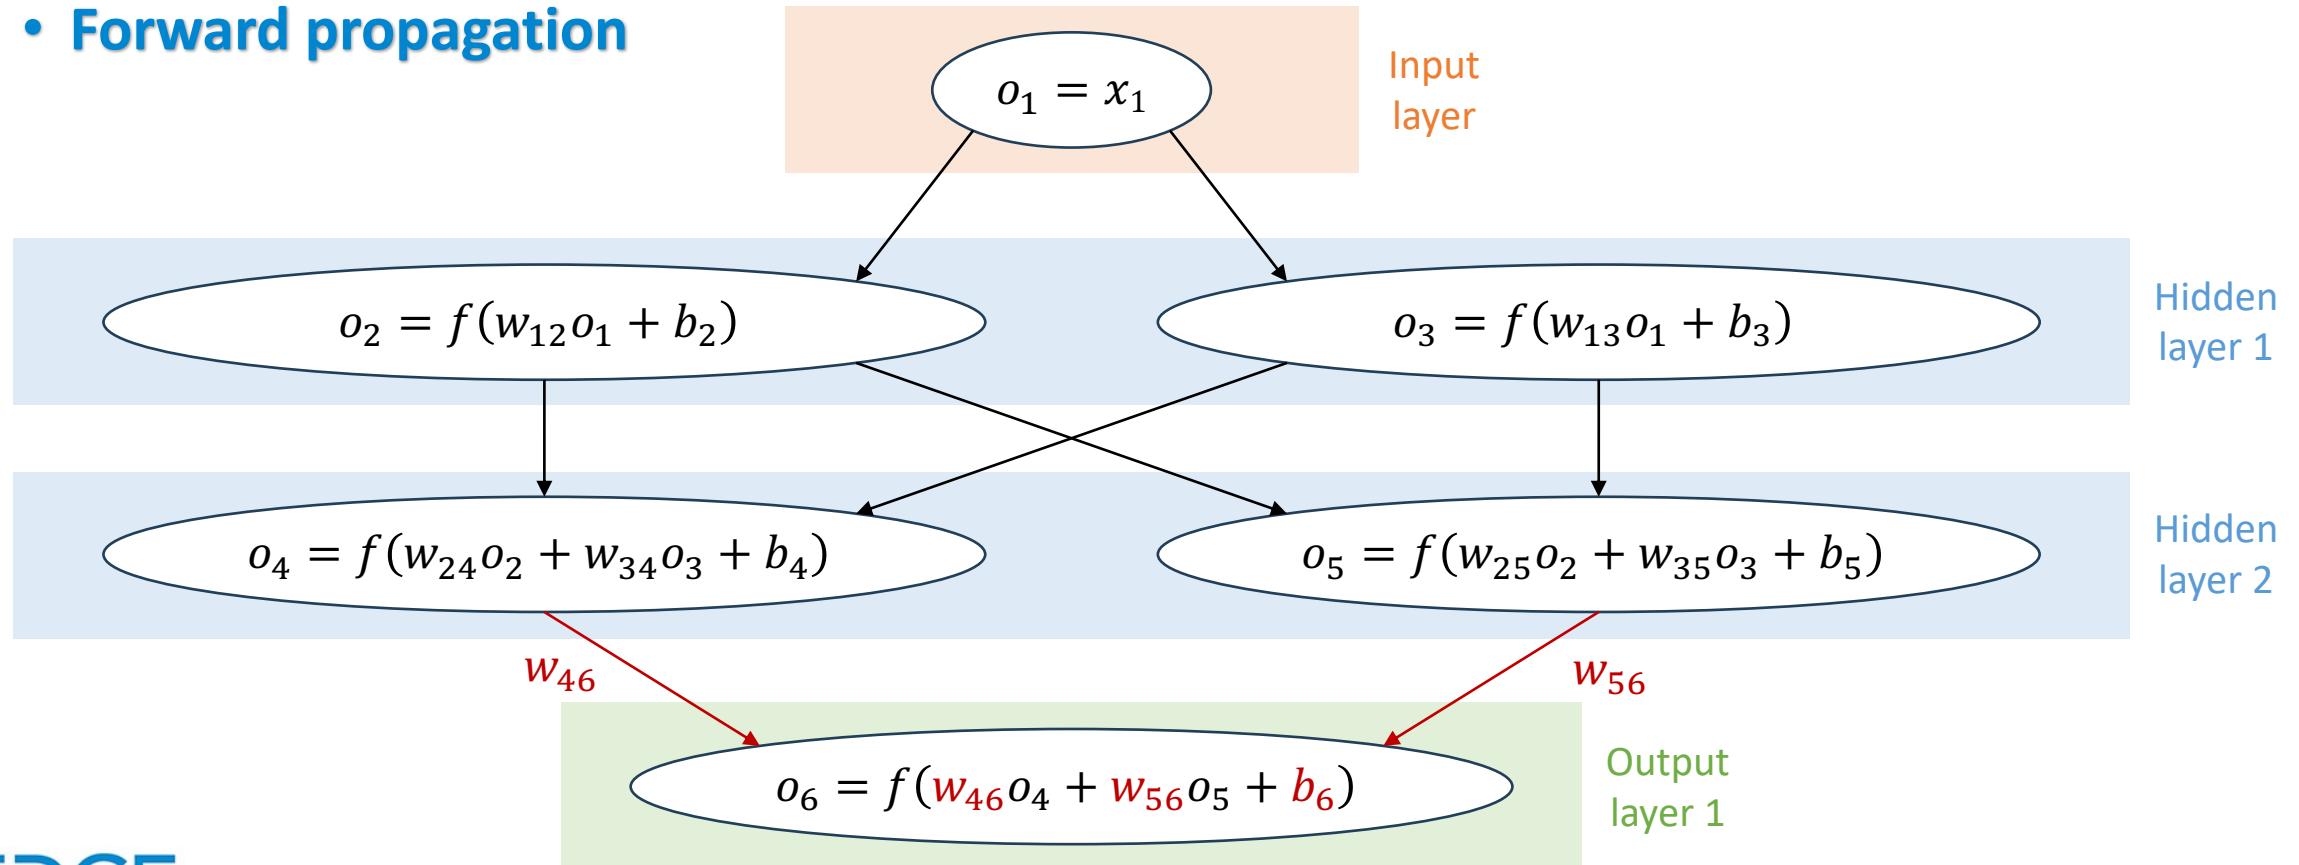
Implement an ANN

- How does ANN perform **inference**?
 - **Forward propagation**



Implement an ANN

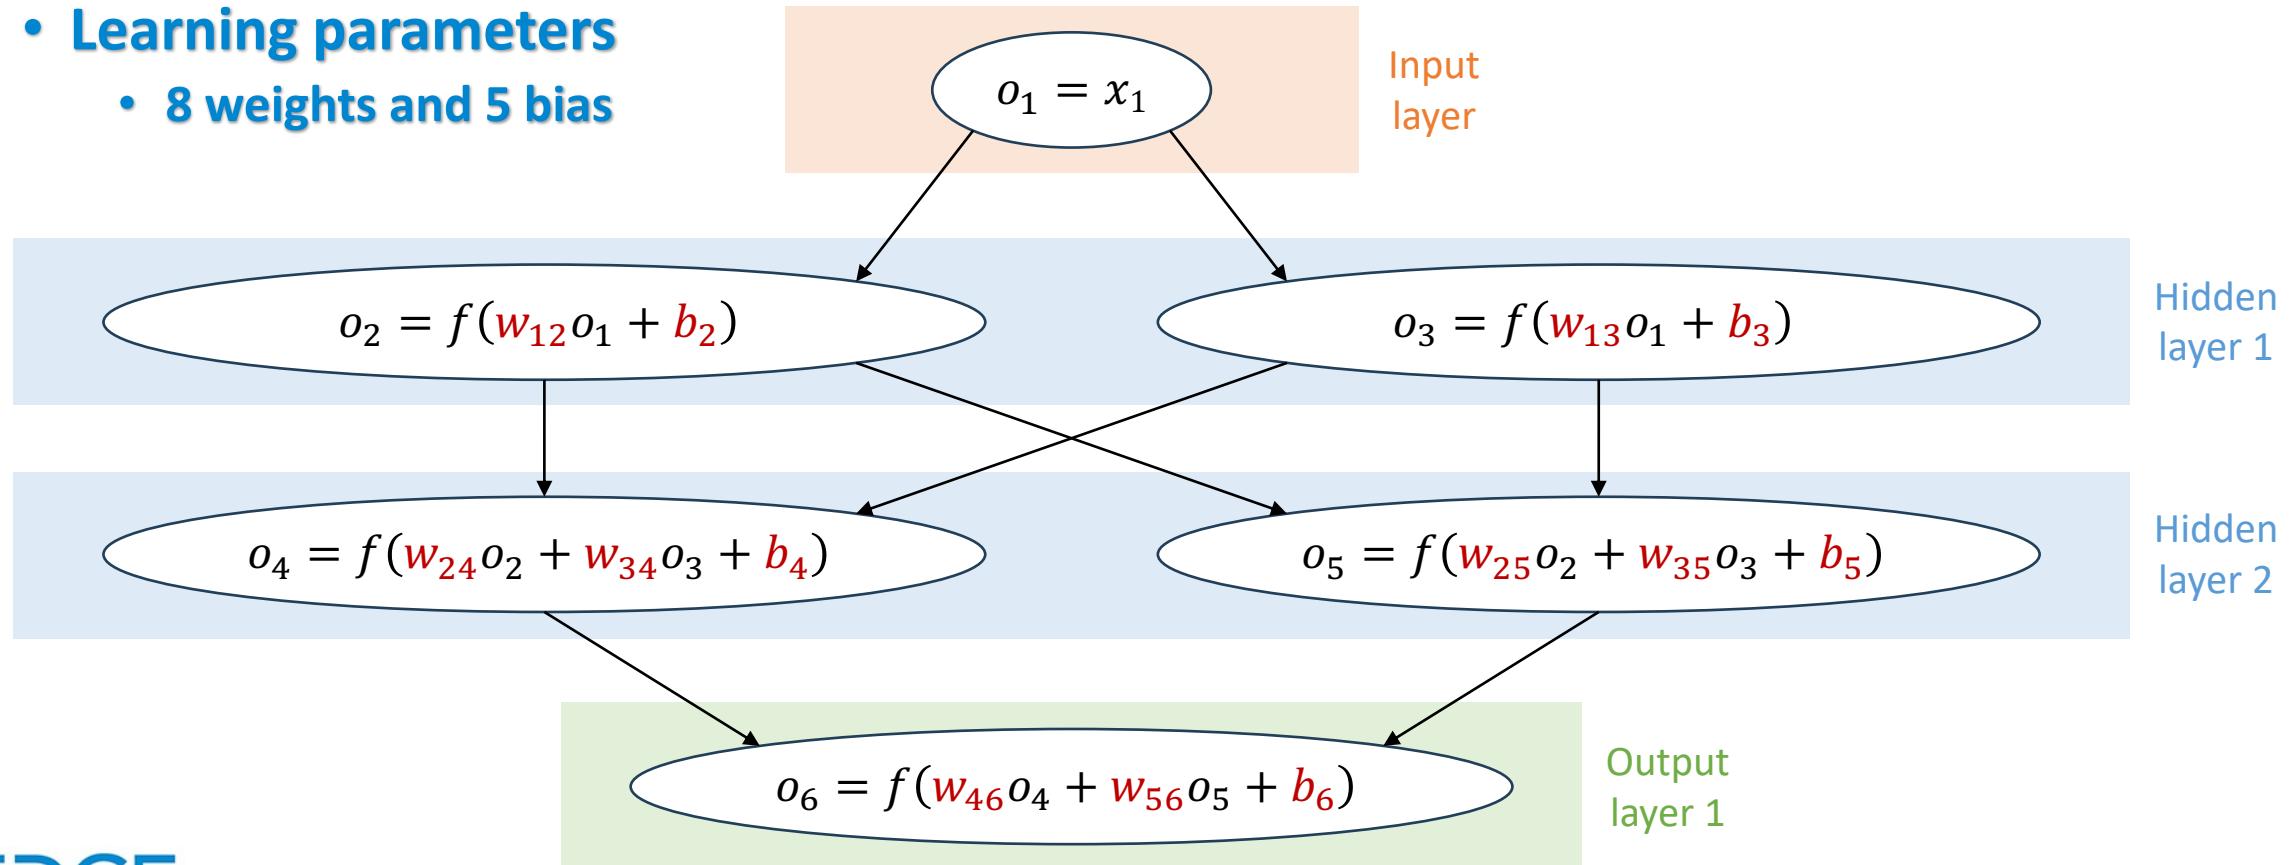
- How does ANN perform **inference**?
 - **Forward propagation**



Implement an ANN

- How does ANN perform **inference**?

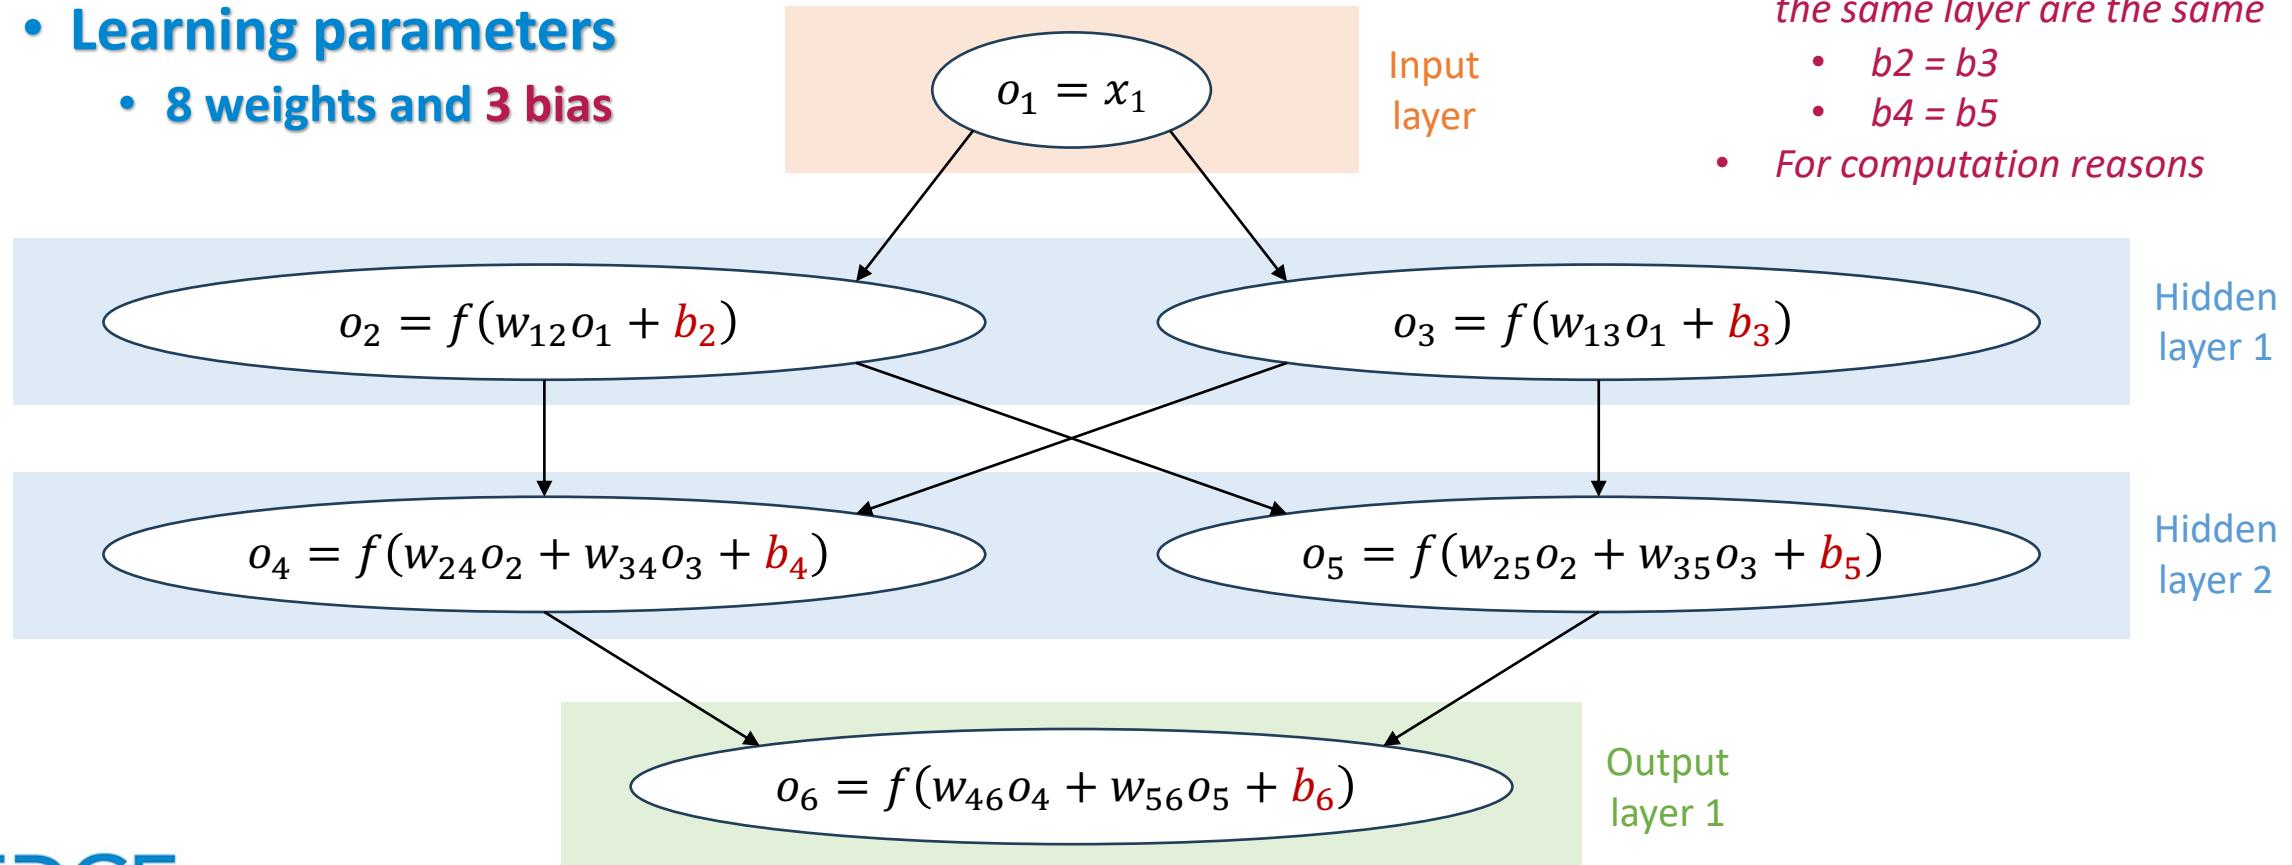
- **Learning parameters**
 - **8 weights and 5 bias**



Implement an ANN

- How does ANN perform **inference**?

- **Learning parameters**
 - **8 weights and 3 bias**

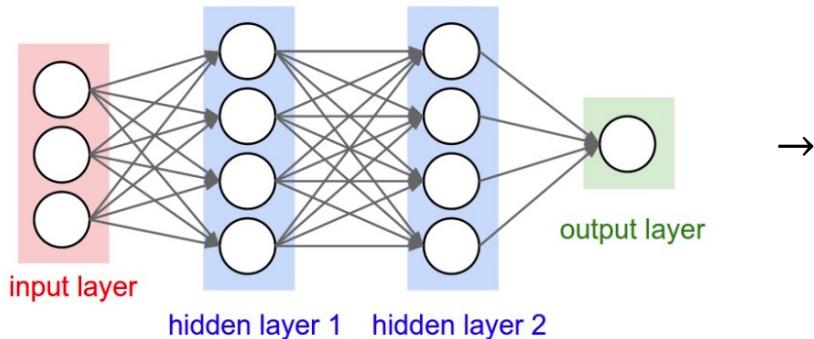


In today's practice:

- The bias values of neurons in the same layer are the same
 - $b_2 = b_3$
 - $b_4 = b_5$
- For computation reasons

Implement an ANN

- How does ANN perform **inference**?
 - Example Python scripts for a 3-layer network to perform inference



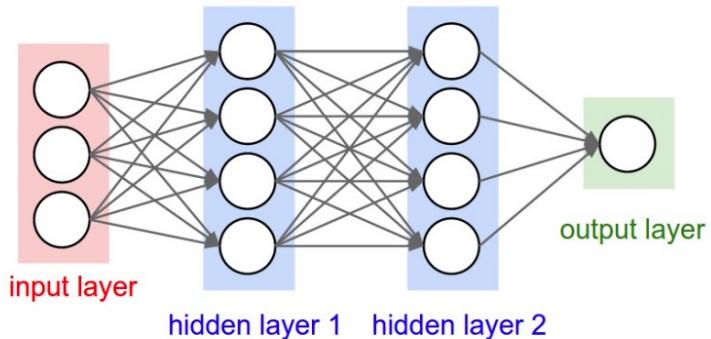
→



```
# forward propagation of a 3-layer neural network:  
f = lambda x: 1.0 / (1.0 + np.exp(-x)) # sigmoid activation function  
x = np.random.rand(3, 1) # random input features of an instance  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer outputs  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidder layer outputs  
out = f(np.dot(W3, h2) + b3) # output prediction
```

Implement an ANN

- How does ANN perform **inference**?
 - Example Python scripts for a 3-layer network to perform inference



→

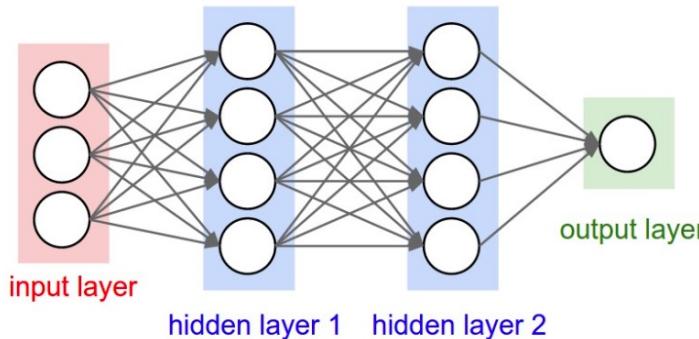
```
# forward propagation of a 3-layer neural network:  
f = lambda x: 1.0 / (1.0 + np.exp(-x)) # sigmoid activation function  
x = np.random.rand(3, 1) # random input features of an instance  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer outputs  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidden layer outputs  
out = f(np.dot(W3, h2) + b3) # output prediction
```

Activation function f:

- In the example scripts, activation functions are the same for all hidden and output neurons
- In practice:
 - Activation functions of the neurons in the same layer are the same
 - Activation functions for different layers can be different

Implement an ANN

- How does ANN perform **inference**?
 - Example Python scripts for a 3-layer network to perform inference



→

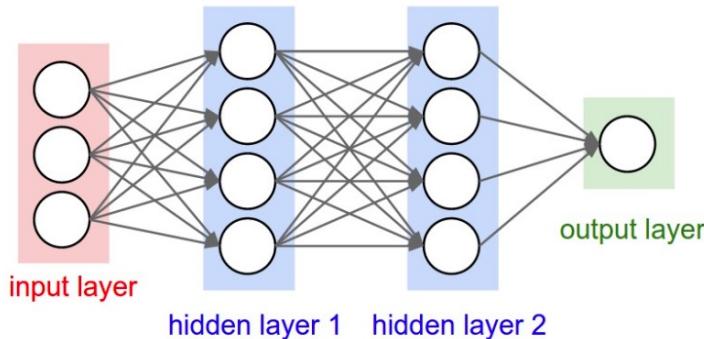


```
# forward propagation of a 3-layer neural network:  
f = lambda x: 1.0 / (1.0 + np.exp(-x)) # sigmoid activation function  
x = np.random.rand(3, 1) # random input features of an instance  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer outputs  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidder layer outputs  
out = f(np.dot(W3, h2) + b3) # output prediction
```

- np.dot()* computes the dot production of two NumPy arrays
- The calculation of forward propagation can be implemented efficiently using *matrix operation*
 - Matrix operation can be *parallelized* → Use *GPU* to accelerate

Implement an ANN

- How does ANN perform **inference**?
 - Example Python scripts for a 3-layer network to perform inference



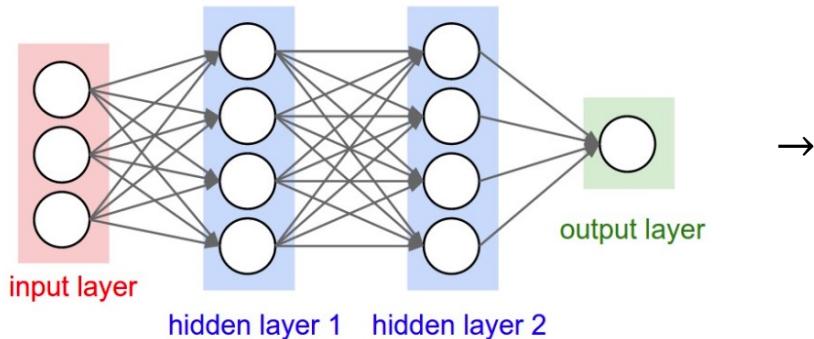
```
● ● ●  
# forward propagation of a 3-layer neural network:  
f = lambda x: 1.0 / (1.0 + np.exp(-x)) # sigmoid activation function  
x = np.random.rand(3, 1) # random input features of an instance  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer outputs  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidder layer outputs  
out = f(np.dot(W3, h2) + b3) # output prediction
```

Be careful of matrix operation, the dimensions should be compatible:

shape(x) = (3, 1)
shape(h₁) = (4, 1)
shape(h₂) = (4, 1)
shape(out) = (1, 1)

Implement an ANN

- How does ANN perform **inference**?
 - Example Python scripts for a 3-layer network to perform inference



```
● ● ●  
# forward propagation of a 3-layer neural network:  
f = lambda x: 1.0 / (1.0 + np.exp(-x)) # sigmoid activation function  
x = np.random.rand(3, 1) # random input features of an instance  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer outputs  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidder layer outputs  
out = f(np.dot(W3, h2) + b3) # output prediction
```

Be careful of matrix operation, the dimensions should be compatible:

$\text{shape}(W_1) = (4, 3)$

$\text{shape}(x) = (3, 1)$

$\text{shape}(W_2) = (4, 4)$

$\text{shape}(h_1) = (4, 1)$

$\text{shape}(W_3) = (4, 1)$

$\text{shape}(h_2) = (4, 1)$

$\text{shape}(out) = (1, 1)$

$\text{shape}(\text{out}) = (1, 1)$

Implement an ANN

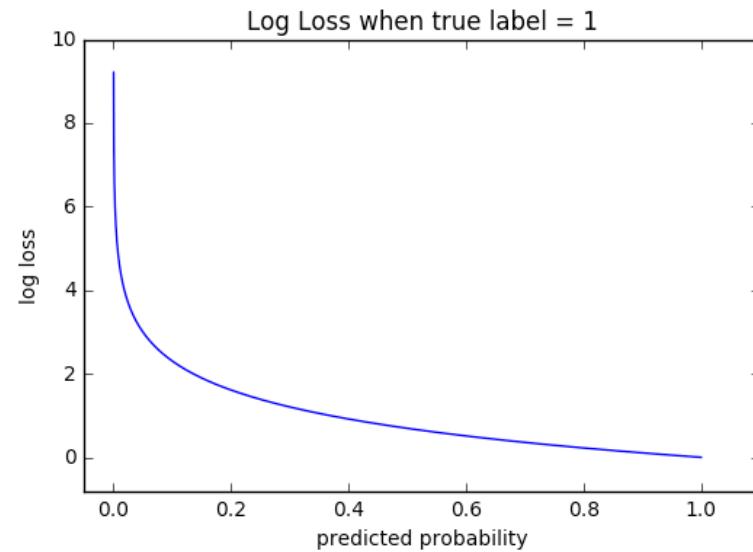
- How to **create** an ANN?
 - Number of layers, neurons and connections, activation functions
- How does ANN perform **inference**?
 - Forward propagation
- How to **train** an ANN?
 - **Learning process** → Decide the values of **learning parameters**
 - **Loss / Objective function**
 - A formal measure of how good (or bad) the models are.
 - **Optimization algorithm**
 - Search for the best values of learning parameters for minimizing the loss function

Implement an ANN

- How to **train** an ANN?
 - **Learning process** → Decide the values of **learning parameters**
 - **Loss / Objective function**
 - A formal measure of how good (or bad) the models are.
 - ‘Loss’ means lower value is better
 - Regression
 - **MSE, RMSE, MAE**
 - Classification
 - **Cross-entropy loss** (also called log loss)

Implement an ANN

- How to **train** an ANN?
 - **Learning process** → Decide the values of **learning parameters**
 - **Cross-entropy loss** (also called log loss)
 - Measures the performance of a classifier whose output is a probability value between 0 and 1
 - Lower is better
 - For binary classification, one instance
 - $L = -(y \log(p) + (1 - y) \log(1 - p))$
 - For M class, one instance
 - $L = -\sum_{c=1}^M y_c \log(p_c)$
 - For M class, N instance
 - $L = -\sum_{i=1}^N \sum_{c=1}^M y_{ci} \log(p_{ci})$



Implement an ANN

- How to **train** an ANN?
 - **Learning process** → Decide the values of **learning parameters**
 - **Loss / Objective function**
 - **Optimization algorithm**
 - Search for the best values of learning parameters for minimizing the loss function
 - **Gradient descent**

Outline

- Classical Machine Learning vs. Deep Learning
- Artificial Neural Network
- Implement an ANN
- **Preliminaries with PyTorch**

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - For Mac

PyTorch Build	Stable (2.2.1)			Preview (Nightly)	
Your OS	Linux	Mac	Windows		
Package	Conda	Pip	LibTorch	Source	
Language	Python				C++ / Java
Compute Platform	CUDA 11.8	CUDA 12.1	ROCM 5.7	Default	
Run this Command:	<pre>conda install pytorch::pytorch torchvision torchaudio -c pytorch</pre>				

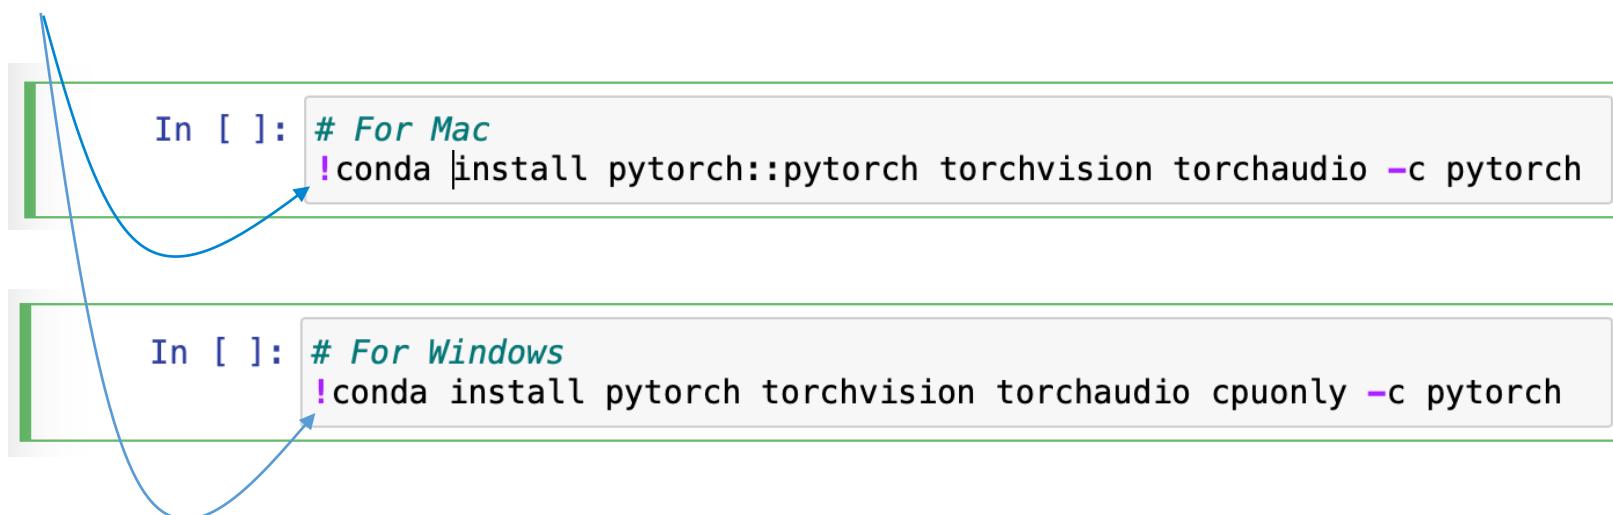
Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - For Windows

PyTorch Build	Stable (2.2.1)	Preview (Nightly)
Your OS	Linux	Mac
Package	Conda	Pip
Language	Python	C++ / Java
Compute Platform	CUDA 11.8	ROCM 5.7
Run this Command:	<pre>conda install pytorch torchvision torchaudio cpuonly -c pytorch</pre>	

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - Execute the corresponding command in a Jupyter Notebook Cell
 - Add ! in front of the command



```
In [ ]: # For Mac  
!conda install pytorch::pytorch torchvision torchaudio -c pytorch
```



```
In [ ]: # For Windows  
!conda install pytorch torchvision torchaudio ccpuonly -c pytorch
```

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - Execute the corresponding command in a Jupyter Notebook Cell
 - Add ! in front of the command

```
In [*]: # For Mac
!conda install pytorch::pytorch torchvision torchaudio -c pytorch
```

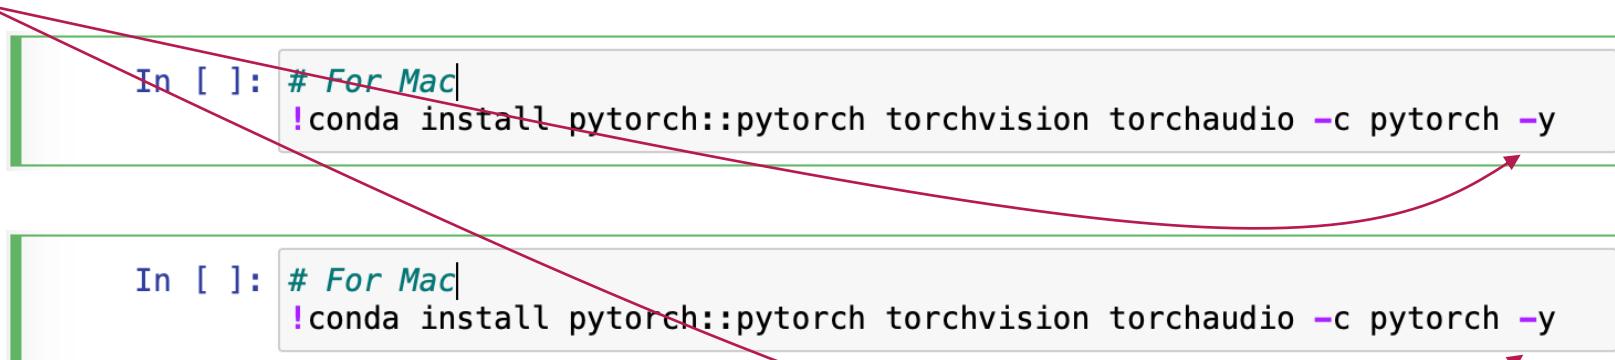
Dependency	Path
mpc	pkgs/main/osx-arm64::mpc-1.1.0-h8c48613_1
mpfr	pkgs/main/osx-arm64::mpfr-4.0.2-h695f6f0_1
mpmath	pkgs/main/osx-arm64::mpmath-1.3.0-py311hca03da5_0
networkx	pkgs/main/osx-arm64::networkx-3.1-py311hca03da5_0
numpy	pkgs/main/osx-arm64::numpy-1.26.4-py311he598dae_0
numpy-base	pkgs/main/osx-arm64::numpy-base-1.26.4-py311hfbfe69c_0
openjpeg	pkgs/main/osx-arm64::openjpeg-2.3.0-h7a6adac_2
pillow	pkgs/main/osx-arm64::pillow-10.2.0-py311h80987f9_0
pysocks	pkgs/main/osx-arm64::pysocks-1.7.1-py311hca03da5_0
pytorch	pytorch/osx-arm64::pytorch-2.2.1-py3.11_0
requests	pkgs/main/osx-arm64::requests-2.31.0-py311hca03da5_1
sympy	pkgs/main/osx-arm64::sympy-1.12-py311hca03da5_0
torchaudio	pytorch/osx-arm64::torchaudio-2.2.1-py311_cpu
torchvision	pytorch/osx-arm64::torchvision-0.17.1-py311_cpu
urllib3	pkgs/main/osx-arm64::urllib3-2.1.0-py311hca03da5_1
zstd	pkgs/main/osx-arm64::zstd-1.5.5-hd90d995_0

Proceed ([y]/n)?

There is a double check to run this command

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - Execute the corresponding command in a Jupyter Notebook Cell
 - Add **!** in front of the command
 - Add **-y** at the end of the command to skip the double check



```
In [ ]: # For Mac  
!conda install pytorch::pytorch torchvision torchaudio -c pytorch -y
```

```
In [ ]: # For Mac  
!conda install pytorch::pytorch torchvision torchaudio -c pytorch -y
```

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - Execute the corresponding command in a Jupyter Notebook Cell
 - Add **!** in front of the command
 - Add **-y** at the end of the command to skip the double check

You only need to install once.

```
In [1]: # For Mac
!conda install pytorch::pytorch torchvision torchaudio -c pytorch -y

pytorch-2.2.1      | 55.5 MB  | #####
torchaudio-2.2.1   | 5.0 MB    | #####
pytorch-2.2.1      | 55.5 MB  | #####
torchaudio-2.2.1   | 5.0 MB    | #####
pytorch-2.2.1      | 55.5 MB  | #####
pytorch-2.2.1      | 55.5 MB  | #####
pytorch-2.2.1      | 55.5 MB  | #####
pytorch-2.2.1      | 55.5 MB  | #####
pytorch-2.2.1      | 55.5 MB  | #####
torchvision-0.17.1  | 6.8 MB    | #####
```

Preparing transaction: done
Verifying transaction: done
Executing transaction: done

These 3 lines indicate you successfully install PyTorch

Preliminaries with PyTorch

- Install PyTorch use Conda
 - <https://pytorch.org/get-started/locally/>
 - Execute the corresponding command in a Jupyter Notebook Cell
 - Add **!** in front of the command
 - Add **-y** at the end of the command to skip the double check
 - Verify whether PyTorch is installed correctly

```
In [2]: import torch
x = torch.rand(5, 3)
print(x)

tensor([[0.6496, 0.9073, 0.0164],
        [0.2405, 0.4017, 0.6177],
        [0.1217, 0.6217, 0.0389],
        [0.4661, 0.9980, 0.1216],
        [0.9674, 0.0551, 0.5011]])
```

Reading Assignment

- Reading assignments
 - <https://d2l.ai/>
 - Chapter 2. Preliminaries
 - 2.1 Data Manipulation
 - 2.3 Linear Algebra
 - 2.4 Calculus
 - 2.5 Automatic Differentiation
 - Read the content and execute the code in the Jupyter Notebook