

ARTIFICIAL INTELLIGENCE NEEDS REAL INTELLIGENCE

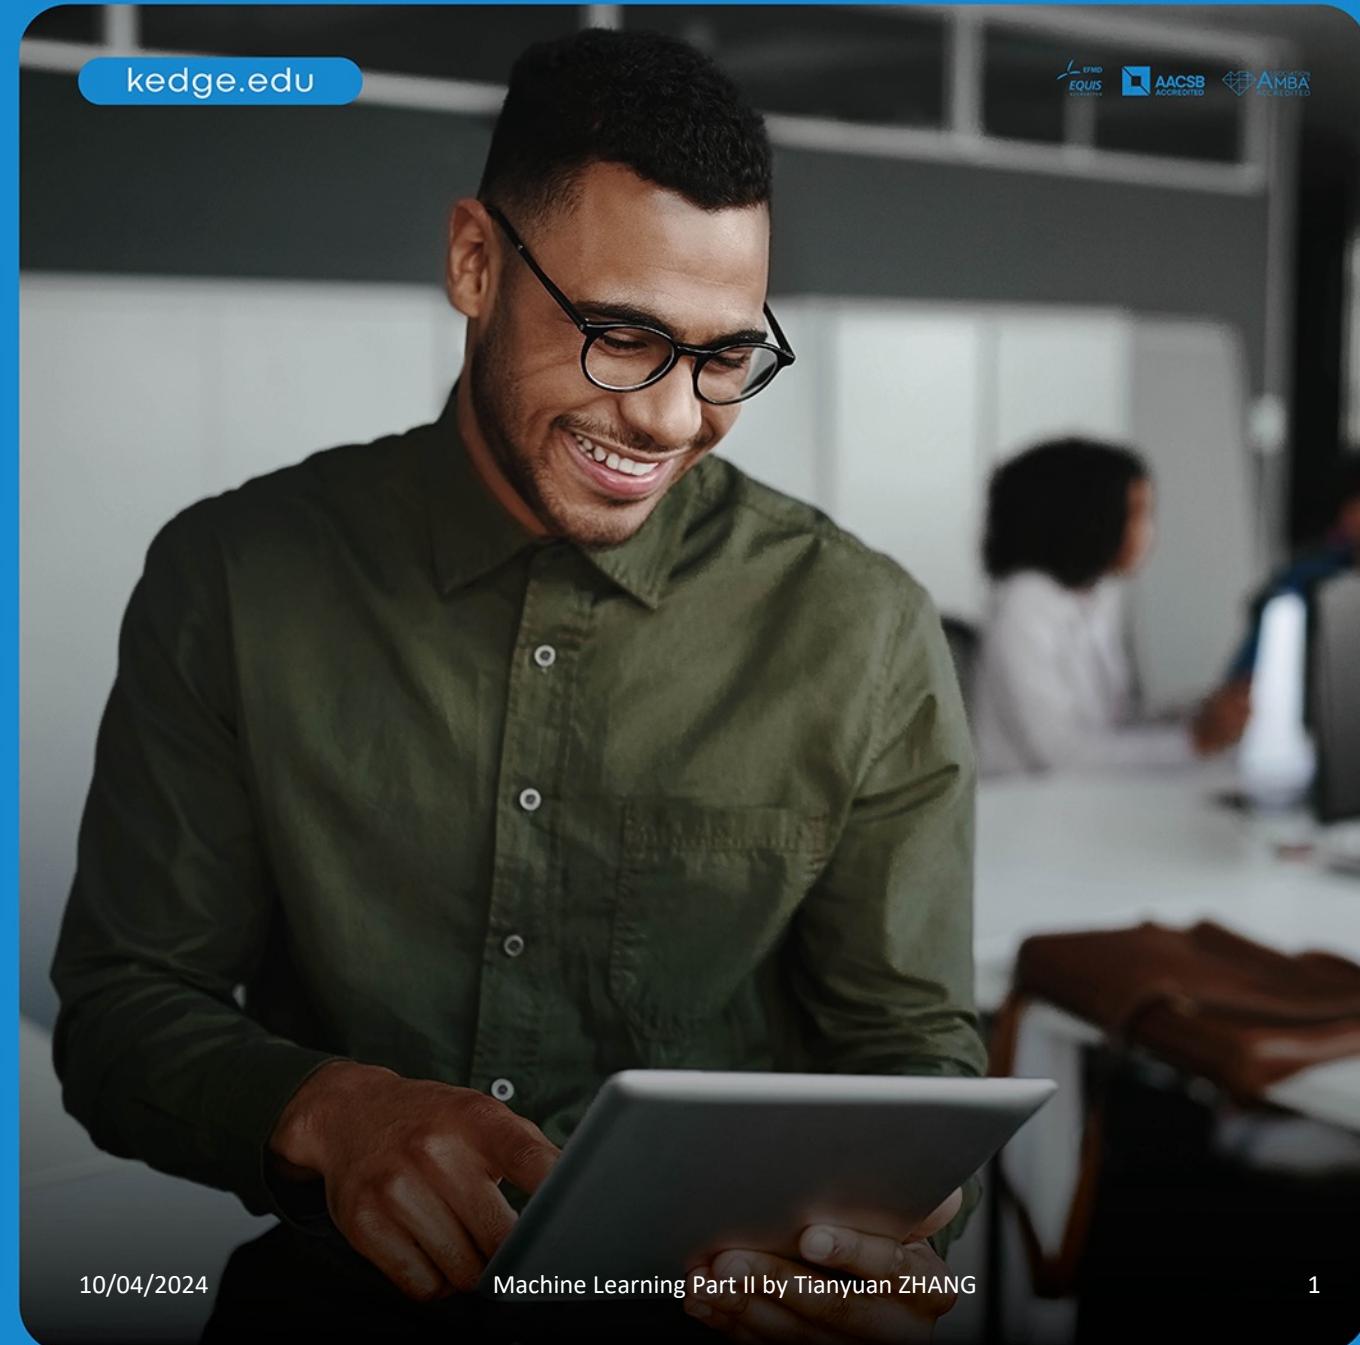
Convolutional Neural Network II

Professor: Tianyuan ZHANG
tianyuan.zhang@kedgebs.com



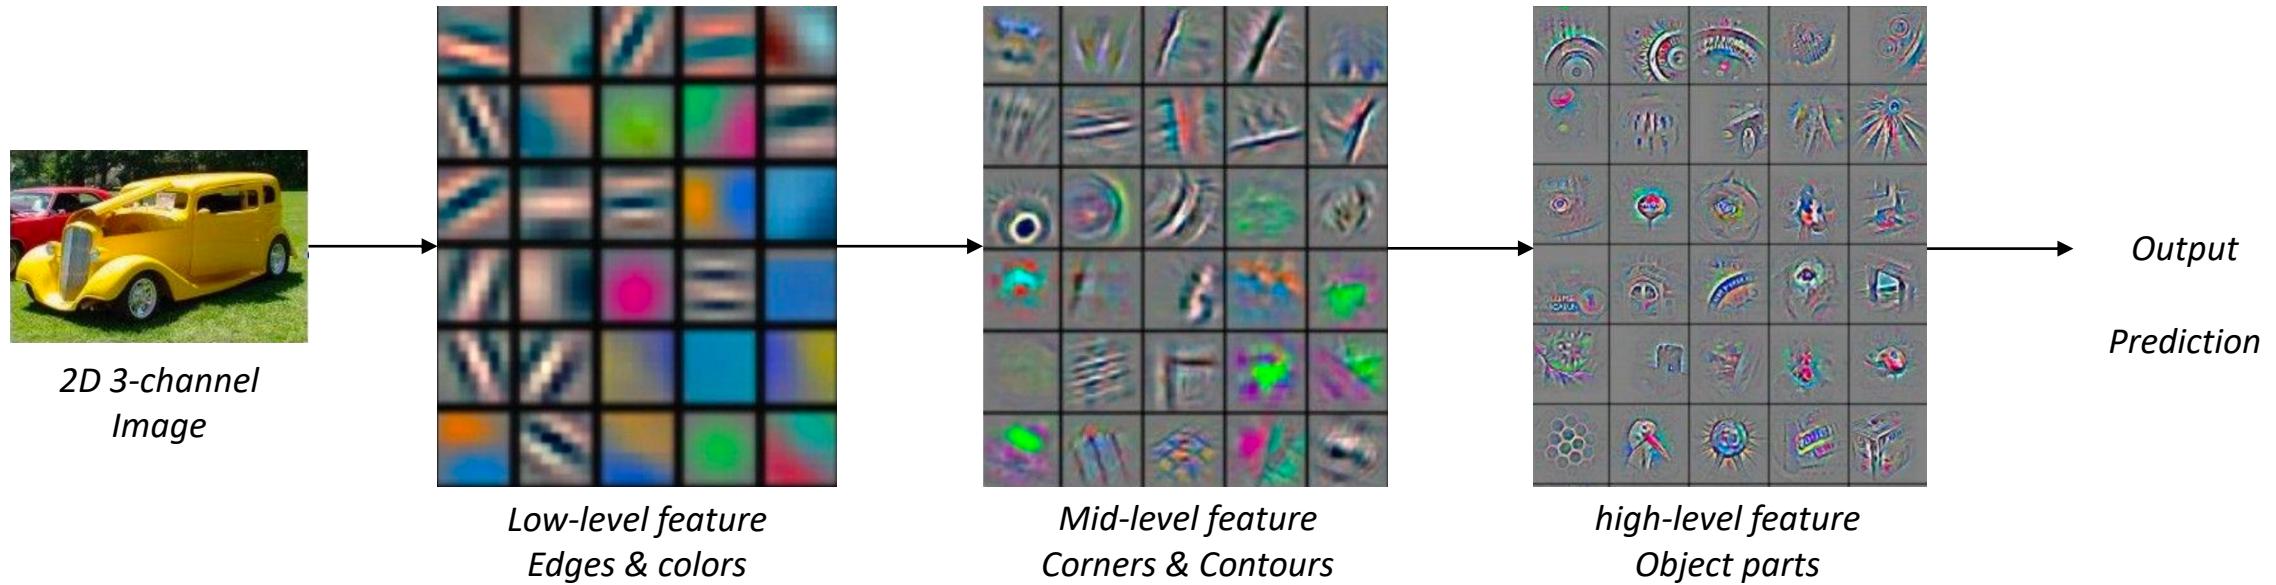
kedge.edu

EFMD EQUIS ACCREDITED AACSB ACCREDITED AMBA ACCREDITED



Recap of Previous Session

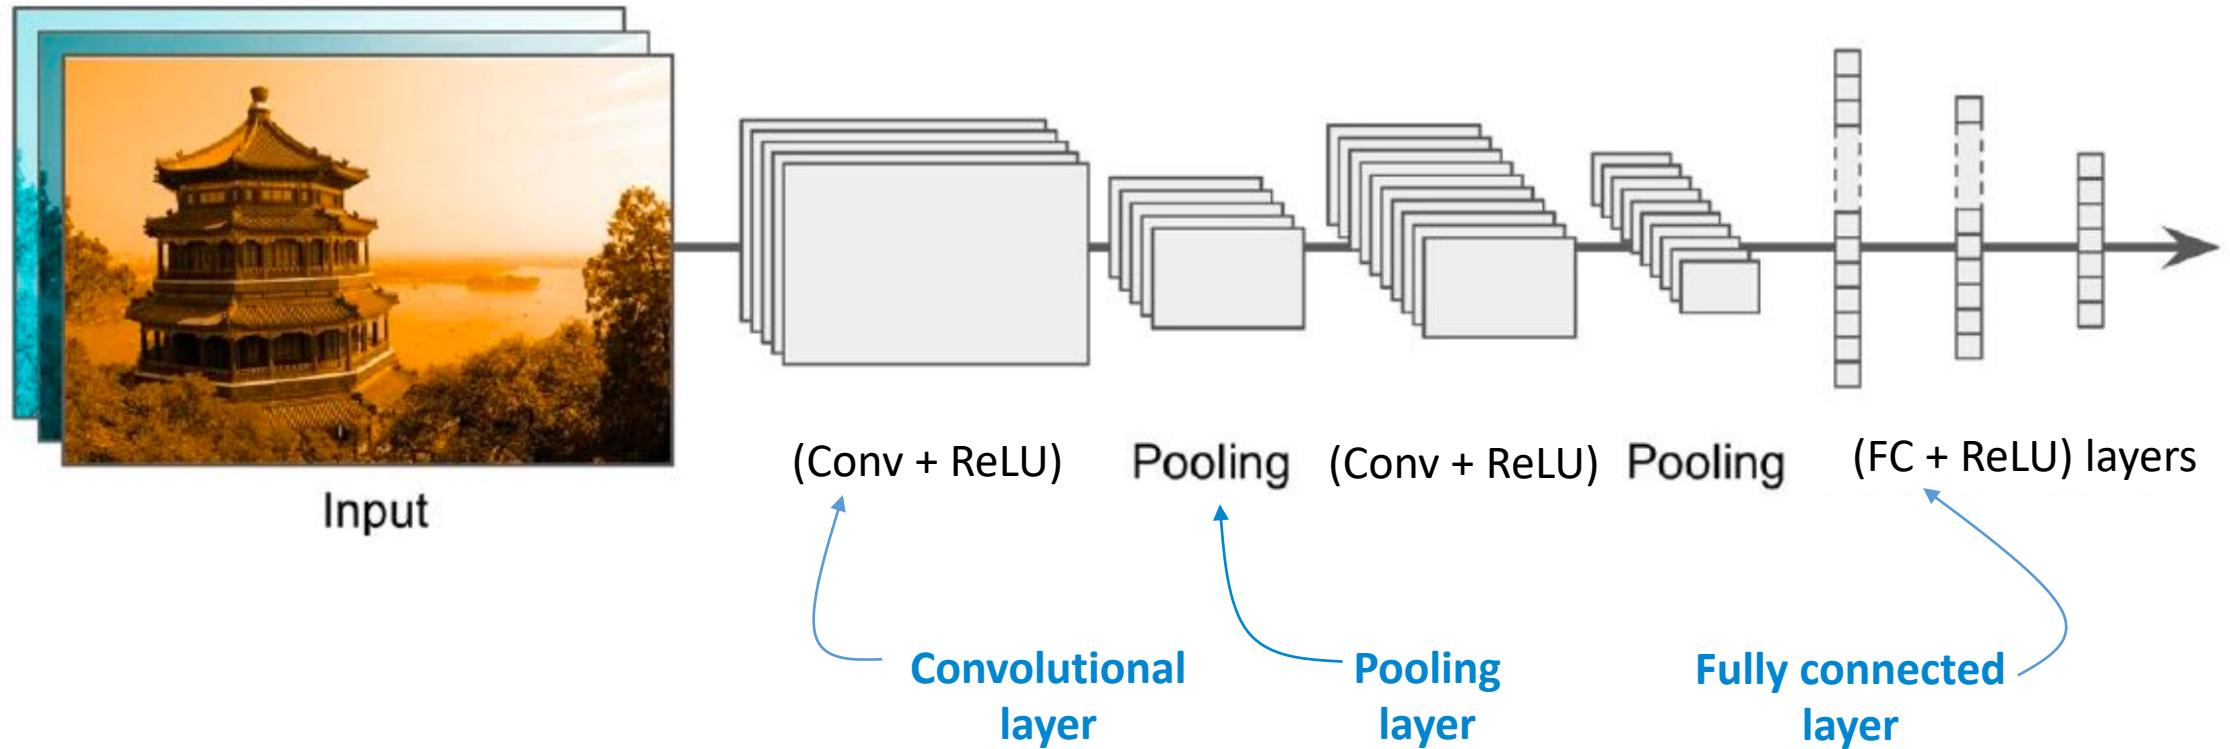
- **Convolution neural network**



- Detect local spatial patterns and extract them as feature maps by convolution
- Stack multiple convolutional layers to aggregate features hierarchically

Recap of Previous Session

- Typical CNN architecture



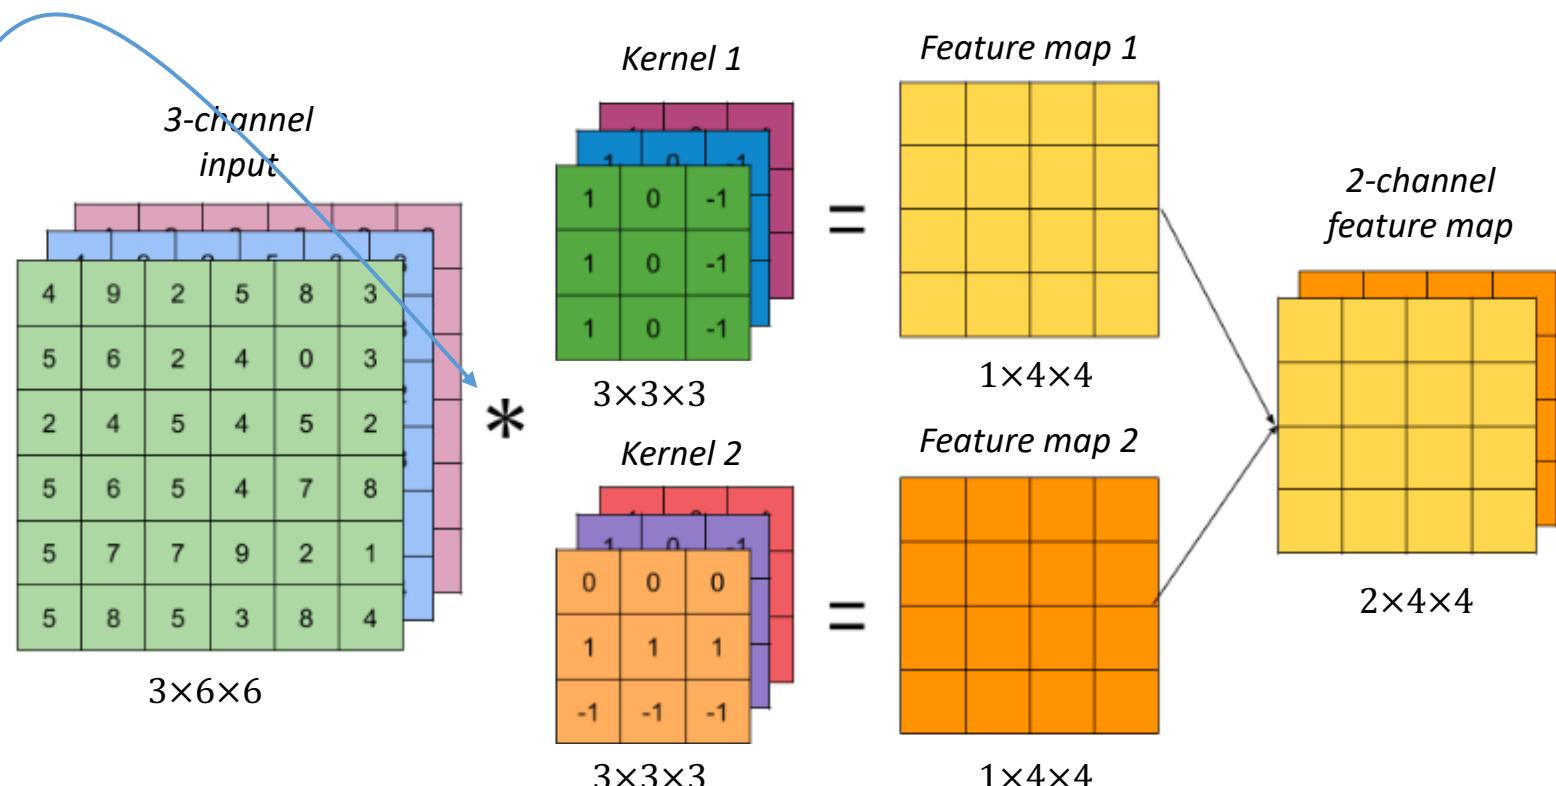
Recap of Previous Session

• Convolutional layer

- Multi-channel input image / feature map → Multi-channel output feature map

1. Convolution operation per input channel:

- Compute the sum of element-wise products in the kernel window
- Slide the kernel window across the input image / feature map
 - From left to right
 - From top to down
- Kernel size
- Padding
- Stride

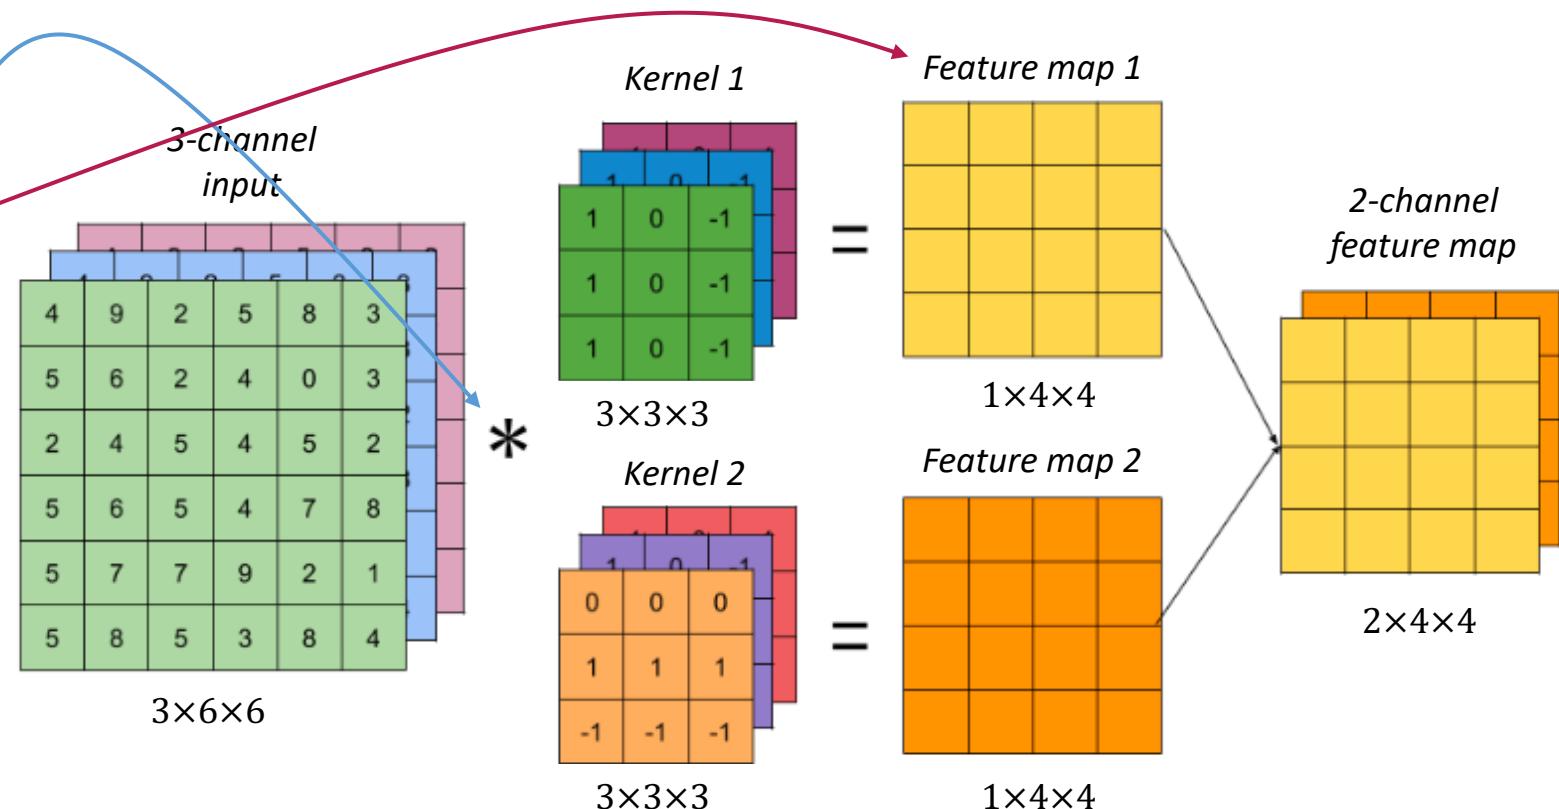


Recap of Previous Session

- **Convolutional layer**

- Multi-channel input image / feature map → Multi-channel output feature map

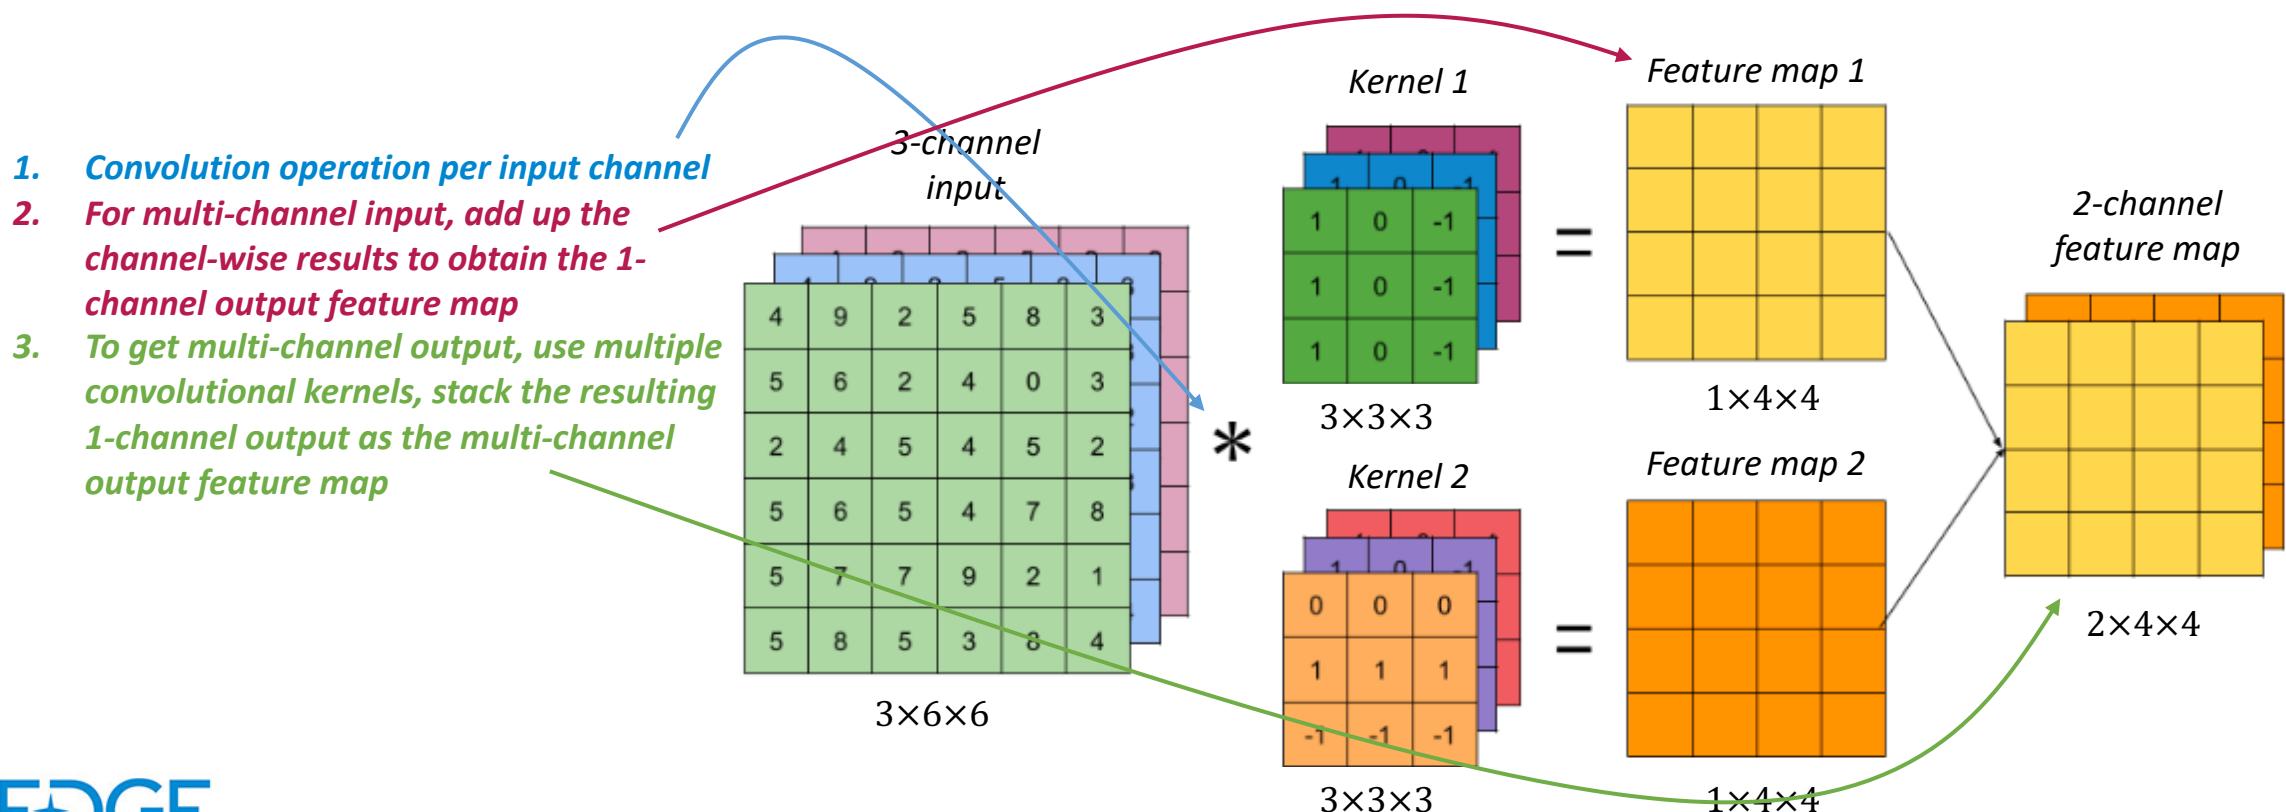
1. *Convolution operation per input channel*
2. *For multi-channel input, add up the channel-wise results to obtain the 1-channel output feature map*



Recap of Previous Session

- **Convolutional layer**

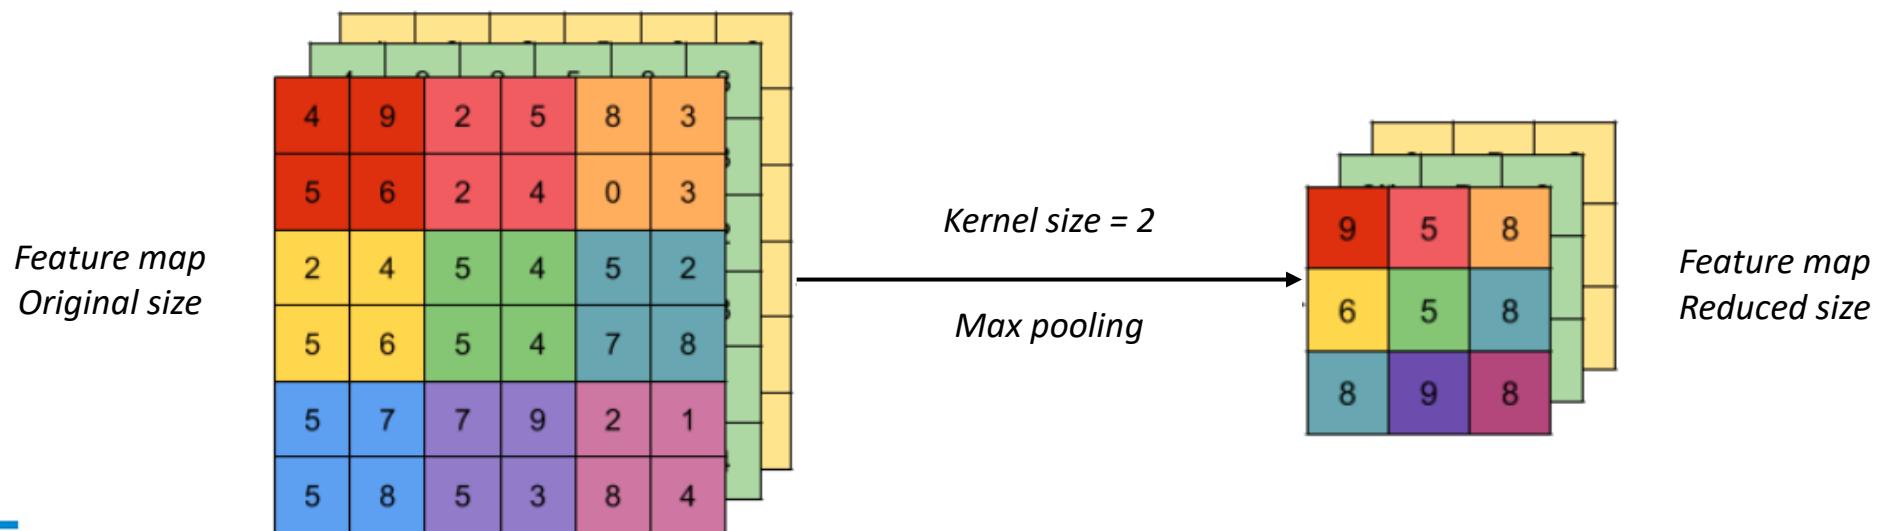
- Multi-channel input image / feature map → Multi-channel output feature map



Recap of Previous Session

- **Pooling layer**

- Reduce the size of feature map for extracting high-level features
- **Max Pooling**
 - Output the maximum value in the kernel window
 - Slide the kernel window across the input feature map (no overlap by default)



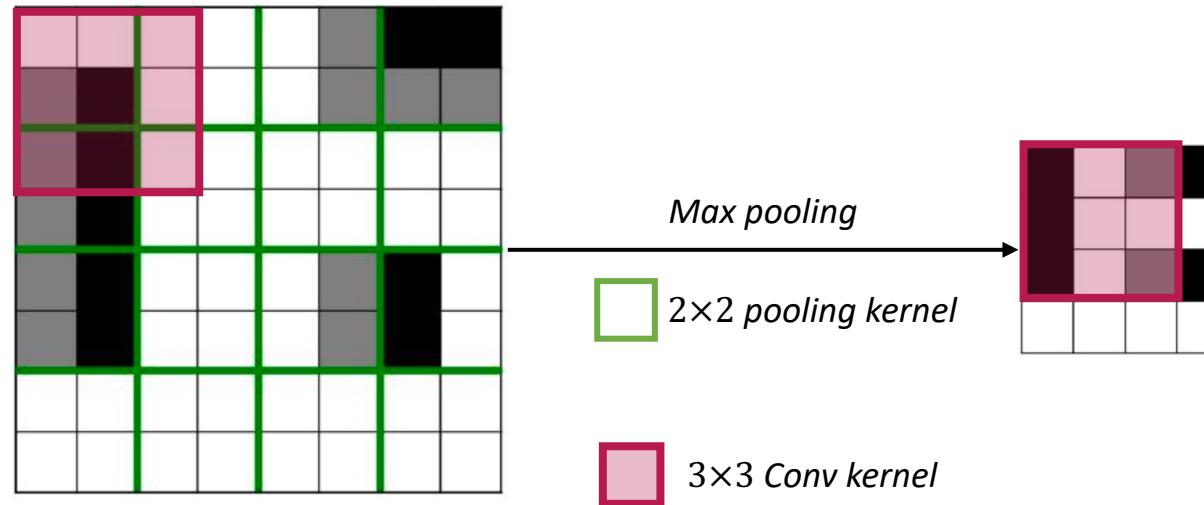
Recap of Previous Session

- **Pooling layer**

- Reduce the size of feature map for extracting high-level features

- **Max Pooling**

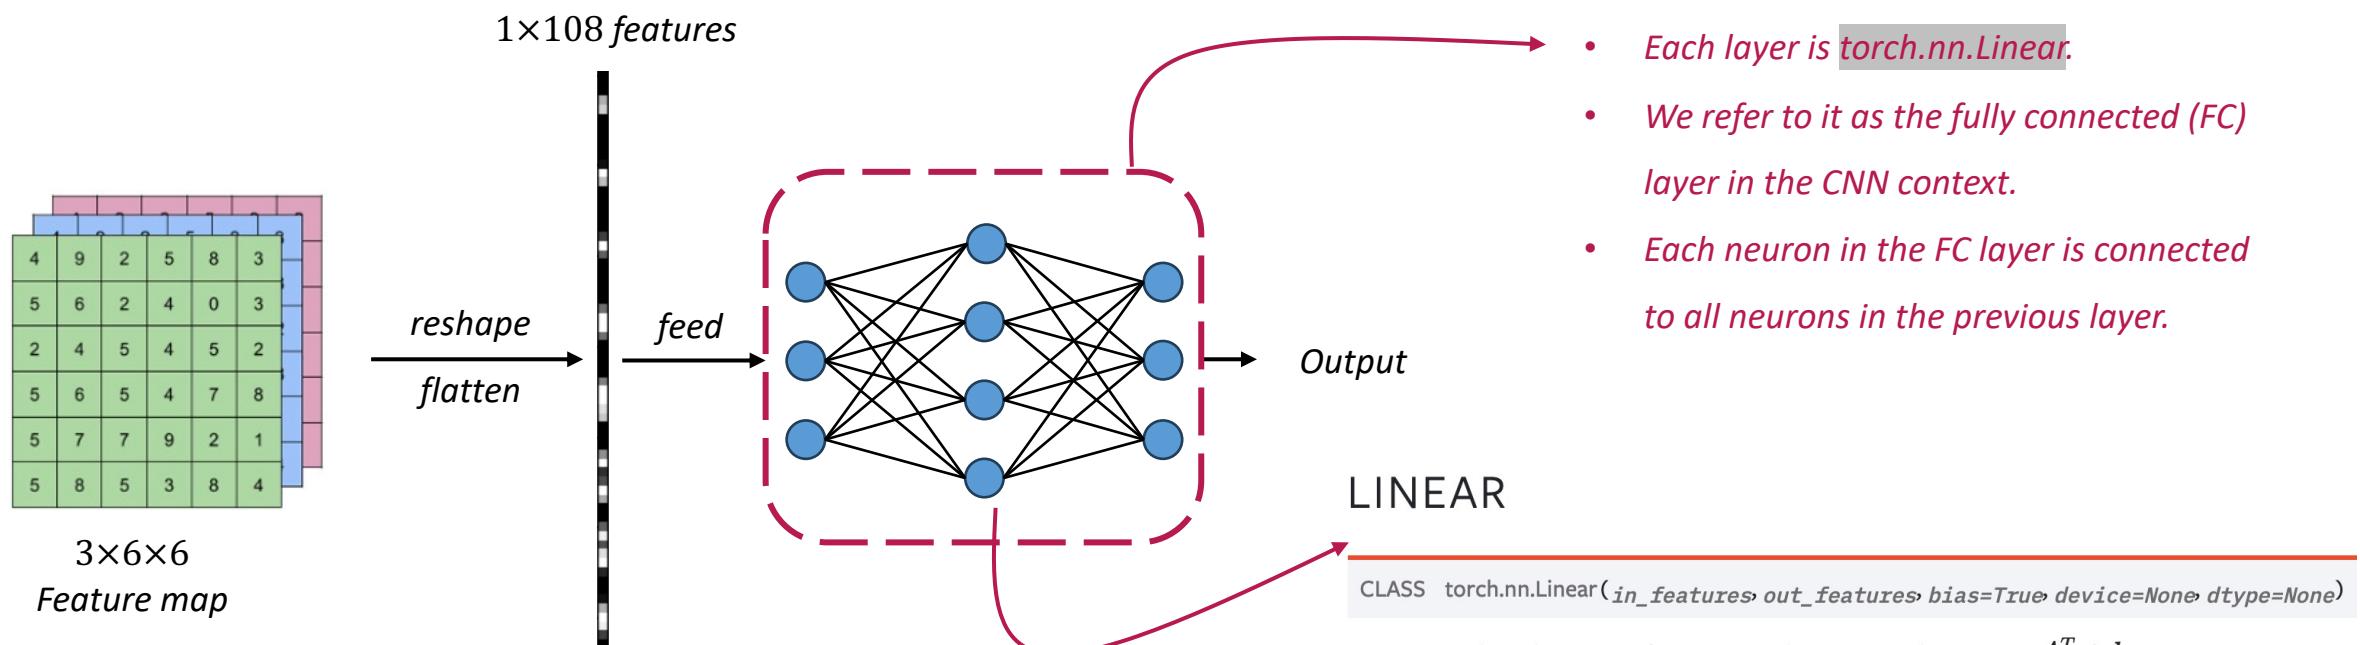
- Preserve the global spatial pattern, ignore the local one
- After pooling, the convolutional kernel can see a larger region to extract more global features



Recap of Previous Session

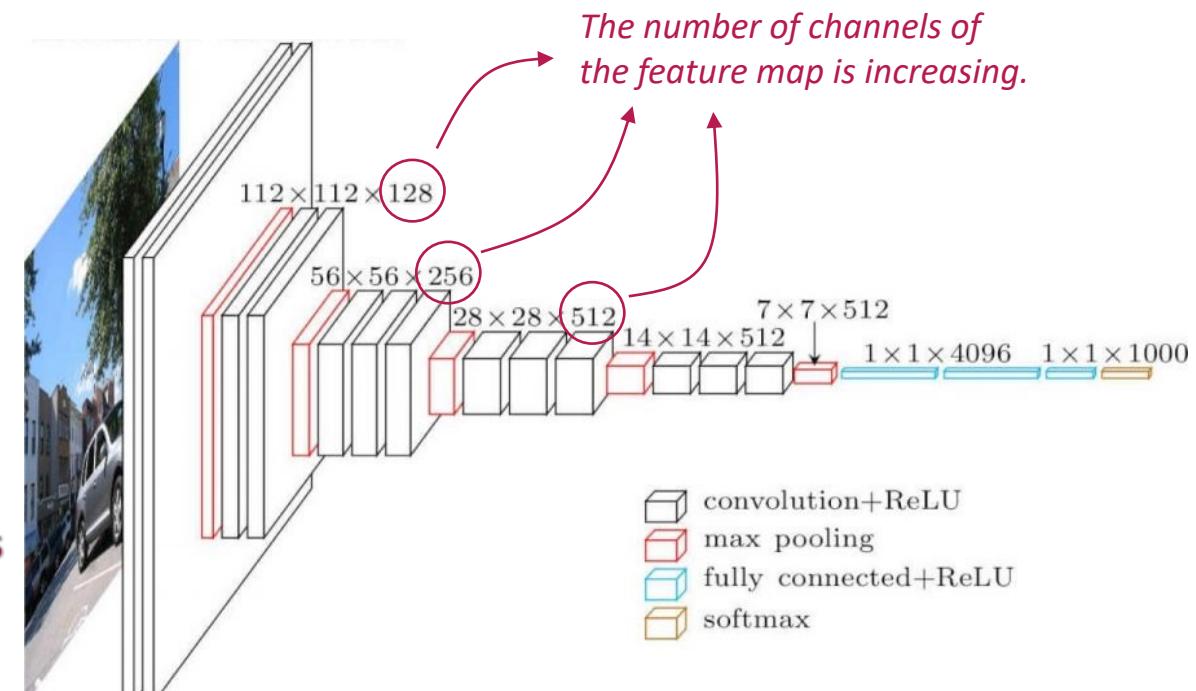
- **Fully connected layer**

- Flatten the feature map as 1D feature vector
- Use this 1D feature vector to make predictions



Recap of Previous Session

- **Pyramid architecture:**
 - As we go deeper in the CNN
 - **The size of feature map is shrinking** by convolution and pooling layers
 - Low-level feature → High-level feature
 - Local spatial pattern → Global spatial pattern
 - The number of channels of feature map is increasing
 - **Increase the number of convolutional kernels**
 - Corresponds to more possible combinations of simple features



Outline

- Transfer learning
- 1D Convolution neural network

Transfer learning

- Two cornerstones of deep learning
 - Deep neural networks
 - The deeper the model depth, the more learning parameters, the larger the capacity, and the stronger the representational power
 - Large-scale training datasets
 - Small datasets are not enough to train complex neural networks
 - Small datasets are not representative enough
 - The network can easily over-fit to the small dataset
 - Large-scale datasets are more representative and contain more general knowledge
 - Models trained on large-scale general datasets can also perform well on specific tasks
 - For example, ChatGPT

Model	GPT-1	GPT-2	GPT-3	GPT-4
Parameter count	117 million	1.5 billion	175 billion	1.7 trillion

Transfer learning

- Two cornerstones of deep learning

- Deep neural networks

- The deeper the model depth, the more learning parameters, the larger the capacity, and the stronger the representational power

- Large-scale training datasets

- Small datasets are not enough to train complex neural networks
 - Small datasets are not representative enough
 - The network can easily over-fit to the small dataset
 - Large-scale datasets are more representative and contain more general knowledge
 - Models trained on large-scale general datasets can also perform well on specific tasks
 - For example, ChatGPT

Training such networks can be very time-consuming.

Collecting and labeling such datasets are labor-intensive, and sometimes impossible for some specific tasks, for example, estimating the house price in a small village that only contains hundreds of houses.

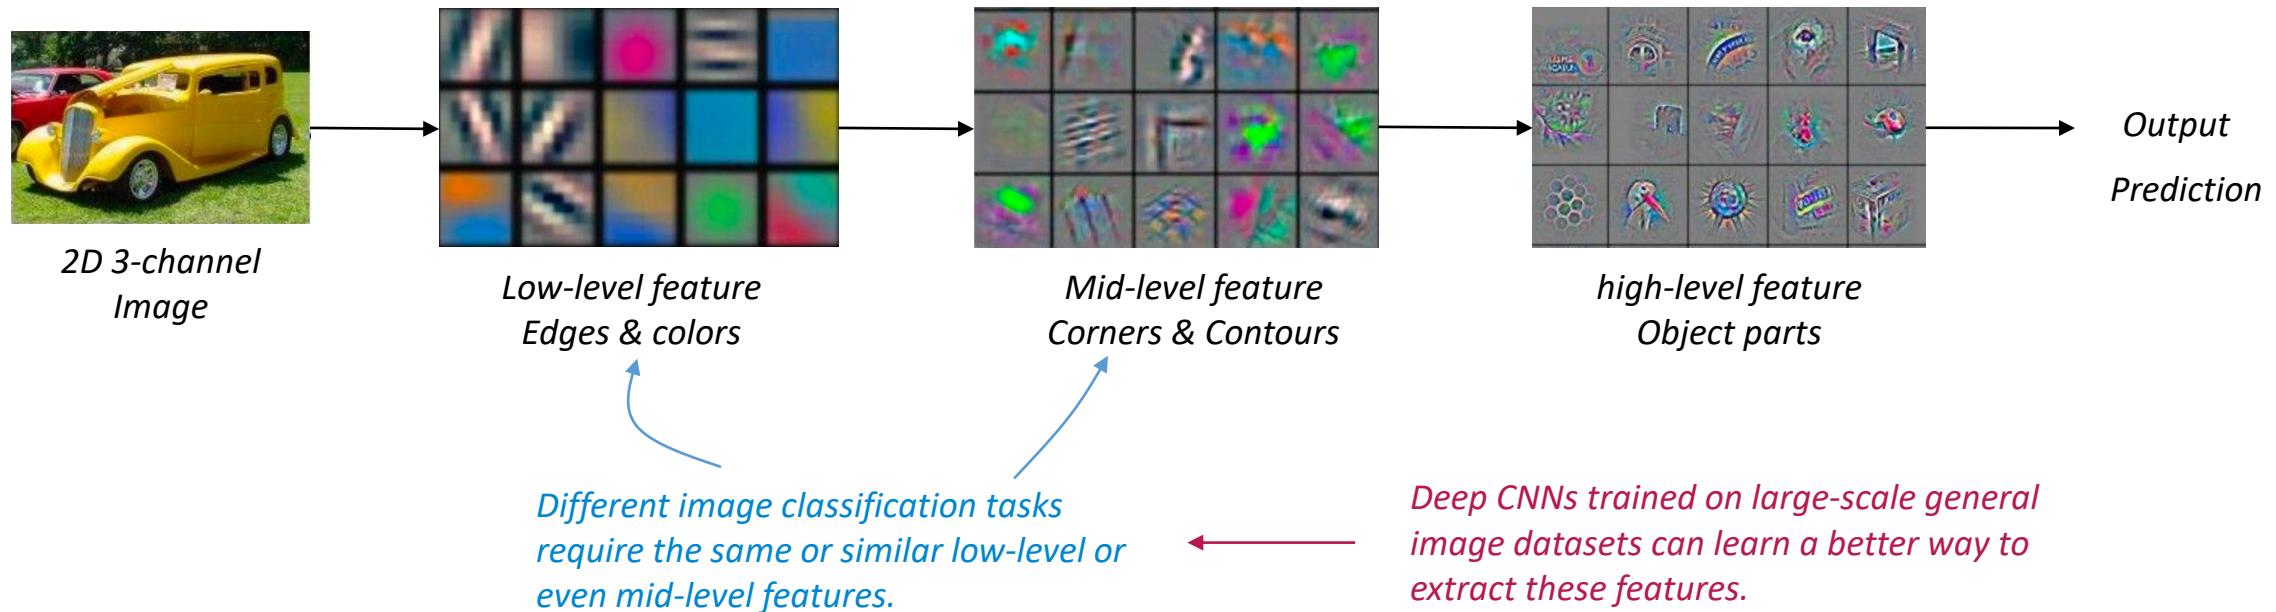
Can we re-use a deep neural network trained on a large-scale dataset and adapt it to a different task without requiring a full training process?

Transfer learning

- General definition of transfer learning:
 - A technique in machine learning in which knowledge learned from a task is re-used in order to boost performance on a related task
 - For example, knowledge gained while learning to recognize cars could be transferred to tasks such as recognizing trucks
- In the CNN context:
 - Transfer some knowledge from one neural network to another
 - Use a pre-trained model which has been trained on some large dataset
 - For example, ImageNet that contains more than 10 million images
 - Customized the pre-trained model to a specific task
 - Like distinguish between cat and dog

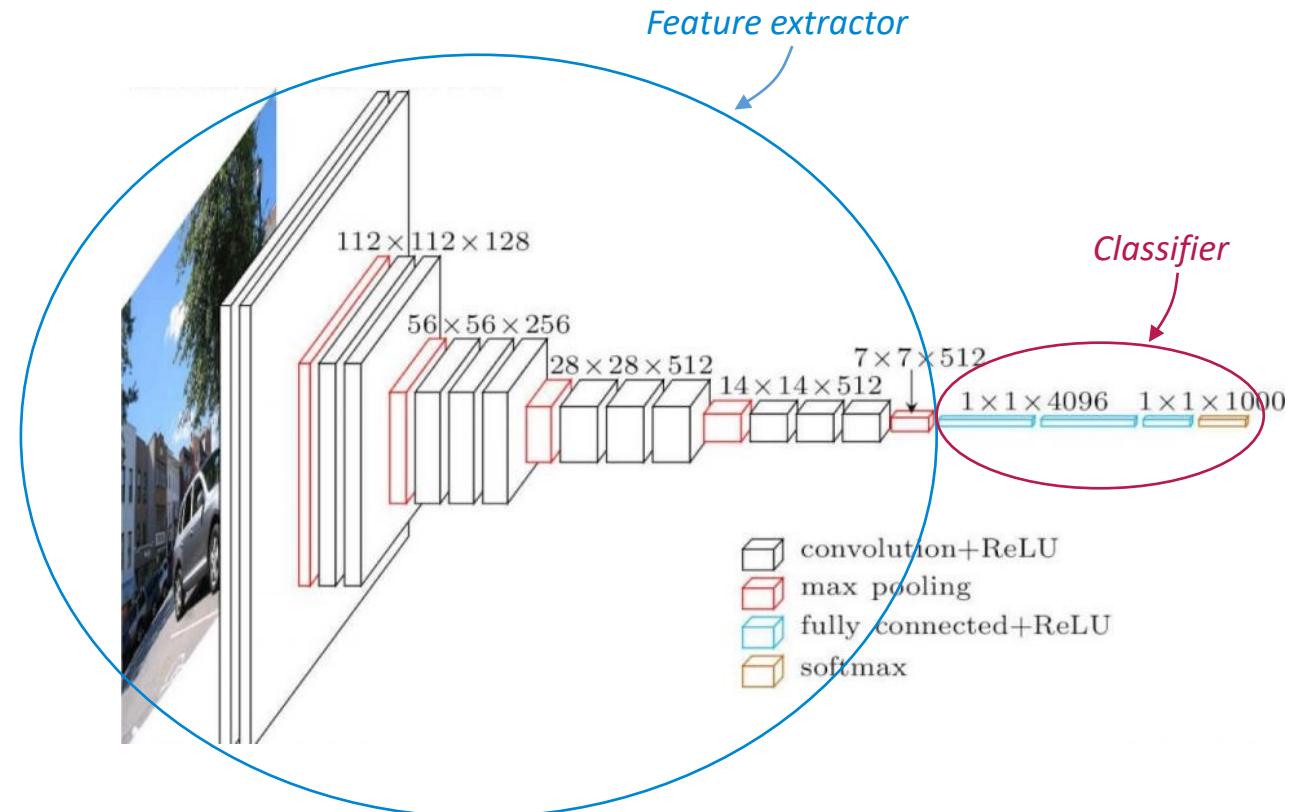
Transfer learning

- Rationale behind transfer learning
 - Domain-specific knowledge (like ML) is based on general knowledge regardless of domains (like mathematics, biology, statistics, etc.)



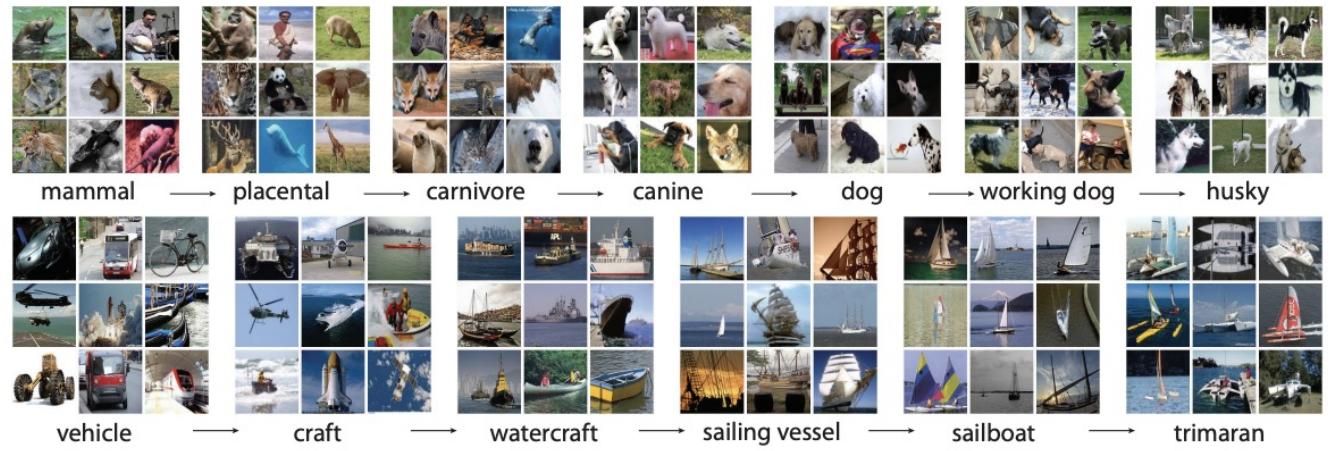
Transfer learning

- Rationale behind transfer learning
 - We can use the pre-trained CNNs to extract features and build a new classifier on top of these extracted features
 - Keep the feature extractor part unchanged
 - Re-build and re-train the classifier part
 - Only this classifier needs to be optimized, thus reduce the training time
 - (Optional) If we are not satisfied with the performance, we can then fine-tune the entire network with a small learning rate



Transfer learning

- Available resources
 - ImageNet
 - 1,281,167 training images
 - 50,000 validation images
 - 100,000 test images
 - 1,000 object classes
 - The object classes are based on the WordNet hierarchy
 - Pre-trained CNNs on ImageNet to predict the 1,000 classes
 - `torchvision.models`



Transfer learning

- Available resources
 - Pre-trained CNNs on ImageNet to predict the 1,000 classes
 - <https://pytorch.org/vision/stable/models.html#table-of-all-available-classification-weights>

Weight	Acc@1	Acc@5	Params	GFLOPS	Recipe
AlexNet_Weights.IMGNET1K_V1	56.522	79.066	61.1M	0.71	link
ConvNeXt_Base_Weights.IMGNET1K_V1	84.062	96.87	88.6M	15.36	link
ConvNeXt_Large_Weights.IMGNET1K_V1	84.414	96.976	197.8M	34.36	link
ConvNeXt_Small_Weights.IMGNET1K_V1	83.616	96.65	50.2M	8.68	link
ConvNeXt_Tiny_Weights.IMGNET1K_V1	82.52	96.146	28.6M	4.46	link
DenseNet121_Weights.IMGNET1K_V1	74.434	91.972	8.0M	2.83	link
DenseNet161_Weights.IMGNET1K_V1	77.138	93.56	28.7M	7.73	link
DenseNet169_Weights.IMGNET1K_V1	75.6	92.806	14.1M	3.36	link
DenseNet201_Weights.IMGNET1K_V1	76.896	93.37	20.0M	4.29	link
EfficientNet_B0_Weights.IMGNET1K_V1	77.692	93.532	5.3M	0.39	link
EfficientNet_B1_Weights.IMGNET1K_V1	78.642	94.186	7.8M	0.69	link
EfficientNet_B1_Weights.IMGNET1K_V2	79.838	94.934	7.8M	0.69	link
EfficientNet_B2_Weights.IMGNET1K_V1	80.608	95.31	9.1M	1.09	link
EfficientNet_B3_Weights.IMGNET1K_V1	82.008	96.054	12.2M	1.83	link

- AlexNet
 - ConvNeXt
 - DenseNet
 - EfficientNet
 - EfficientNetV2
 - GoogLeNet
 - Inception V3
 - MaxVit
 - MNASNet
 - MobileNet V2
 - MobileNet V3
 - RegNet
 - ResNet
 - ResNeXt
 - ShuffleNet V2
 - SqueezeNet
 - SwinTransformer
 - VGG
 - VisionTransformer
 - Wide ResNet
- The simplest one
- 

Outline

- Transfer learning
- **1D Convolution neural network**

1D Convolution neural network

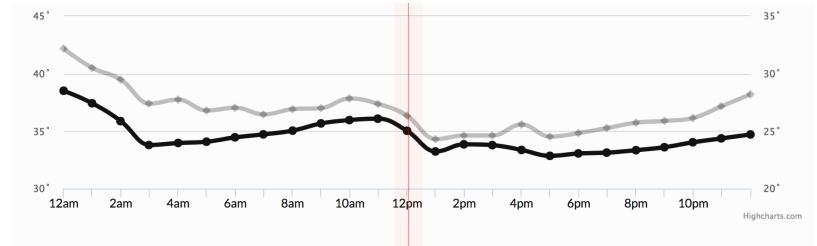
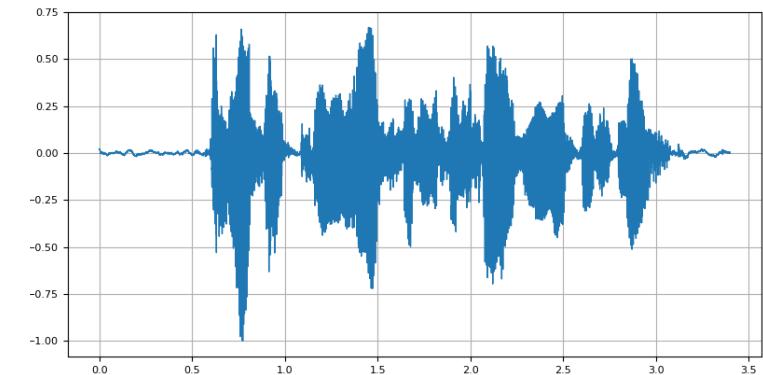
- Convolutional operation
 - Detect local spatial pattern / structure as feature
 - Not limited to 2D
- 1D Convolutional operation

<i>1D input</i>	<i>1D kernel</i>	<i>1D output</i>													
<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 25%;">1</td><td style="width: 25%;">3</td><td style="width: 25%;">2</td><td style="width: 25%;">6</td><td style="width: 25%;">4</td></tr></table>	1	3	2	6	4	*	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">-1</td><td style="width: 33.33%;">0</td><td style="width: 33.33%;">1</td></tr></table>	-1	0	1	=	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">1</td><td style="width: 33.33%;"></td><td style="width: 33.33%;"></td></tr></table>	1		
1	3	2	6	4											
-1	0	1													
1															
<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 25%;">1</td><td style="width: 25%;">3</td><td style="width: 25%;">2</td><td style="width: 25%;">6</td><td style="width: 25%;">4</td></tr></table>	1	3	2	6	4	*	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">-1</td><td style="width: 33.33%;">0</td><td style="width: 33.33%;">1</td></tr></table>	-1	0	1	=	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">1</td><td style="width: 33.33%;">3</td><td style="width: 33.33%;"></td></tr></table>	1	3	
1	3	2	6	4											
-1	0	1													
1	3														
<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 25%;">1</td><td style="width: 25%;">3</td><td style="width: 25%;">2</td><td style="width: 25%;">6</td><td style="width: 25%;">4</td></tr></table>	1	3	2	6	4	*	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">-1</td><td style="width: 33.33%;">0</td><td style="width: 33.33%;">1</td></tr></table>	-1	0	1	=	<table border="1" style="width: 100%; border-collapse: collapse;"><tr><td style="width: 33.33%;">1</td><td style="width: 33.33%;">3</td><td style="width: 33.33%;">2</td></tr></table>	1	3	2
1	3	2	6	4											
-1	0	1													
1	3	2													

1D Convolution neural network

- Use traditional ANN when features are isolated
 - The order of features doesn't matter
 - For example, the flight price dataset
 - Airline name, source city, departure time, stops, ...
- Use 1D CNN when the order of features matters
 - There are local spatial patterns between adjacent features

Audio data → 1D CNN for speech recognition



Time series data → 1D CNN for predicting temperature / stock price ...

1D Convolution neural network

- The architecture of 1D CNNs is the same as the 2D ones
 - `torch.nn.Conv2d` → `torch.nn.Conv1d`
 - `torch.nn.LazyConv2d` → `torch.nn.LazyConv1d`
 - `torch.nn.MaxPool2d` → `torch.nn.MaxPool1d`
 - `torch.nn.BatchNorm2d` → `torch.nn.BatchNorm1d`
 - `torch.nn.LazyBatchNorm2d` → `torch.nn.LazyBatchNorm1d`

Hands-on Exercise

- **Exercise 07 CNN & Transfer Learning - Instruction**
- Exercise 07 CNN & Transfer Learning - Assignment